

# Integrating RISC-V SIMT and Scalar Cores: Loosely to Tightly Coupled

**Presenter:** Clay Hughes, Sandia National Labs

**Authors:** Sooraj Chetput, Anusuya Nallathambi

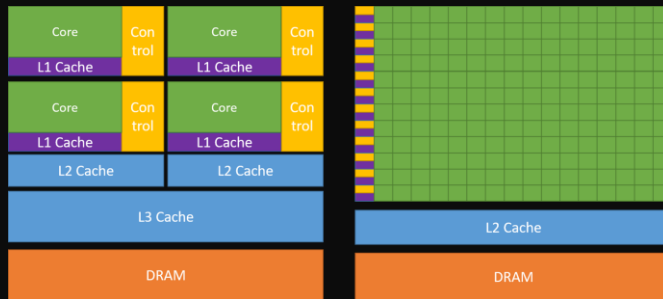


Sandia  
National  
Laboratories



# Introduction: SIMT Architectures


- Single Instruction Multiple Thread
- Computing model designed to handle highly parallel workloads.
- Highly optimized for overall throughput over thread latency.
- **Lockstep Synchronization:** threads within a warp (a collection of threads) execute same instruction simultaneously, all threads progress together.



CPU

GPU/SIMT

Source: <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>



# Introduction: **Vortex** GPGPU

- Open-Source RISC-V GPU, written in SystemVerilog
- **Instruction Handling:** Treats RISC-V scalar instructions as standard SIMT instructions.
- **Simplified Code Conversion:** Code requires only recompilation, not a complete rewrite.
- Custom Instructions to Support a SIMT Stack
  - + **WSPAWN:** Spawn a new warp
  - + **SPLIT/JOIN:** Used to handle the immediate post dominator stack for managing control flow divergence
  - + **TMC:** Thread Mask Control to control which threads within a warp are active.
  - + **CSRR/CSRW:** Special Instructions used to write to the control status registers.

## Source:


- + **Vortex: An Open Source Reconfigurable RISC-V GPGPU Accelerator for Architecture Research**
  - + *Authors: Fares Elsabbagh, Blaise Tine, Apurve Chawda, Will Gulian, Yaotian Feng, Da Eun Shim, Priyadarshini Roshan, Ethan Lyons, Lingjun Zhu, Sung Kyu Lim, Hyesoon Kim*
  - + *Hot Chips 32 (2020)*

# Problems with SIMT Architectures:

## Control Flow Divergence

- Occurs when threads within a warp have different execution paths through the code
  - Happens if there are if statements or loops in the instruction streams
- Breaks the [Lockstep Synchronization](#)
- This is a problem for the SIMT model because the execution of threads within a warp becomes serialized


# Existing Solution: Immediate Post Dominators




- The above figure demonstrates how code translates to basic blocks (BB).
- We can say that BB4 immediately post dominates BB1
- Let's see what happens in the hardware if thread 0 and 2 decide to go through BB0 and threads 1 and 3 decide to go through BB3




BB1: Everything executes in Lockstep



BB2: Thread 0, 2 is masked off, to allow Threads 1 and 3 to execute.



BB3: Thread 1, 3 is masked off to allow thread 0 and 2 to execute



BB4: All threads come back together

# Problem Statement

- **Control Flow Divergence:**
  - Leads to serialization of execution paths, increasing total execution time
  - **Post Dominators:** aligning execution by merging paths at a control point. This is implemented in the hardware using a stack.

## Previously Proposed Solutions

- **Dynamic Warp Reshuffling:**
  - Proposed solution for maximizing SIMT lane utilization
  - Re-shuffling warps and rearranging threads based on Control Flow Divergence Patterns
  - Not used due to hardware overhead
- **2-way/N-way Lane Splitting:**
  - Suggests allowing parallel execution of divergent paths
  - Not used due to hardware overhead.

# The idea behind our **Solution**

- We have 2 cores: one SIMT, one Scalar.
- Both cores run the same extended RISC-V ISA,
  - Thread transfer shouldn't require any recompilation or translation

**If we can send the divergent threads to the Scalar Core, and the less divergent ones to the SIMT Core, we can effectively reduce serialization on the SIMT Core.**

# Roadmap

- 1 The Vortex Tapeout: Loose Integration with a Scalar core
- 2 SIMT-Scalar V1 and V2 Architectures for mitigating Control Flow Divergence: Tight Integration
- 3 Future Works



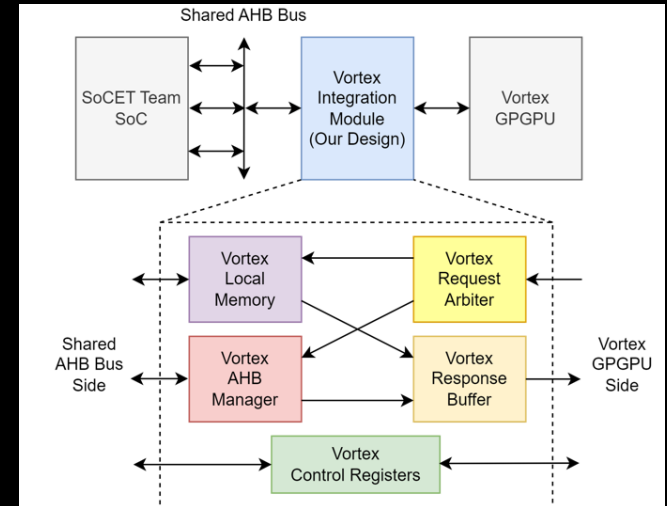
# The Vortex Tapeout

- **Integration Overview**

- Integration follows a host-device CPU-GPU model on the AHB bus.
- Enhances a RISC-V based microcontroller with parallel processing capabilities.

- **Software Interface**

- Three control/status registers: Start, Start Address, and Status Register.
- CPU independently handles tasks while GPU is active, enhancing multitasking.




# The Vortex Tapeout Contd.

- **Results and Physical Design Insights**


- We used flip-flops for memory due to PDK constraints.
- Memory consumption is significant, impacting overall area efficiency.
- The table below shows 4 configurations that we considered.

**Table 1.** Netlist Configuration Summary

Netlist Configuration Summary				
Metrics	Config 1	Config 2	Config 3	Config 4
Cores-Threads-Warps	1-4-4	1-4-4	1-4-4	1-4-2
i-cache size	1kB	2kB	1kB	1kB
d-cache size	2kB	4kB	2kB	4kB
Shared memory	1kB/thread	1kB/thread	512B/thread	1kB/thread
Local memory	4kB	4kB	4kB	4kB
FPU	No	Yes	Yes	Yes
Area estimate	29mm <sup>2</sup> (Netlist)	34mm <sup>2</sup> (PD)	20mm <sup>2</sup> (Netlist)	21mm <sup>2</sup> (PD)




SoC+Vortex Amoeba View



SoC+Vortex Physical View

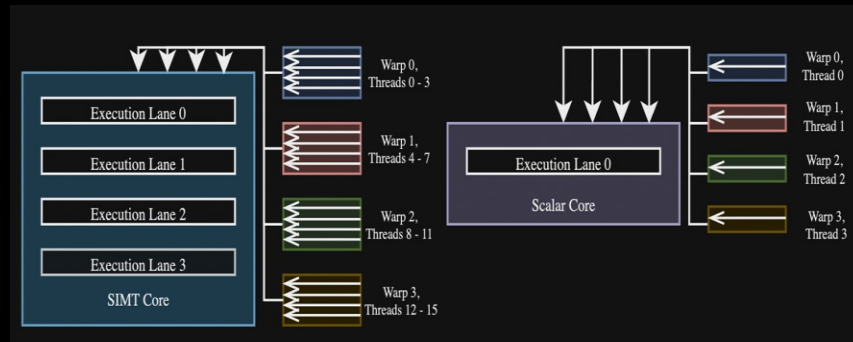
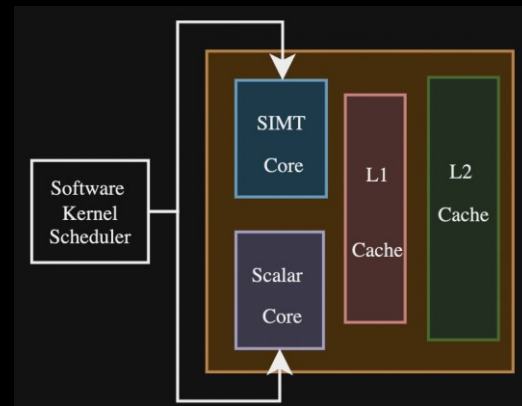
# The SIMT-Scalar **V1** and **V2** Architectures for mitigating **Control Flow Divergence**

- **V1 architecture:** a simple in-order Scalar core created by parameterizing Vortex's SIMT core
- **V2 architecture:** Scalar core with logic to transfer threads from the SIMT core during runtime + simple not-taken branch predictor to reduce stalling of threads



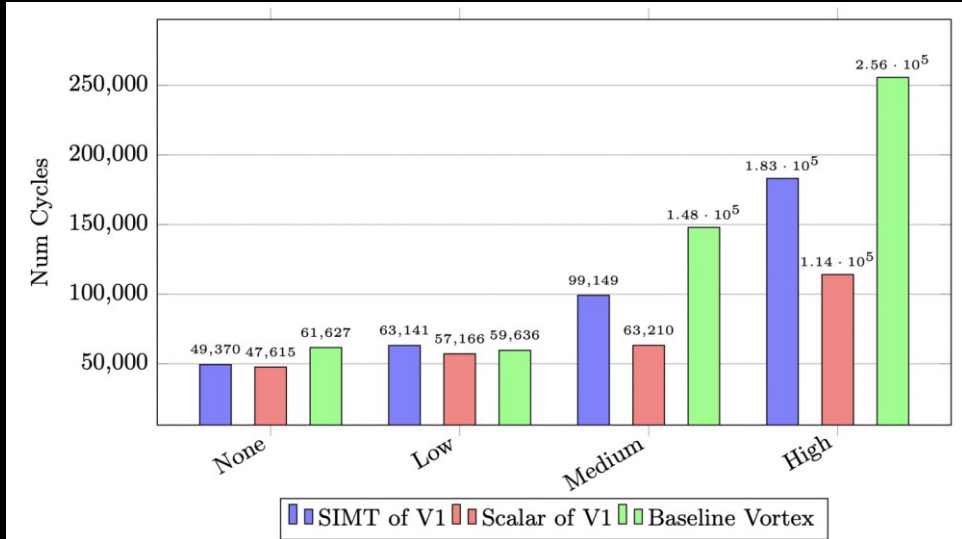
# SIMT-Scalar V1 Architecture

- Kernel scheduler assigns threads separately to the cores according to priority.
- Priority is arbitrarily assigned by software.
- Low priority -> SIMT core ( 4 threads/warp) (4 warps)
- High priority -> scalar core ( 4 hyperthreaded threads )
- No Datapath hardware is shared between the SIMT and the Scalar core.
- L1 and L2 caches are shared
- Modified the Vortex GPGPU to implement this architecture
  - Scalar core is a parametrized single-lane version of the SIMT core



# V1 Architecture – Simulations, Results

- Tested with simple vector add type kernels
- We ran a total of 20 threads on The V1 Arch and the baseline vortex.
- 4 different degrees of divergence
  - **No divergence:**
    - All threads Add
  - **Min Divergence:**
    - odd threads Add
    - even threads Subtract
  - **Mid Divergence:**
    - Odd threads: 4th bit == 1? Add else Subtract
    - Even threads: Subtract
  - **High Divergence**
    - Complete serialization of all the threads
    - (each thread takes its own path)



**Fig. 10.** Num Cycles Comparison for SIMT and Scalar Cores of V1 Architecture Across Divergence Levels vs. the Baseline Vortex.

# ***Discussion: V1 Architecture***

- **Pros**


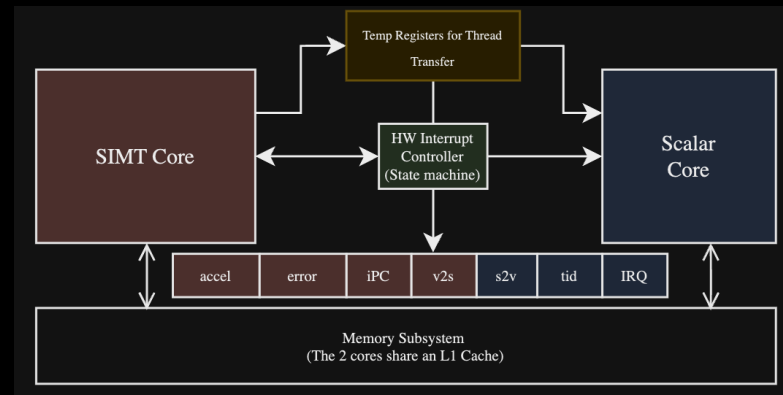
- Simple design, ideal for proof-of-concept.

- **Cons**

- We don't know which threads are more divergent and thus should be marked as high priority before run-time. This process is implemented arbitrarily by the software currently.
- The Scalar core is not optimized for branches (e.g. it stalls every time it runs into a branch).

# SIMT-Scalar V2 Architecture

- **SIMT Core: 4 Warps, 4 threads/warp**
- **Scalar Core: 1 Thread, 1 Warp**
- **Coupling**
  - § At the hardware interrupt controller level (shared control status registers)
- **Thread Priority Assignment**
  - § Software still assigns a relative priority index to each thread (does not classify it as high priority or low priority like V1 did).
  - § The Scalar core kernel scheduler uses this information to decide what thread to pull from the SIMT core.
- **Dynamic Switching:**
  - § Introducing hardware interrupt controller to facilitate dynamic thread transfer from the SIMT to the scalar upon the scalar core's request
- **Memory Subsystem**
  - § Shared L1 Instruction/Data cache for scalar and SIMT cores, Unified L2 cache
- Shared "status registers" inside of interrupt controller to allow communication between SIMT and scalar cores
  - § Thread context is dumped into dedicated control status registers



# V2 Arch - Discussion

- This still does not solve the problem of how to select the thread that is to be transferred in an intelligent way.
- All we did here is build a mechanism to transfer the thread from the SIMT to the Scalar core.
- Consider a case when none of the threads divergence
  - It would be best to execute the entire thing on the SIMT
- Consider when all the threads are completely serialized
  - It would be best to just execute the entire thing on the Scalar core.
- The development of the RTL is still under progress, and so is benchmarking.



# Future Work

- Complete V2 Architecture and Benchmarking
- V3 Architecture: We can use the IPDOM analysis, and branch history to predict branch divergence. The threads that will cause the most serialization will be sent to the scalar core.
- V4 Architecture: We want to consider a cluster of M SIMT and N Scalar cores. This could be better for handling more complex divergence patterns.

# Purdue SoCET Team

- The goal of Purdue SoCET (system-on-chip extension technologies) is to provide students hands on experience with a fully developed industry quality SoC design flow. Members of the group engage with RTL design, physical design, chip bringup, verification methods, an array of EDA tools and software development.
- Checkout the cool work the team does at the following link:
  - <insert QR code for this URL: <https://engineering.purdue.edu/SoC-Team#about>>
-