

# Advancing the RISC-V Performance Simulation Ecosystem with Data Prefetching

**Author:** Luís Crespo

**Email:** [luís.miguel.crespo@tecnico.ulisboa.pt](mailto:luís.miguel.crespo@tecnico.ulisboa.pt)

**Co-authors:**

Nuno Neves

Pedro Tomás

Nuno Roma



# Overview

- 1. Motivation**
- 2. Background**
- 3. Proposed Simulation Flow**
- 4. Prefetcher Module**
- 5. Prefetching Integration**
- 6. Evaluation**
- 7. Conclusion**

# Motivation

- **The adoption of RISC-V for HPC**
  - Performance improvements for HPC
  - Performance Modelling (and Simulation) SIG
  - Olympia: performance model of an OoO Superscalar RISC-V Core
- **Olympia lacks support for:**
  - Data prefetching mechanisms, crucial technique for mitigating the “memory wall”
  - Flexible and full ISA trace generation (currently generated with an emulator restricted to RV64GC)
- **Contributions:**
  - Parser to convert Spike-generated traces into Olympia’s format
  - Integration of data prefetching mechanisms into Olympia
  - Implementation of multiple prefetching algorithms within Olympia

# Background: Simulation

- **Functional Simulation**

- Simulate only the outcome of instructions.
- No cycle-level detail; fast and lightweight.



- **Spike**

- Functional simulator
- Golden reference RISC-V simulator

- **Timing Simulation**

- Evaluate processor design choices.
- Identify bottlenecks and estimate IPC

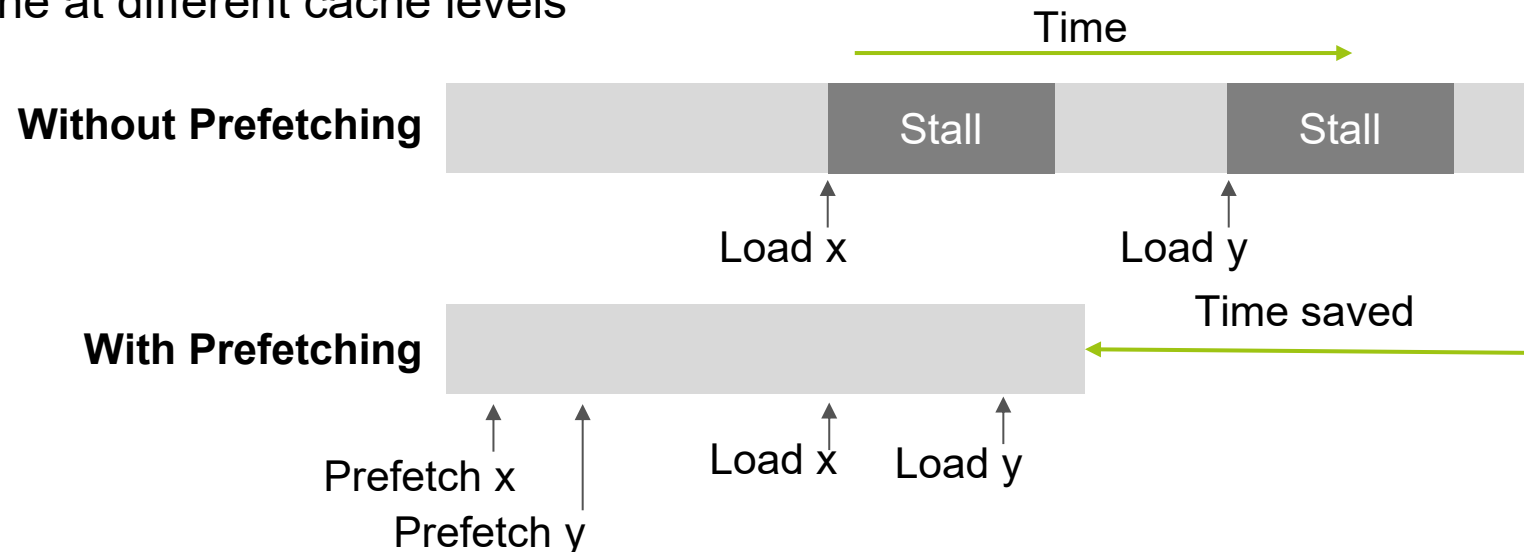


- **Olympia**

- Trace-driven timing simulator
- Reads Simulation Trace Format (STF) or JSON
- Built on Sparta framework
- Simulate OoO Superscalar RISC-V Cores
- Designed for RISC-V performance modeling

# Background: Data Prefetching

- Software Prefetching
  - Compiler or programmer inserts explicit prefetch instructions
- **Hardware Prefetching** (focus of this work)
  - Dedicated module that detects memory access patterns and issue prefetch requests
- Proactively load data into caches before the processor demands it
- Reduces stall cycles and improving throughput
- Can be done at different cache levels



# Background: Data Prefetching

- **Prefetching Evaluation Metrics**


- Issued prefetches
- Demand misses
- Useful prefetches
- Useless prefetches
- Late prefetches

- $Accuracy = \frac{Useful}{Issued}$


- $Coverage = \frac{Useful}{Useful + demand\ misses}$

- **Aggressiveness parameters:**

- **Distance** - how far ahead to prefetch
- **Degree** - number of blocks fetched



# Proposed Simulation Flow



## Integrating Spike and Olympia

### Trace Generation:

- Custom STF parser
- Converts Spike logs into STF
- Extract:
  - PC (next PC)
  - Instruction type (load/store)
  - Encoding
  - Memory address
  - Data size
  - Instruction width


**START**

0x10264	(0x00004033)		
0x10268	(0x4629) x12	0x0000000000000000a	<b>Memory Address</b>
0x1026a	(0x0005a007) f0	0xffffffffc2ac75c3 mem	0x00000003fffffb10 4
0x1026c	(0x00052087) f1	0xffffffff00000000 mem	0x00000003fffffae8 4
<b>Encoding</b>	(0x08a07043) f0	0xffffffffc3815852	
0x10270	(0x00052027) mem	0x00000003fffffae8 0xc3815852	<b>4 Data size</b>
0x1027a	(0x0511) x10	0x00000003fffffaec	
0x1027c	(0x167d) x12	0x00000000000000009	
0x1027e	(0x0591) x11	0x00000003fffffb14	
<b>PC</b>	0x10280	(0xf66d)	
0x1026a	(0x0005a007) f0	0xffffffff4252999a mem	0x00000003fffffb14 4
...			
0x10282	(0x0010c033)		
	<b>END</b>		

**Spike log**

# Prefetcher Module


- Implemented as a standalone module
  - Can be instantiated at multiple cache levels (L1D, L2)
  - Easily extendable to support other algorithms
- Prefetch Engine:
  - Generates prefetches based on the selected algorithm
- Prefetch Queue:
  - Buffers pending requests
- Parameterizable
  - Prefetch Degree, Prefetch Distance, Prefetch Store addresses, etc
- Three prefetching algorithms were implemented:
  - Next-Line prefetcher
  - Stride prefetcher
  - Best-Offset prefetcher







# Prefetcher Integration

- One prefetcher module for each cache level
- By default, both the L1 and L2 prefetchers are disabled
- Activate a prefetcher in the architecture configuration file
- Each prefetcher snoops its cache requests and fill responses
- Squash unnecessary requests in the cache
  - Cache line is already present
  - Already requested (allocated entry in the MSHR)



# Prefetcher Integration



## L1D-Cache:

- Three stage cache pipeline:
  - Arbitration and Lookup
  - Data Read
  - Deallocation
- Extended to non-blocking
- Notifies the prefetcher to send requests

## L2-Cache:

- Configurable size pipeline:
  - Arbitration
  - Lookup (several cycles)
  - Data Read
- Prefetchs arbitrated with other requests (lowest priority)

# Evaluation: Setup


CPU	Functional Units	Caches
RV64GC 1 OoO core 8-wide fetch/issue/commit 8 LSQ, 30 ROB	5 Int ALU (1 cycle) 2 Int Mult (3 cycles) 1 Int Div (23 cycles) 2 FP ALU (2 cycles) 2 FP MUL/DIV (4/63 cycles)	<b>L1-I:</b> - 16KB / 2-way - 8 MSHRs / 2-cycle latency <b>L1-D:</b> - 64KB / 2-way - 8 MSHRs / 2-cycle latency <b>L2:</b> - 512KB / 16-way - 16 MSHRs / 10-cycle latency

- CPU configuration available for Olympia (big\_arch)
- 8 applications from the MachSuite
- Compiled using Clang with -O3
- Spike and proposed STF parser to generate traces
- Different prefetcher combination for each cache level

Prefetcher	L1D Cache	L2 Cache
Next-line	✓	✓
Stride	✓	✗
Best-Offset	✗	✓


# Evaluation: L1D Cache

- Speedup (over the baseline)
- Accuracy, and Coverage



# Evaluation: L1D Cache


- Speedup (over the baseline)
- Accuracy, and Coverage
- Highest speedups:
  - BFS (bulk & queue)
  - SPMV (ellpack & csr)
  - No spatial locality benefit






# Evaluation: L1D Cache

- Speedup (over the baseline)
- Accuracy, and Coverage
- Highest speedups:
  - BFS (bulk & queue)
  - SPMV (ellpack & csr)
  - No spatial locality benefit
- High accuracy w/ low coverage:
  - GEMM (ncubed)
  - Stencil (2D)
  - Memory hierarchy is overloaded




# Evaluation: L1D Cache

- Speedup (over the baseline)
- Accuracy, and Coverage
- Highest speedups:
  - BFS (bulk & queue)
  - SPMV (ellpack & csr)
  - No spatial locality benefit
- High accuracy w/ low coverage:
  - GEMM (ncubed)
  - Stencil (2D)
  - Memory hierarchy is overloaded
- Stride vs Next-line in Gemm (blocked)



# Evaluation: L2 Cache

- Same metrics as before
- Configurations:
  - L2 next-line
  - L2 best-offset
  - L2 best-offset + L1 next-line
  - L2 best-offset + L1 stride
- Next-line:
  - Insufficient to predict the patterns
  - Lower amount of cache requests
- Best-offset prefetcher:
  - Long training latency
  - Only works for GEMM (dataset size)





# Conclusion

## Key Limitations Identified:

- Limited Flexibility in Memory Hierarchy:
  - Aggressive prefetching exposes bottlenecks in cache and bus structures
  - Blocking bus model between L2 and memory stalls memory requests, causing pipeline bubbles
- Need for Store Queue:
  - Decoupling store operations from LSU can reduce cache pressure and improve efficiency
- Load-Store Unit (LSU) Bottleneck:
  - LSU generates redundant memory requests, increasing congestion
  - Immediate processing of store operations can delay/block loads and prefetches

## Summary:

- Developed a new trace generation (via Spike)
- Developed a new extensible prefetcher module
- Integrated data prefetching in Olympia
- Validated performance gains with next-line, stride, and best-offset prefetchers



DEFINING TECHNOLOGY

**Thank you!**