



University of Stuttgart

Germany

IPVS, Scientific Computing



Parallel FFTW on RISC-V

A Comparative Study
including OpenMP, MPI, and
HPX

June 13, 2025

Alexander Strack,
Christopher Taylor,
and Dirk Pflüger

About Me

- Background in Computational Engineering
- Third-year PhD student in the Scientific Computing department at the University of Stuttgart
- Focused on asynchronous task-based parallelization of numerical algorithms across diverse hardware architectures

Motivation



- The MILK-V Pioneer Box: first practical RISC-V desktop system
- Portability is necessary — but performance is critical

Motivation

Architecture	x86-64	RISC-V
CPU	AMD EPYC 7742	SOPHON SG2042
Cores	64	64
Base clock	2.25GHz	2.00GHz
Process	7nm/14nm	12nm
L3 Cache	256MB	64MB
RAM	2TB DDR4	128GB DDR4
SIMD	AVX2	RVV 0.7.1
NUMA domains	4	4
TDP	225W	120W
Release	2019	2023



Software Stack

The Fastest Fourier Transform in the West (FFTW)

- De-facto standard for over 20 year
- Backend for current state-of-the-art libraries
- Performance secret: black magic

The Fastest Fourier Transform in the West (FFTW)

- De-facto standard for over 20 year
- Backend for current state-of-the-art libraries
- Performance secret: ~~black magic~~
 - Small FFT codelets generated at compile time
 - Combination of codelets for large transforms
 - Planning to find best combination for given system

HPX-FFT

- Leveraging FFTW as backend for 1D FFTs
- Parallelization using the asynchronous many-task runtime HPX
- Origin: tool to evaluate different parallelization approaches with HPX

HPX-FFT


- Leveraging FFTW as backend for 1D FFTs
- Parallelization using the asynchronous many-task runtime HPX
- Origin: tool to evaluate different parallelization approaches with HPX
- Here: performance optimization on x86 and RISC-V, baseline for evaluation of FFTWs parallel plans



Parallel FFT


Two-dimensional FFT

- Strided access




Two-dimensional FFT

- Strided access



- Data transpose



Parallel FFTW

Threading backends:


- OpenMP and pthreads
- Choose the block size and number of threads in order to minimize the critical path

MPI backend:

- All-to-all collectives
- Slab decomposition
- Compatible with threading backends to run MPI+X


HPX-FFT variants

- HPX-FFT naive




HPX-FFT variants

- HPX-FFT opt




HPX-FFT variants

- HPX-FFT sync



HPX-FFT variants

- HPX-FFT for_loop





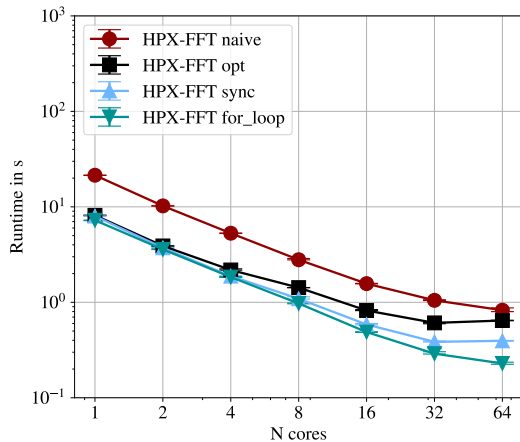
Results

Benchmark

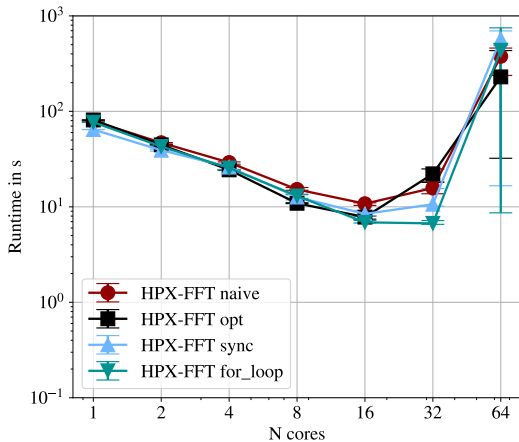
Strong scaling from one to 64 cores:

- Two-dimensional real-to-complex FFT
- No vectorization
- Problem size: $2^{14} \times 2^{14}$
- Data points: median out of 10 runs
- Errorbars: min/max out of 10 runs

HPX-FFT optimization




AMD EPYC 7742




SOPHON SG2042

HPX-FFT partial (for_loop)

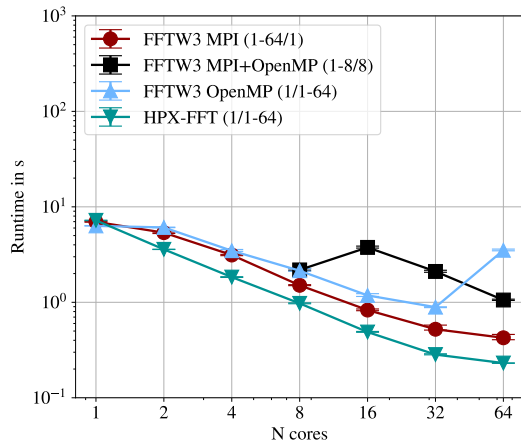


AMD EPYC 7742




SOPHON SG2042

FFTW estimated planning

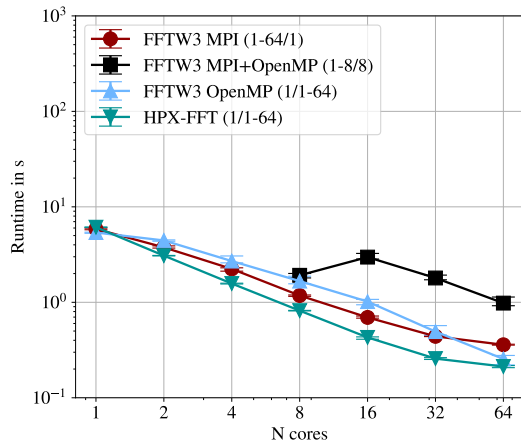


AMD EPYC 7742




SOPHON SG2042

FFTW measured planning



AMD EPYC 7742



SOPHON SG2042



Conclusion and Outlook

Conclusion

- Removing synchronization barriers does not guarantee better performance
- Across different backends performance delta between factor four to eight
- Memory bottleneck on SG2042 chip

Outlook

- Distributed FFTW and HPX-FFT
- FFTW with RVV1.0
- Comparison against ARM architecture



University of Stuttgart
Germany

Thank You!

Alexander Strack, Christopher Taylor, and Dirk Pflüger
IPVS



Mail	alexander.strack@ipvs.uni-stuttgart.de
Phone	+49 711 685 88308
Internet	https://www.ipvs.uni-stuttgart.de