


An Empirical Comparison of the RISC-V and AArch64 Instruction Sets

Daniel Weaver & Simon McIntosh-Smith

University of Bristol

Motivation




A Fair Comparison

- Power, Performance, Area
- Feature sets
 - Armv8-a+nosimd vs rv64g
- Compilers
- Benchmarks

Problem Definition

$$\textit{Program Execution Time} = \textit{Path Length} \times \textit{CPI} \times \textit{Time per Cycle}$$



The Workloads

- STREAM
 - Sustained memory bandwidth
 - 4 simple kernels
- CloverLeaf_Serial
 - High energy physics simulation on 2D cartesian grid
 - Many kernels each working over entire grid
- miniBUDE
 - Approximation of molecular docking simulation
- Lattice Boltzmann
 - Computational fluid dynamics
 - d2q9-bgk serial code
- minisweep
 - Radiation transportation mini app

$$\text{Program Execution Time} = \text{Path Length} \times \text{CPI} \times \text{Time per Cycle}$$

The Compilers

GCC 9.2 GCC 12.2

`-march=armv8-a+nosimd -mtune=cortex-a55 -static`

`-march=rv64g -mtune=sifive-7-series -static`

$$\textit{Program Execution Time} = \textit{Path Length} \times \textit{CPI} \times \textit{Time per Cycle}$$

SimEng

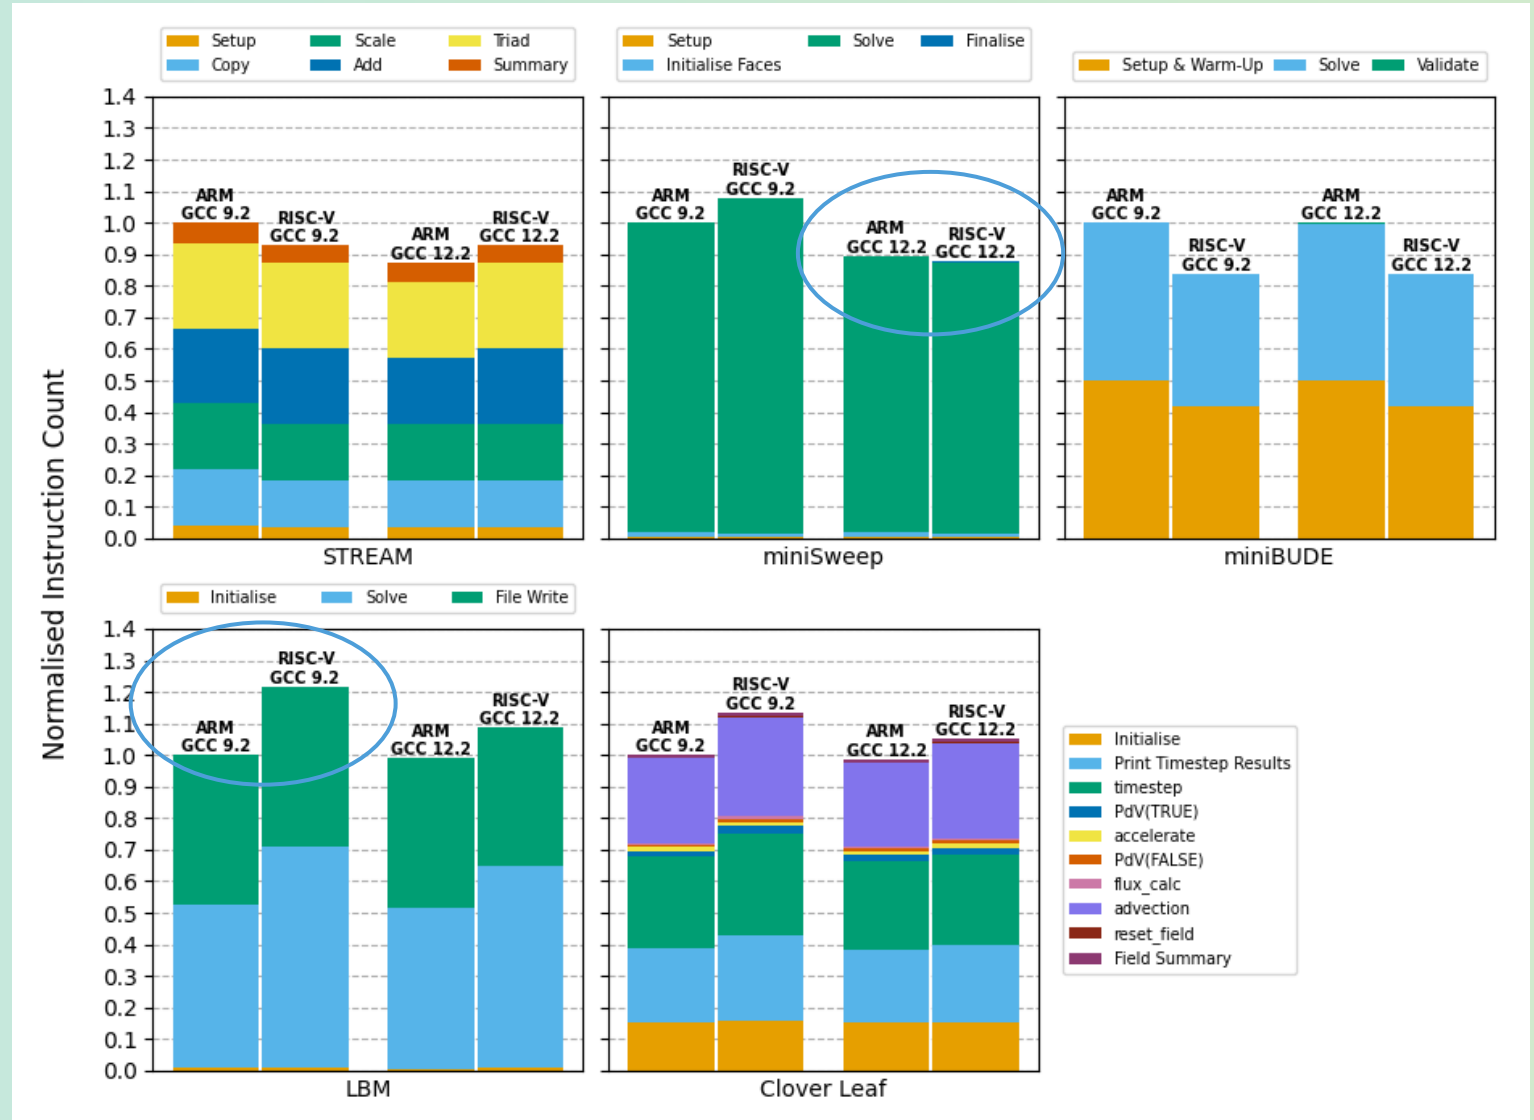


- CPU microarchitecture simulator
- Open source
- Fast
- Easily modifiable
- Accurate
- Supported ARMv8
- RISC-V support added

$$\text{Program Execution Time} = \text{Path Length} \times \text{CPI} \times \text{Time per Cycle}$$

Path Length

- Benchmark source code modified to add instrumentation – NOPs around kernels
- Modified SimEng emulation core



$$\text{Program Execution Time} = \text{Path Length} \times \text{CPI} \times \text{Time per Cycle}$$

STREAM Analysis

```
ldr    d1, [x22, x0, lsl #3]
str    d1, [x19, x0, lsl #3]
add    x0, x0, #1
cmp    x0, x20
b.ne   0x400abc <main+0x248>
```

```
fld    fa5,0(a5)
fsd    fa5,0(a4)
add    a5,a5,8
add    a4,a4,8
bne    a5,s0,10c58 <main+0x25c>
```

$$\text{Program Execution Time} = \text{Path Length} \times \text{CPI} \times \text{Time per Cycle}$$

Ideal CPI

- Ideal processor
 - Execute entire instruction in 1 cycle
 - Execute any number of instructions in 1 cycle
 - Perfect branch prediction
 - MUST obey RAW (true) dependencies between instructions
- How many cycles to complete a program? The longest chain of true dependencies
- Critical path (CP)
- Another modified SimEng emulation core

$$\text{Program Execution Time} = \text{Path Length} \times \text{CPI} \times \text{Time per Cycle}$$

Critical Path Length

- Instruction level parallelism (ILP) measure of no. instructions required to be executed per cycle to achieve ideal CPI

Table 1: Critical Paths and ILP per Benchmark

	STREAM				CloverLeaf			
	GCC 9.2		GCC 12.2		GCC 9.2		GCC 12.2	
	AArch64	RISC-V	AArch64	RISC-V	AArch64	RISC-V	AArch64	RISC-V
Path Length	3,350,107,615	3,110,150,358	2,930,114,073	3,110,139,144	12,832,452	14,553,390	12,647,061	13,481,498
CP	10,000,234	10,005,341	10,000,234	10,004,815	46,933	191,538	46,658	228,036
ILP	335	311	293	311	273	76	271	59
2GHz Run time (ms)	5.00	5.00	5.00	5.00	0.0235	0.0958	0.0233	0.114

	LBM				miniBUDE			
	GCC 9.2		GCC 12.2		GCC 9.2		GCC 12.2	
	AArch64	RISC-V	AArch64	RISC-V	AArch64	RISC-V	AArch64	RISC-V
Path Length	380,391,346	463,305,683	376,329,390	412,979,829	137,280,541	115,064,988	137,183,536	114,897,049
CP	10,910,427	5,196,321	4,660,144	4,873,467	196,357	197,285	196,331	196,722
ILP	35	89	81	85	699	583	699	584
2GHz Run time (ms)	5.46	2.60	2.33	2.44	0.0982	0.0986	0.0982	0.0984

	minisweep			
	GCC 9.2		GCC 12.2	
	AArch64	RISC-V	AArch64	RISC-V
Path Length	2,162,866,809	2,332,356,452	1,934,709,957	1,894,737,614
CP	263,120	263,327	280,567	272,444
ILP	8,220	8,857	6,896	6,955
2GHz Run time (ms)	0.132	0.132	0.140	0.136

$$\text{Program Execution Time} = \text{Path Length} \times \text{CPI} \times \text{Time per Cycle}$$

Table 2: Scaled Critical Paths and ILP per Benchmark

	STREAM				CloverLeaf			
	GCC 9.2		GCC 12.2		GCC 9.2		GCC 12.2	
	AArch64	RISC-V	AArch64	RISC-V	AArch64	RISC-V	AArch64	RISC-V
Scaled CP	60,000,545	60,005,845	60,000,545	60,005,845	94,983	191,538	81,925	244,103
ILP	56	52	49	52	135	76	154	55
2GHz Run time (ms)	30.0	30.0	30.0	30.0	0.0475	0.0958	0.0410	0.122
	LBM				miniBUDE			
	GCC 9.2		GCC 12.2		GCC 9.2		GCC 12.2	
	AArch64	RISC-V	AArch64	RISC-V	AArch64	RISC-V	AArch64	RISC-V
CP	42,344,992	5,888,686	4,660,233	5,565,925	685,839	685,842	685,680	685,291
ILP	9.0	79	81	74	168	168	168	168
2GHz Run time (ms)	21.2	2.94	2.33	2.78	0.343	0.343	0.343	0.343
	minisweep							
	GCC 9.2		GCC 12.2					
	AArch64	RISC-V	AArch64	RISC-V				
CP	1,577,198	1,586,189	1,592,550	1,577,099				
ILP	1,371	1,470	1,215	1,201				
2GHz Run time (ms)	0.790	0.793	0.796	0.789				

Scaled CPI

- Scale CP by instruction latencies
- Latencies roughly follow that of TX2
- Do not scale for loads and stores

$$\text{Program Execution Time} = \text{Path Length} \times \text{CPI} \times \text{Time per Cycle}$$

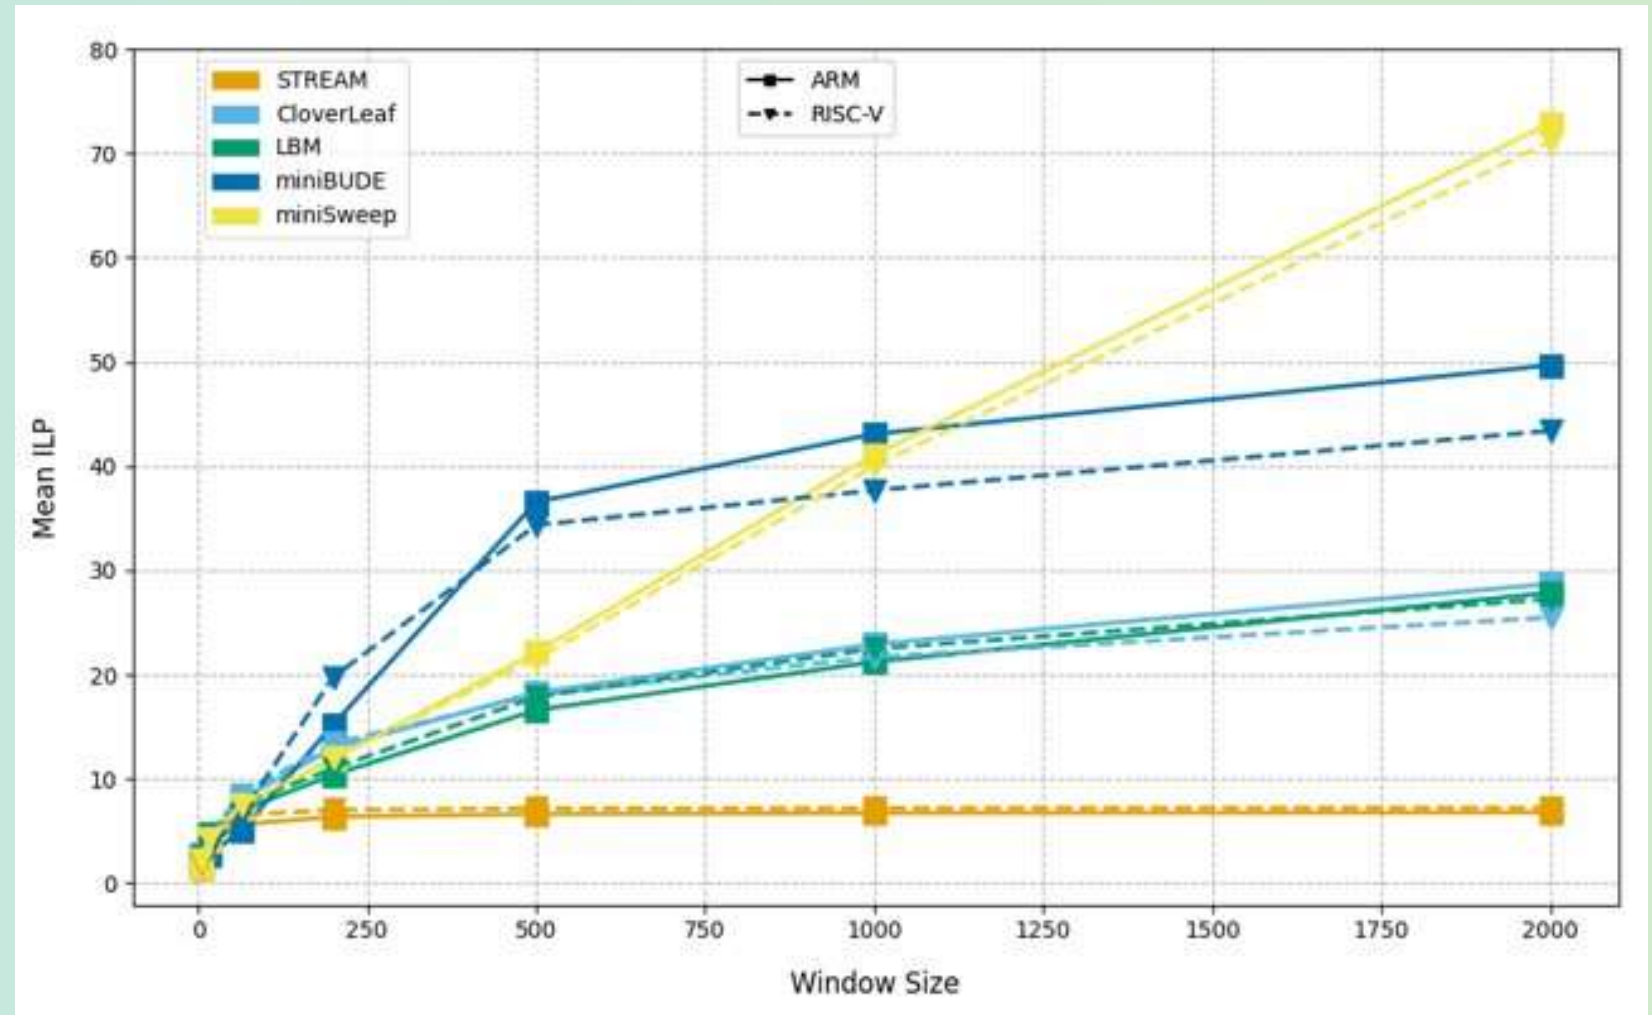
Windowed CP

- Slide window down dynamic instruction trace
- Determine CP length per window
- Window size and stride parameterisable
- We choose values ranging from 4 – 2000
- Stride 50% of size
- GCC 12.2 only

```
bl    4003f4 <mysecond>
fmov  d9, d0
mov   x0, #0x0           // #0
ldr   d1, [x22, x0, lsl #3]
str   d1, [x19, x0, lsl #3]
add   x0, x0, #0x1
cmp   x0, x20
b.ne  400abc <main+0x248> // b.any
ldr   d1, [x22, x0, lsl #3]
str   d1, [x19, x0, lsl #3]
add   x0, x0, #0x1
cmp   x0, x20
b.ne  400abc <main+0x248> // b.any
ldr   d1, [x22, x0, lsl #3]
str   d1, [x19, x0, lsl #3]
add   x0, x0, #0x1
cmp   x0, x20
b.ne  400abc <main+0x248> // b.any
ldr   d1, [x22, x0, lsl #3]
str   d1, [x19, x0, lsl #3]
add   x0, x0, #0x1
cmp   x0, x20
b.ne  400abc <main+0x248> // b.any
ldr   d1, [x22, x0, lsl #3]
str   d1, [x19, x0, lsl #3]
```

$$\text{Program Execution Time} = \text{Path Length} \times \text{CPI} \times \text{Time per Cycle}$$

Windowed CP Results



$$\text{Program Execution Time} = \text{Path Length} \times \text{CPI} \times \text{Time per Cycle}$$

Conclusion and Future Work

- Full OoO superscalar microarchitecture simulation
- Varying available resources
- More benchmarks
- More compilers
- More ISA extensions

$$\textit{Program Execution Time} = \textit{Path Length} \times \textit{CPI} \times \textit{Time per Cycle}$$

Important Links

- Daniel Weaver – ra18837@bristol.ac.uk
- Paper - <https://dl.acm.org/doi/10.1145/3624062.3624233>
- Artifact - <https://github.com/UoB-HPC/Arm-RISCV-Empirical-Comparison-Artifact>
- SimEng repository - <https://github.com/UoB-HPC/SimEng>
- SimEng Documentation - <https://uob-hpc.github.io/SimEng/>