



Why Esperanto picked RISC-V for HPC Computing

Dave Ditzel
Founder and CTO
Esperanto Technologies

November 18, 2024
RISC-V Workshop Keynote
Supercomputing24

email: dave@esperanto.ai

History of instruction sets leading us to RISC-V

Why RISC-V was the right choice for Esperanto

Esperanto ET-SoC-1 at the beginning of RISC-V

Esperanto 4nm ET-SoC-2 and 2nm ET-SoC-3 roadmap with RVV Vector support

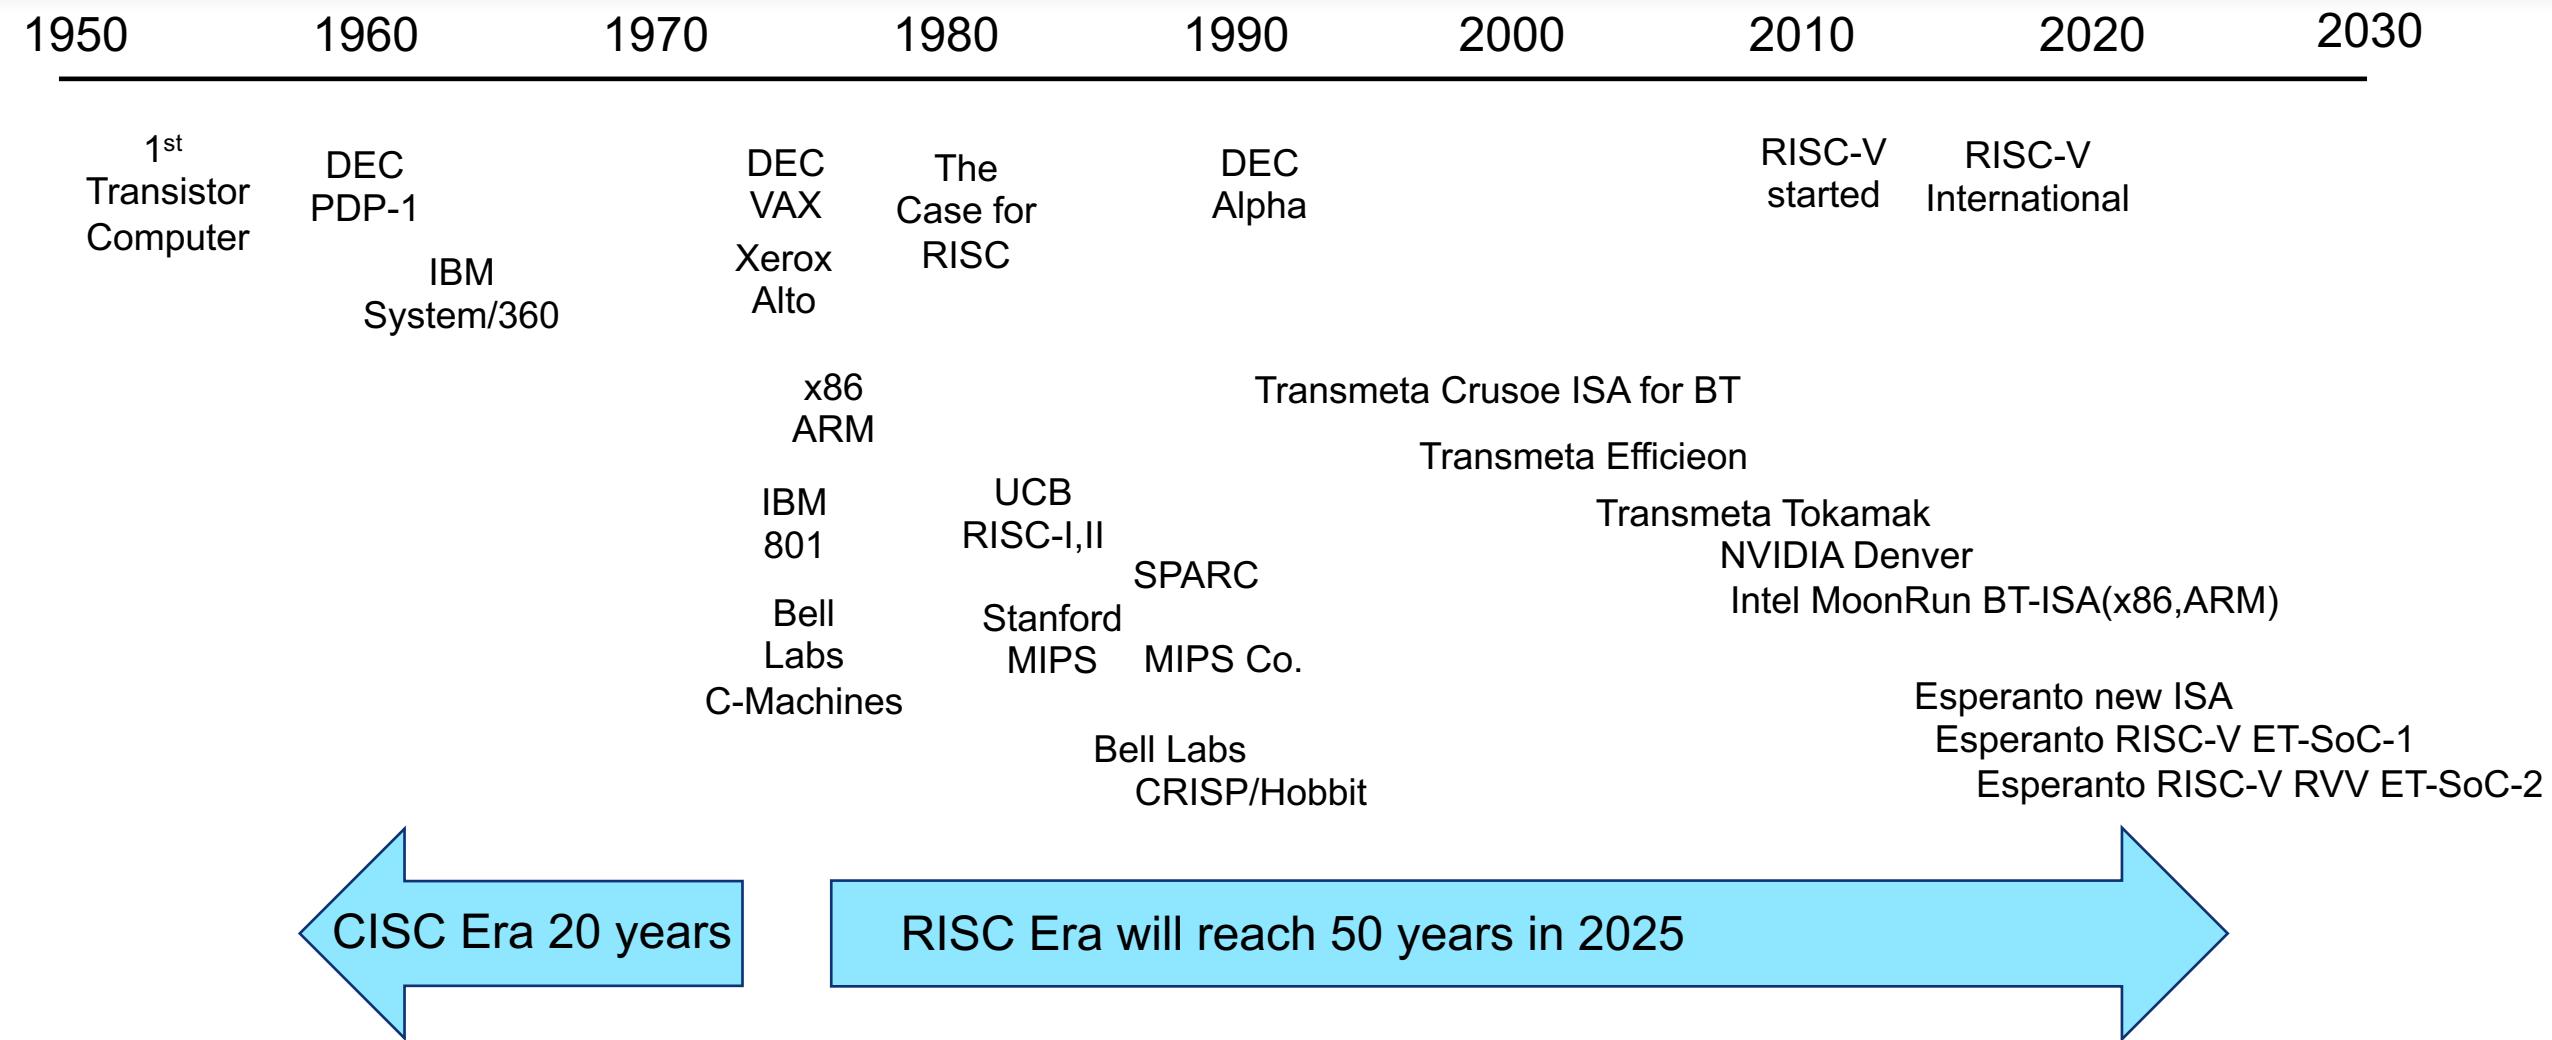
Conclusion: RISC-V is the right instruction set if you want to run both AI and HPC.

Useful quote for computer architects

“Those who cannot remember the past
are condemned to repeat it.”

**George Santayana,
The Life of Reason,
1905.**

My Computer Development Timeline



Era of High Level Language Computer Architecture

later called CISC (Complex Instruction-set Computers)

DEC VAX-11/780: a 32-bit CISC

A real 70's architecture:

- VAX: Transitioned from 16-bit PDP-11 assembly to 32-bit addresses and compiled code
- ISA made it easy to take high level language statements and cast into VAX assembly code

Instruction set heavily influenced by availability of 8-bit wide integrated circuits

VAX ISA composed of variable length instructions

- Instruction opcode (1 or 2 bytes) followed by up to 6 operands
- First operand descriptor (1 byte)
- First operand data (0-4 bytes)
- Second operation descriptor
- Second operand data
- Third operand descriptor
- Third operand data (0-4 bytes)
- ...

This was great when decoding serially one byte at a time,
but a nightmare for pipelined or later superscalar implementations



Xerox Palo Alto Research Center: Alto



Earliest Graphical User Interface with mouse

5.8 MHz microcoded CPU

User loadable instruction sets led to several HLLCA type bytecoded instruction sets for languages like BCPL, Smalltalk and MESA

Enabled exploration of highly tailored instruction sets

But raised the question

“Why not compile to lowest microcode level?”

SYMBOL:

The Ultimate CISC

SYMBOL: The ultimate Complex Instruction Set Computer

SYMBOL was a High Level Language Computer announced in 1971 by Fairchild Semiconductor

Supported by Gordon Moore and Robert Noyce at Fairchild until they left to form Intel

Goal was to reduce cost of software by using hardware instead, literally “programmers cost too much”.

Implemented Compiler, Text Editor and Operating System entirely in logic gates

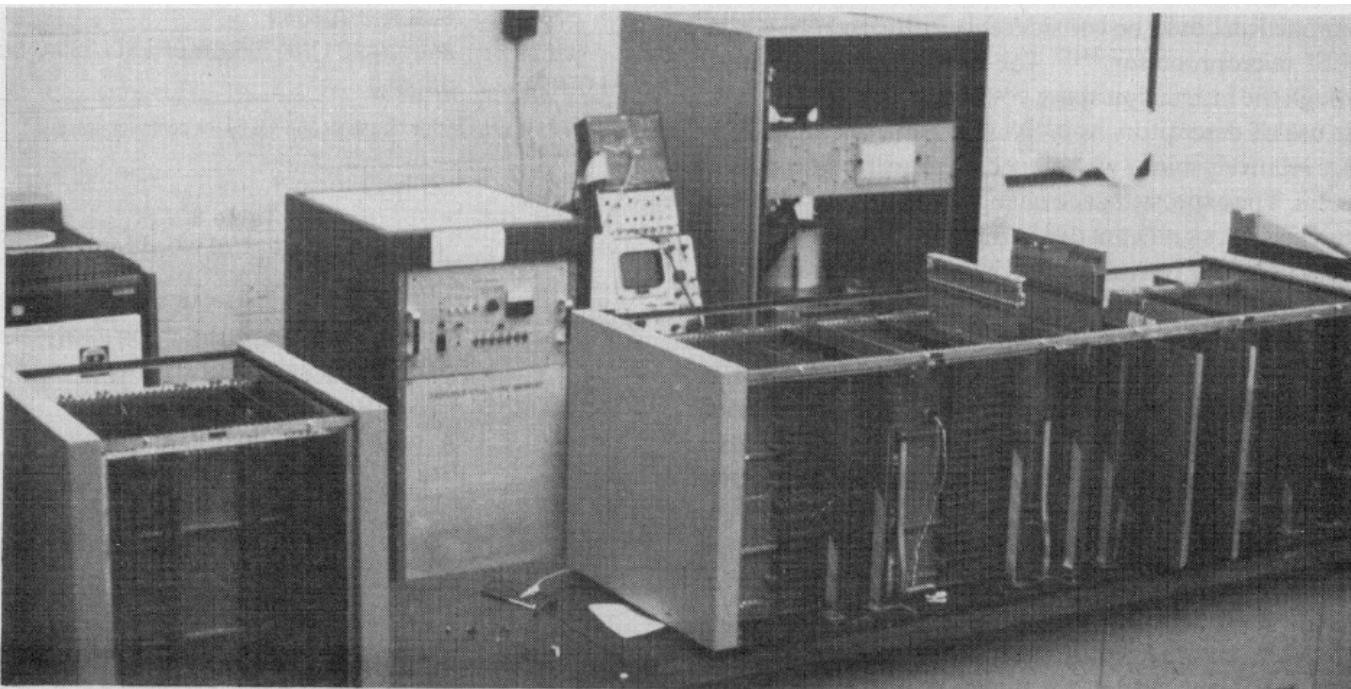
- 2 gates or 1 Flip-Flop per 14-pin DIP package, no ROM
- 20 thousand chips – took several years to debug at ISU
- Instruction set was bytecode mapped almost 1-1 from high level ALGOL/LISP like typeless language
- Tagged architecture with descriptors, and “logical memory” that was not linearly addressable.
- Five years of my life thinking about CISC vs RISC, i.e. better ways to use 20K chips, programming this machine.

I got to work on this computer while a student for 5 years, after it was donated to ISU

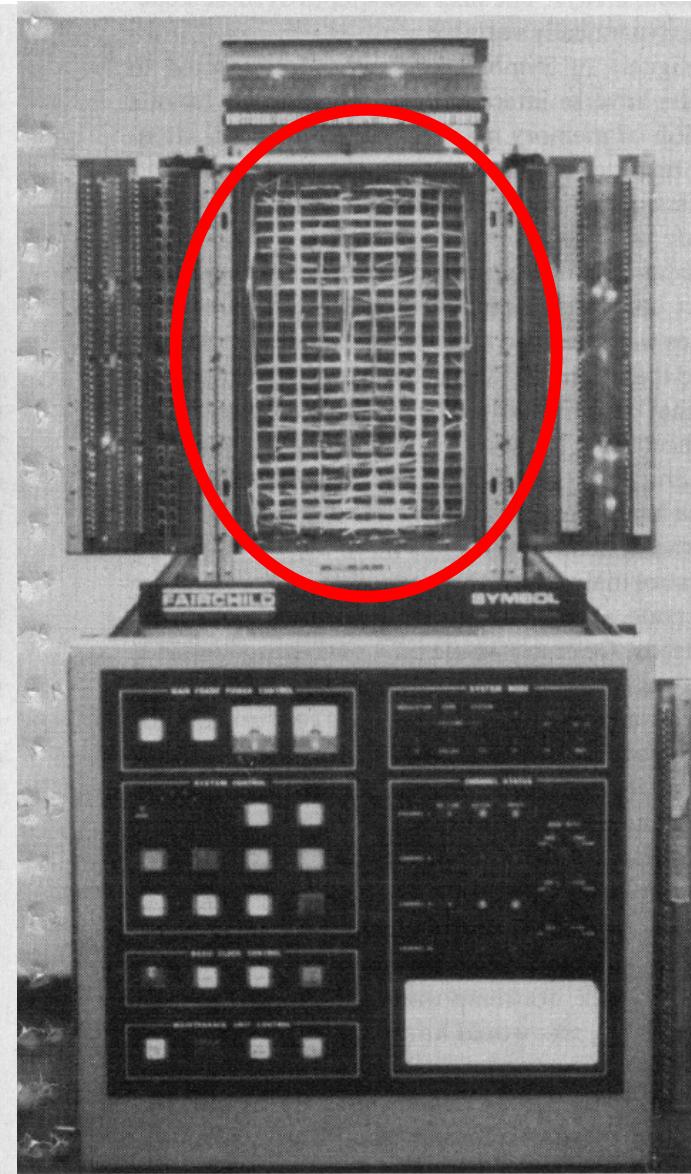
Lessons:

- Don't build your compiler or OS in logic gates
- Use right combination of Hardware, Software or Microcode

Debugging SYMBOL, the inspiration for RISC



A Symbol "terminal" could consist of up to 99 physical I/O devices. The terminal shown at left used a modified IBM Selectric typewriter, editing keyboard, status display, card reader, and line printer. The book next to the typewriter contains operating system and utility source listings. The photo above, a side view of the Symbol mainframe, shows the maintenance processor (far left), power supplies for + 4.5 volts at 1000 amps (below the mainframe), and disk memory, core memory, and paging drum (background). The front view of the mainframe (photo at right) shows a printed circuit card on an extender for testing. Each card held 200 ICs. Wing panels indicate the value of bus signals—100 on the left, 100 on the right, and 50 on top. Additional wire on the PCB was used for "bug" fixes. Each processor could be monitored from a "processor active" lamp on the system's control panel.



Inside red circle
are 220 chips.

In order to fix a bug,
white wires added
on top of blue pc
board to make logic
changes.

Many of the 100
boards were covered
with white bug fixes.

Source: D. Ditzel, Reflections on the High Level Language SYMBOL Computer Systems, IEEE Computer, July 1981.

1978 Photo: My career started at the same time as the first microprocessors



My personal hobby computer

8-bit 6502 Microprocessor

4 MHz

4K Bytes of main memory

Hand “wire-wrapped”

Paper tape reader

Hand wound transformer on
power supply.

Note – base of SYMBOL computer
in background

IBM 801: In my view, the first true RISC

IBM Introduces the 801 Minicomputer, the First Computer Employing RISC

1974

[Permalink](#)



Image Source: www.ibm.com

John Cocke with the computer incorporating RISC architecture that he invented.

In 1974 IBM built the first prototype computer employing [RISC \(Reduced Instruction Set Computer\)](#) architecture. Based on an invention by IBM researcher [John Cocke](#), the RISC concept simplified the instructions given to run computers, making them faster and more powerful. It was implemented in the experimental [IBM 801](#) minicomputer. The goal of the 801 was to execute one instruction per cycle.

In 1987 John Cocke received the A. M. Turing Award for significant contributions in the design and theory of compilers, the architecture of large systems and the development of reduced instruction set computers (RISC); for discovering and systematizing many fundamental transformations now used in optimizing compilers including reduction of operator strength, elimination of common subexpressions, register allocation, constant propagation, and dead code elimination.

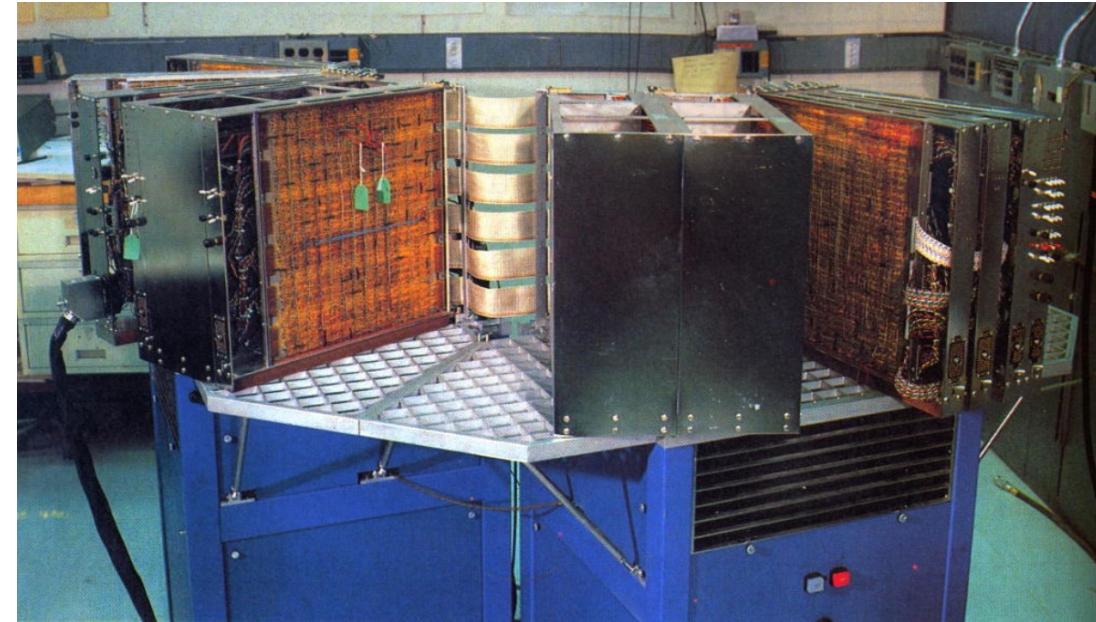


Photo of IBM 801 minicomputer.

The name 801 was from the IBM building number of the T.J. Watson Research Center in Yorktown Heights

First described at ASPLOS-1 in 1982

1980: The Case for RISC Published

The Case for the Reduced Instruction Set Computer

David A. Patterson

Computer Science Division
University of California
Berkeley, California 94720

David R. Ditzel

Bell Laboratories
Computing Science Research Center
Murray Hill, New Jersey 07974

INTRODUCTION

One of the primary goals of computer architects is to design computers that are more cost-effective than their predecessors. Cost-effectiveness includes the cost of hardware to manufacture the machine, the cost of programming, and costs incurred related to the architecture in debugging both the initial hardware and subsequent programs. If we review the history of computer families we find that the most common architectural change is the trend toward ever more complex machines. Presumably this additional complexity has a positive tradeoff with regard to the cost-effectiveness of newer models. In this paper we propose that this trend is not always cost-effective, and in fact, may even do more harm than good. We shall examine the case for a Reduced Instruction Set Computer (RISC) being as cost-effective as a Complex Instruction Set Computer (CISC). This paper will argue that the next generation of VLSI computers may be more effectively implemented as RISC's than CISC's.

WORK ON RISC ARCHITECTURES

At Berkeley. Investigation of a RISC architecture has gone on for several months now under the supervision of D.A. Patterson and C.H. Séquin. By a judicious choice of the proper instruction set and the design of a corresponding architecture, we feel that it should be possible to have a very simple instruction set that can be very fast. This may lead to a substantial net gain in overall program execution speed. This is the concept of the Reduced Instruction Set Computer. The implementations of RISC's will almost certainly be less costly than the implementations of CISC's. If we can show that simple architectures are just as effective to the high-level language programmer as CISC's such as VAX or the IBM S/38, we can claim to have made an effective design.

At Bell Labs. A project to design computers based upon measurements of the C programming language has been under investigation by a small number of individuals at Bell Laboratories Computing Science Research Center for a number of years. A prototype 16-bit machine was designed and constructed by A.G. Fraser. 32-bit architectures have been investigated by S.R. Bourne, D.R. Ditzel, and S.C. Johnson. Johnson used an iterative technique of proposing a machine, writing a compiler, measuring the results to propose a better machine, and then repeating the cycle over a dozen times. Though the initial intent was not specifically to come up with a simple design, the result was a RISC-like 32-bit architecture whose code density was as compact as the PDP-11 and VAX [Johnson79].

At IBM. Undoubtedly the best example RISC is the 801 minicomputer, developed by IBM Research in Yorktown Heights, N.Y.[Electronics76] [Datamation79]. This project is several years old and has had a large design team exploring the use of a RISC architecture in combination with very advanced compiler technology. Though many details are lacking their early results seem quite extraordinary. They are able to benchmark programs in a subset of PL/I that runs about five times the performance of an IBM S/370 model 168. We are certainly looking forward to more detailed information.

Early RISC Principles

RISC: Reduced Instruction Set Computer

- RISC and CISC terms coined by UC Berkeley professors Dave Patterson and Carlo Sequin

Simple orthogonal instructions

- Often fixed 32-bit length
- Easy for compiler code generation
- Easy to implement with small (10-20) gates per cycle
- Easily pipelined

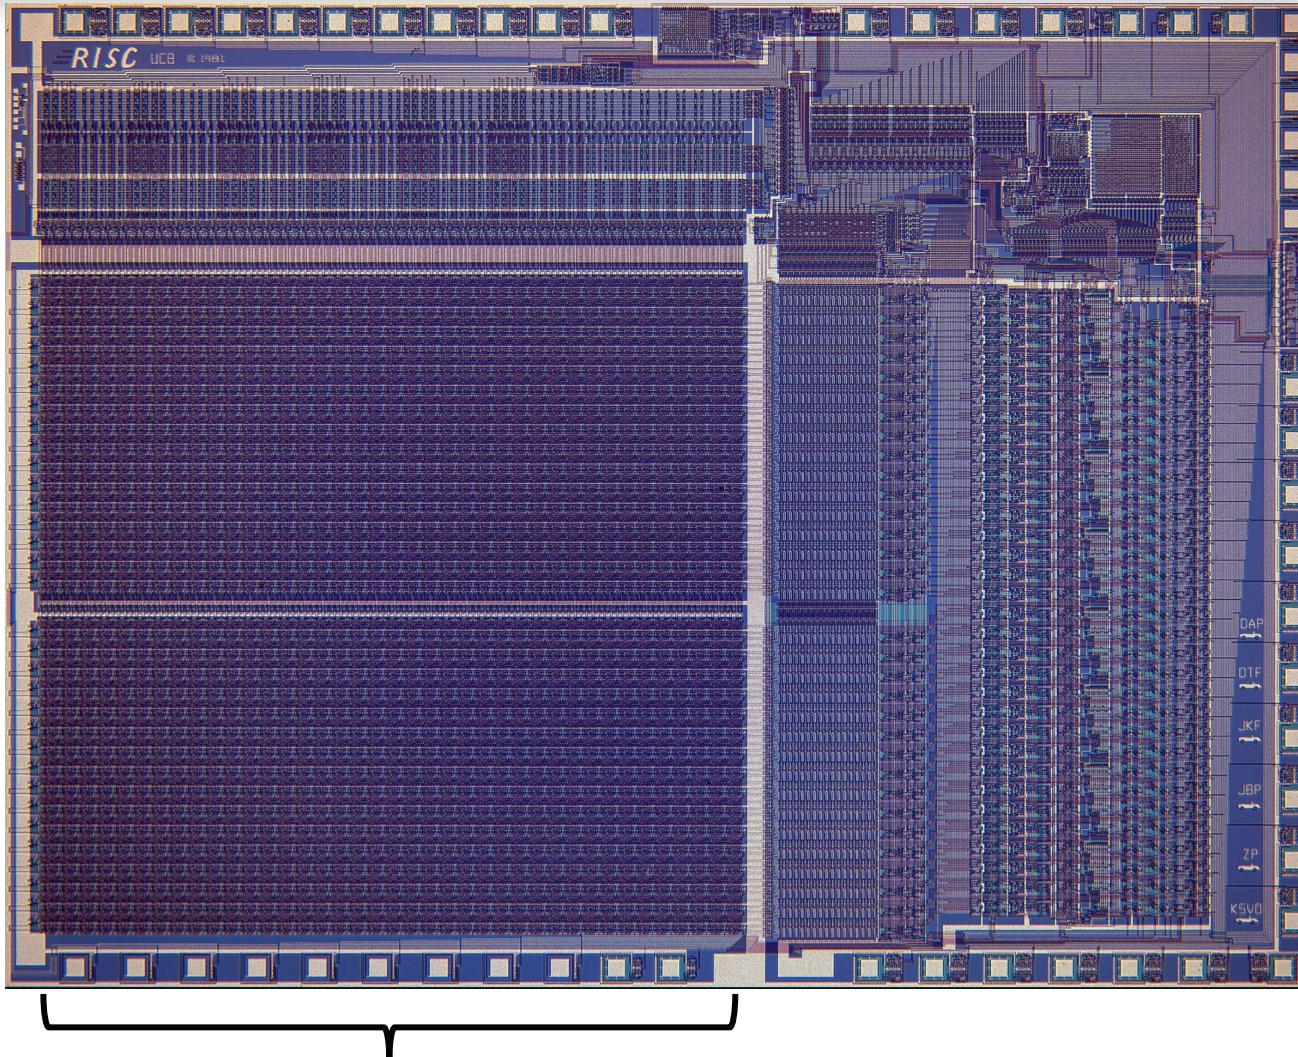
General purpose register file

- Enough to keep data in registers for current and another called subroutine
- Variation in register file style: Flat, Windowed, or invisible (stack cache)

Single load or store per instruction

Register to register arithmetic operations

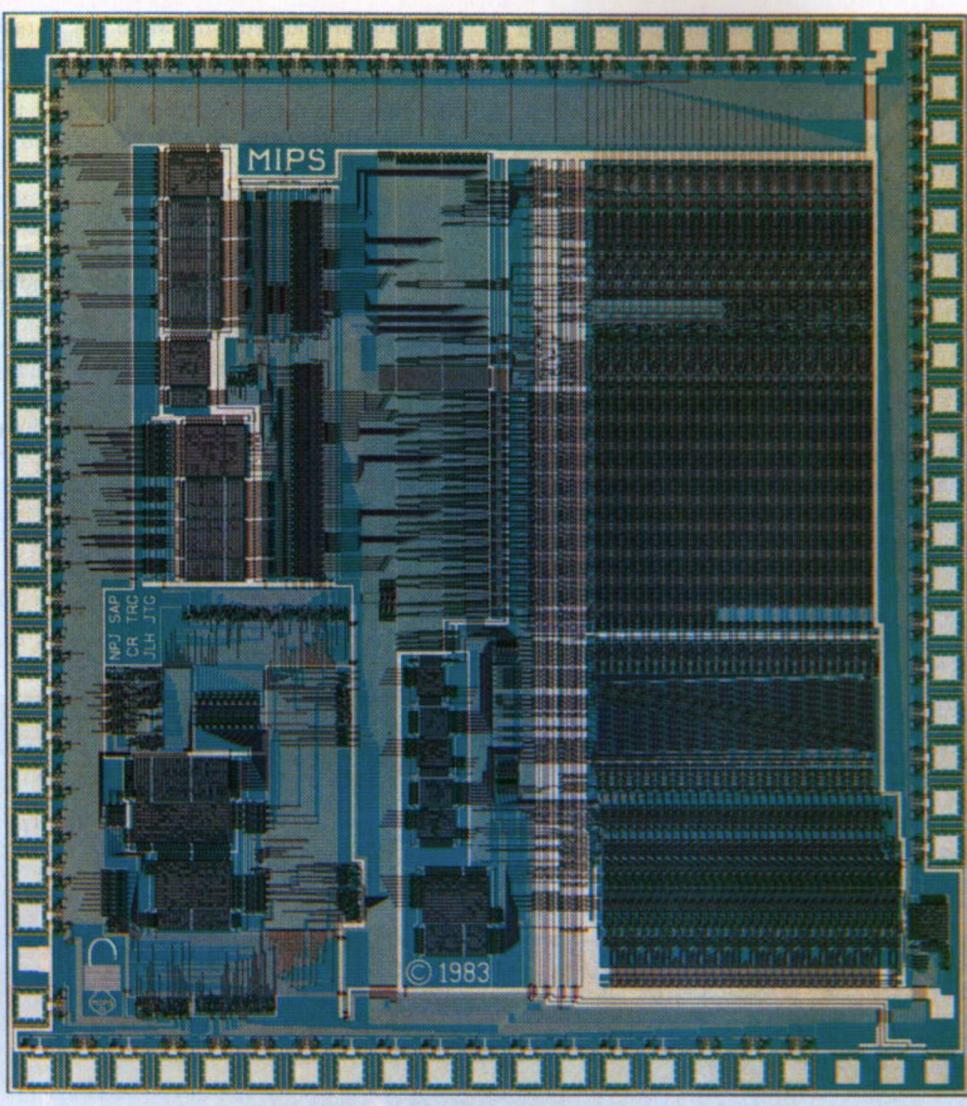
UC Berkeley RISC-I ~1981



Register Windows

- Most of die consumed with register windows
- More registers to keep more operands on die
- Single load or store to/from a register
- 3-address register to register operations
- Integer only
- No caches
- But simple enough for students to design!

Stanford MIPS ~1983



The MIPS instruction set consists of about 111 total instructions

Each instruction encoded in 32-bits

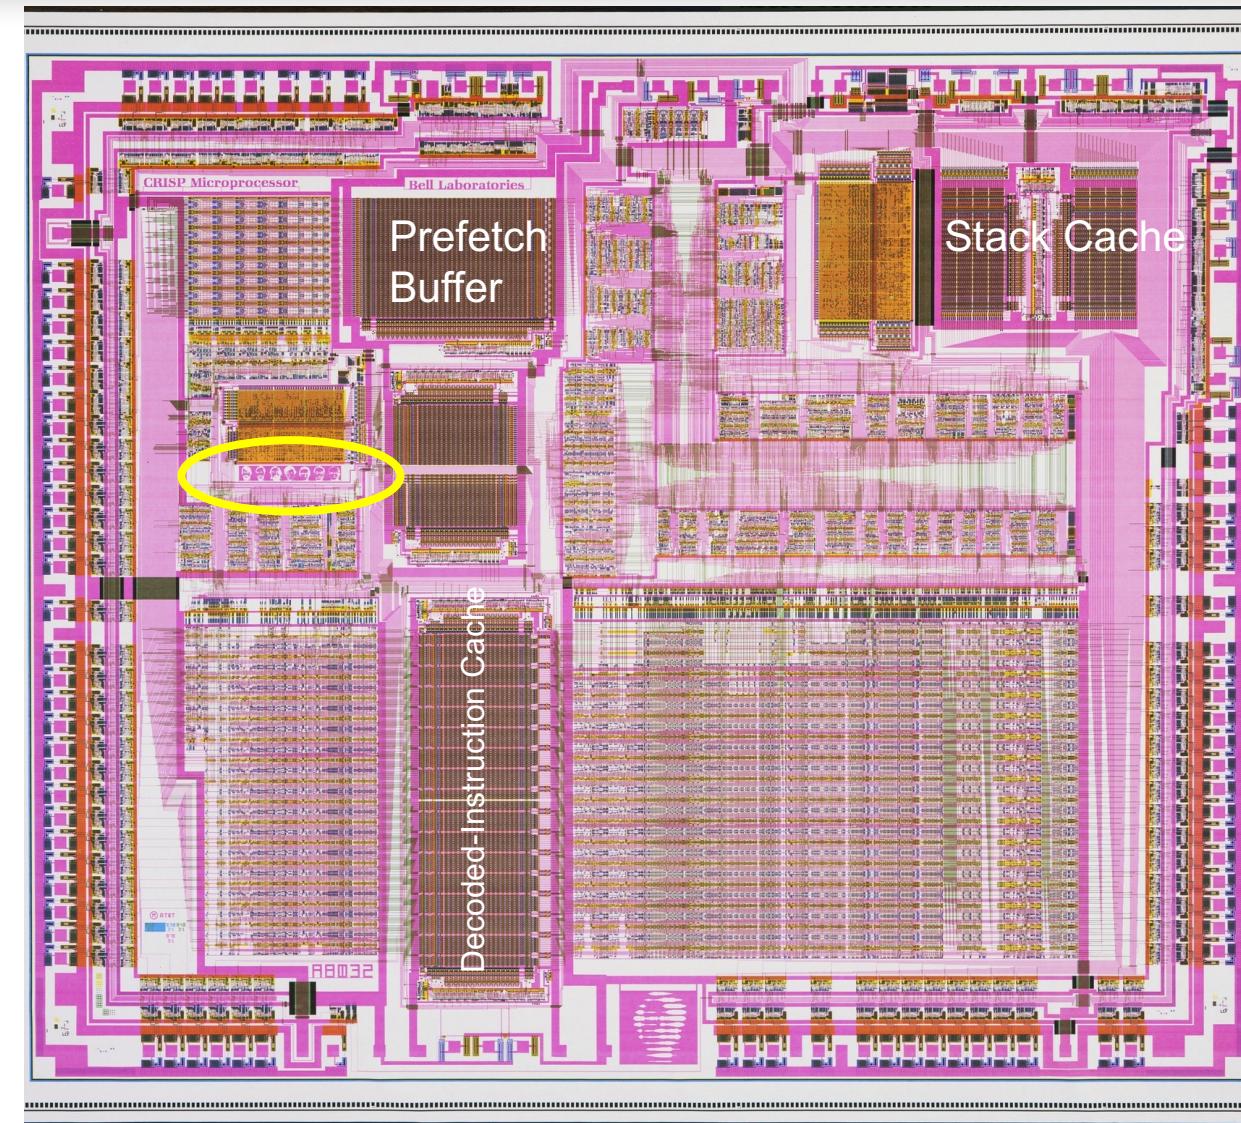
Pipelined to execute one instruction per clock

The instruction set includes:

- 21 arithmetic instructions (+, -, *, /, %)
- 8 logic instructions (&, |, ~)
- 8 bit manipulation instructions
- 12 comparison instructions (>, <, =, >=, <=, ≠)
- 25 branch/jump instructions
- 15 load instructions
- 10 store instructions
- 8 move instructions
- 4 miscellaneous instructions

Delayed branch with 2 delay slots

AT&T Bell Labs CRISP: C language Reduced Instruction Set Processor



C-Language Reduced Instruction Set Processor

1.75 micron CMOS

Announced 1987

1st CMOS superscalar (more than 1 instr/clock) chip

Hardware translated compact instructions into 180-bit wide decoded micro-Op cache

Branch folding eliminated the overhead of branches

“Stack Cache” performed automatic register allocation in hardware

Low power chip, used in EO Tablet computer

Software binary translator used to move software from CISC WE32100

Lessons

- External to internal translation worked well
- Decouple external and internal ISA



So why did Esperanto pick RISC-V and why you should too

Esperanto initially had its own "universal" instruction set, but we switched to RISC-V as it gained momentum

- We needed freedom to use and innovate with an ISA

RISC-V scalar ISA not that different from early RISC designs

- UC Berkeley team did a good job of taking the best from prior RISC chips to create RISC-V
- Notably missing: delayed branches, fixed length instructions (mix of 16/32-bit), no register windows
- Now the committees are adding features so complexity will grow... but still cleaner than alternatives.

The real IMPORTANT reasons to use RISC-V

- RISC-V is the only free and open instruction set that currently has critical mass of adoption
 - Intel and ARM work to keep non-licensed outsiders from creating new CPU implementations
 - RISC-V is where the innovation will occur
- RISC-V allows and encourages innovation with both private and official extensions
- RISC-V has a cleaner instruction set that enables simpler more energy efficient implementations
- RISC-V is good enough that inventing yet another RISC instruction set doesn't make sense
- For anyone wanting to do a new CPU, RISC-V is not merely the best choice, its the only practical choice.



Esperanto's CPU Roadmap

7nm

4nm

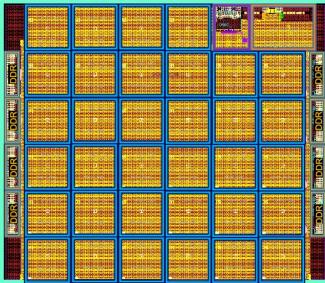
2nm



**November 2024
Subject to change**

Silicon Roadmap from SoC to Chiplets

ET-SoC-1 7nm



Product shipping today

7nm Single chip SoC 15-60W
1K ET-Minion1 RISC-V cores
100 TOp Int8
25 TFlops FP16
12.5 TFlops FP32

PCIe Gen4
Up to 32GB LPDDR4x DRAM

ET-SoC-2x 4nm



Production 2026

4nm Compute Fabric Chiplet 15-60W
256 ET-Minion2 CPU/Chiplet @2 GHz
256 TFlops FP8/Int8 per chiplet
128 TFlops FP16 per chiplet
32 TFlops FP32 per chiplet
16 TFlops FP64 per chiplet
256 to 2048 RISC-V cores per package
Up to 8 ET-SoC-2x chiplets per package

ET-SoC-3x 2nm



Production 2027

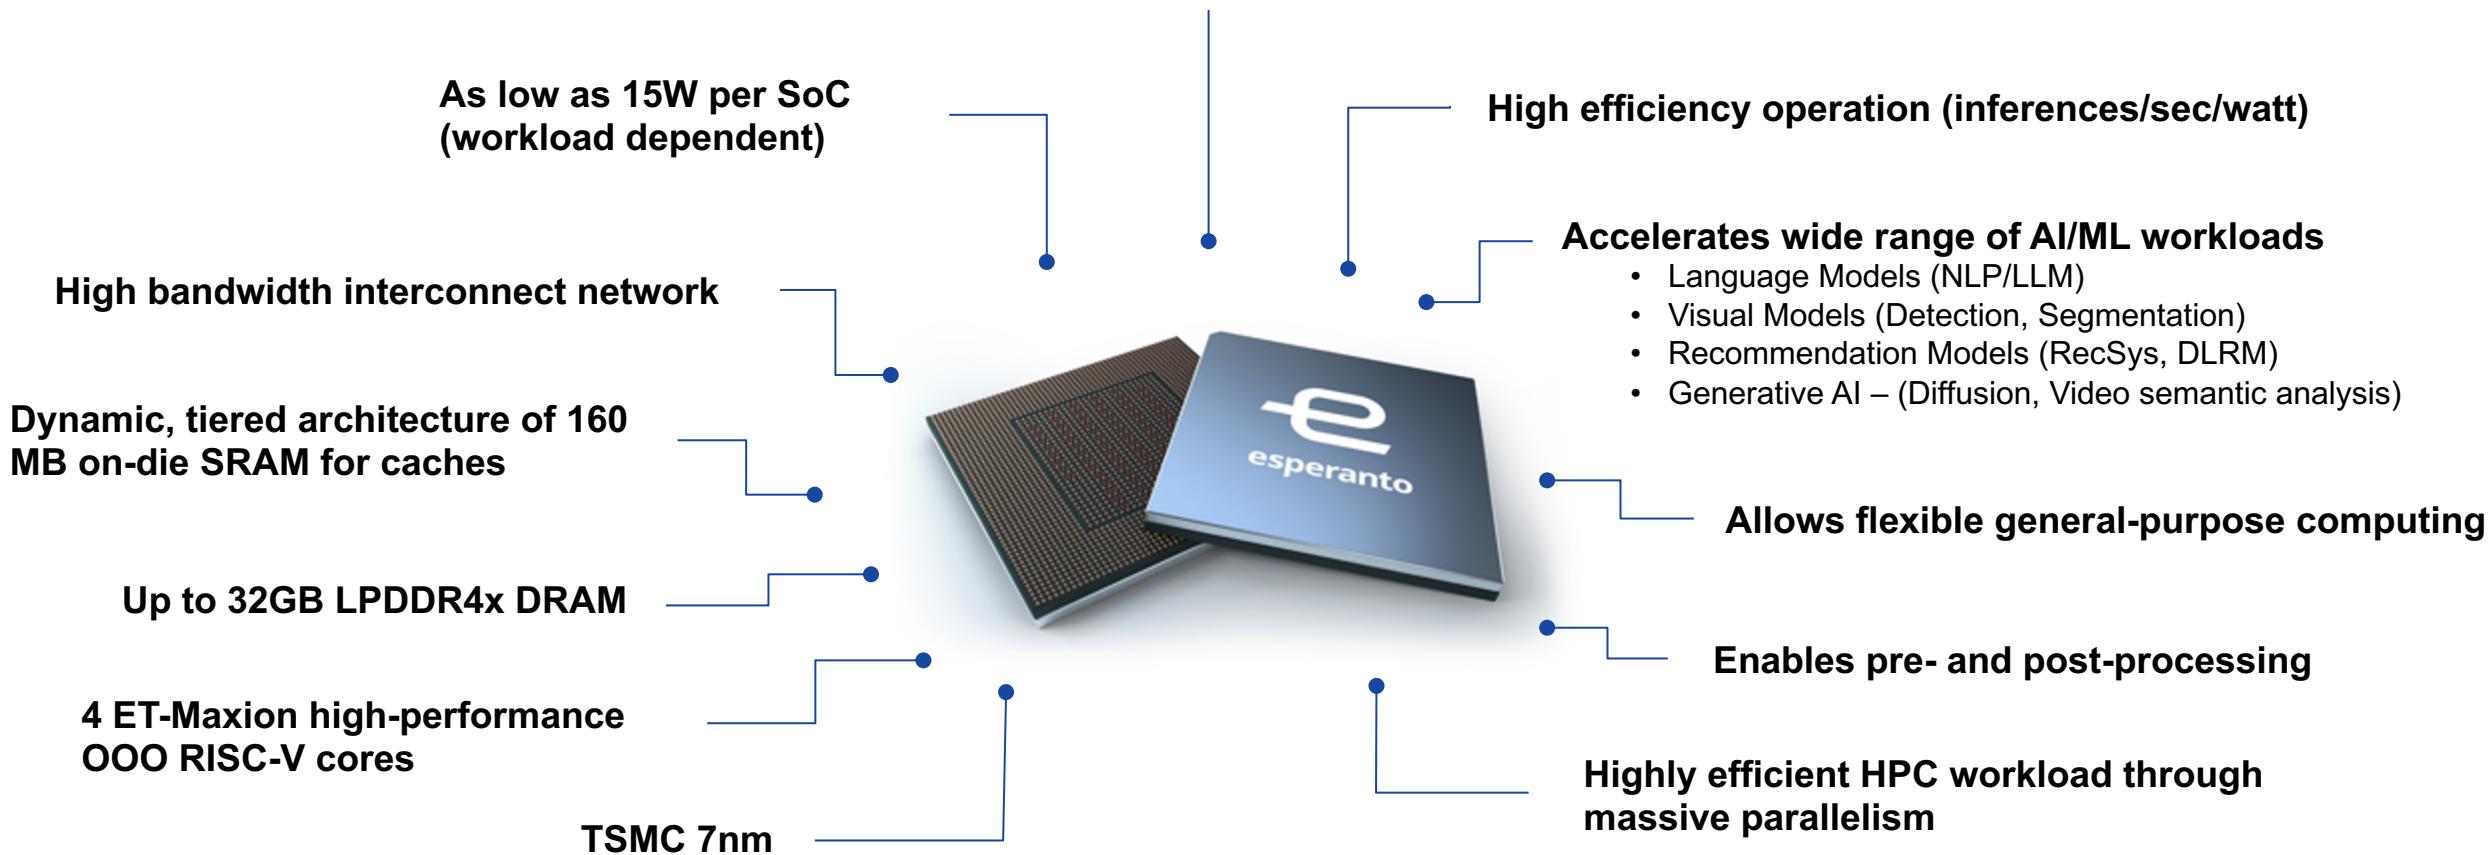
2nm Compute Fabric Chiplet 10-60W
512 ET-Minion3 CPU/Chiplet @2GHz
512 TFlops FP8/Int8 per chiplet
256 TFlops FP16 per chiplet
64 TFlops FP32 per chiplet
32 TFlops FP64 per chiplet
512 to 4096 RISC-V cores per package
Up to 8 ET-SoC-3x chiplets per package

ET-SoC-1

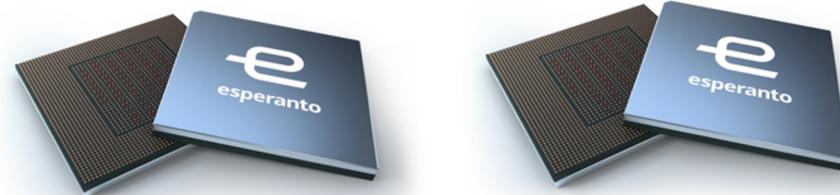
7nm

**Intended for Machine Learning Inference
Prototyping low power advances**

Over 1,000 64-bit RISC-V CPUs per Chip



ET-SoC-1: 7nm First Generation Customer Roadmap



Key Parameters	ET-SoC-1: 600 MHz	ET-SoC-1: 800 MHz
Typical application	<ul style="list-style-type: none">AI inferenceHPC (Single precision)	<ul style="list-style-type: none">AI inferenceHPC (Single precision)
Int8 Peak TOps	77	102
FP16 Peak TFlops	19	26
FP32 Peak TFlops	10	13
Max DRAM Capacity	32 GB	32 GB
Memory Bandwidth	120 GB/s	137 GB/s
Typical Power running ML	< 30 Watts	< 40 Watts
Sample date	Shipping Now	now
FCS	Shipping Now	2025
FP16 Relative Performance including software improvements	1x	1.4x

ET-Minion1 is an Energy-Efficient RISC-V CPU with a Vector/Tensor Unit

CPU is tailored for Massively Parallel ML Applications

ET-MINION1 IS A CUSTOM BUILT 64-BIT RISC-V PROCESSOR

- In-order pipeline with low gates/stage to improve MHz at low voltages
- Architecture and circuits optimized to enable low-voltage operation
- Two hardware threads of execution
- Software configurable L1 data-cache and/or scratchpad

ML OPTIMIZED VECTOR/TENSOR UNIT

512-bit wide integer per cycle

- 128 8-bit integer operations per cycle, accumulates to 32-bit Int

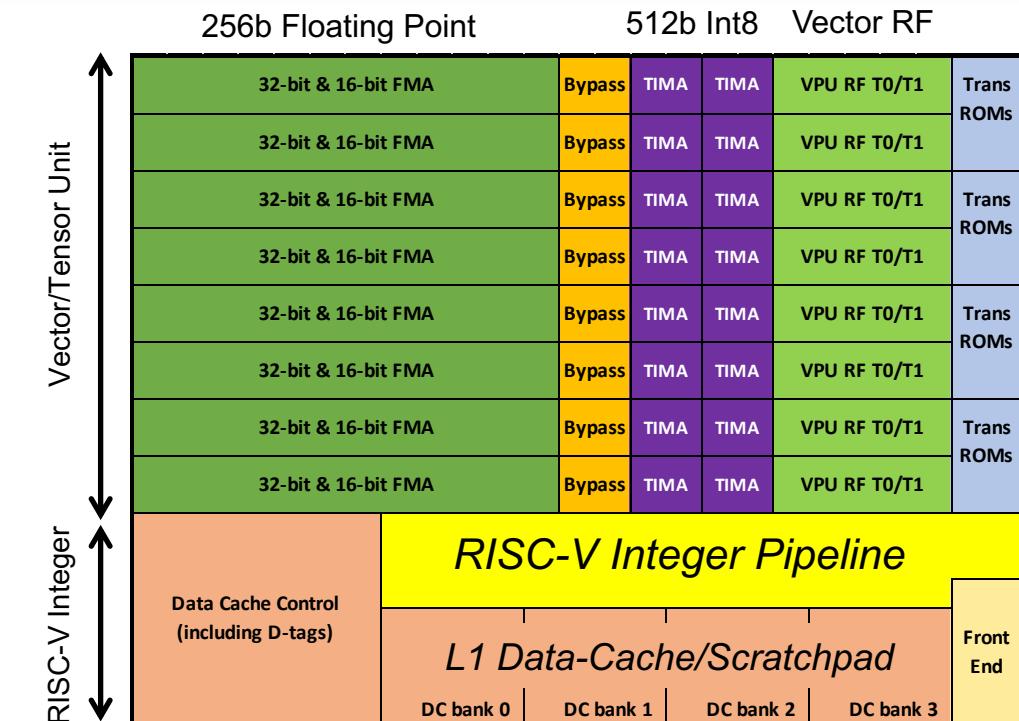
256-bit wide floating point per cycle

- 16 32-bit single precision operations per cycle
- 32 16-bit half precision operations per cycle

New multi-cycle Tensor Instructions

- Can run for up to 512 cycles (up to 32K operations) with one tensor instruction
- Reduces instruction fetch bandwidth and reduces power
- RISC-V integer pipeline put to sleep during tensor instructions

Vector transcendental instructions



ET-Minion1 RISC-V Core and Tensor/Vector unit
optimized for low-voltage operation
to improve energy-efficiency

Optimized for energy-efficient ML operations. Each ET-Minion1 can deliver peak of 128 Int8 GOPS per GHz

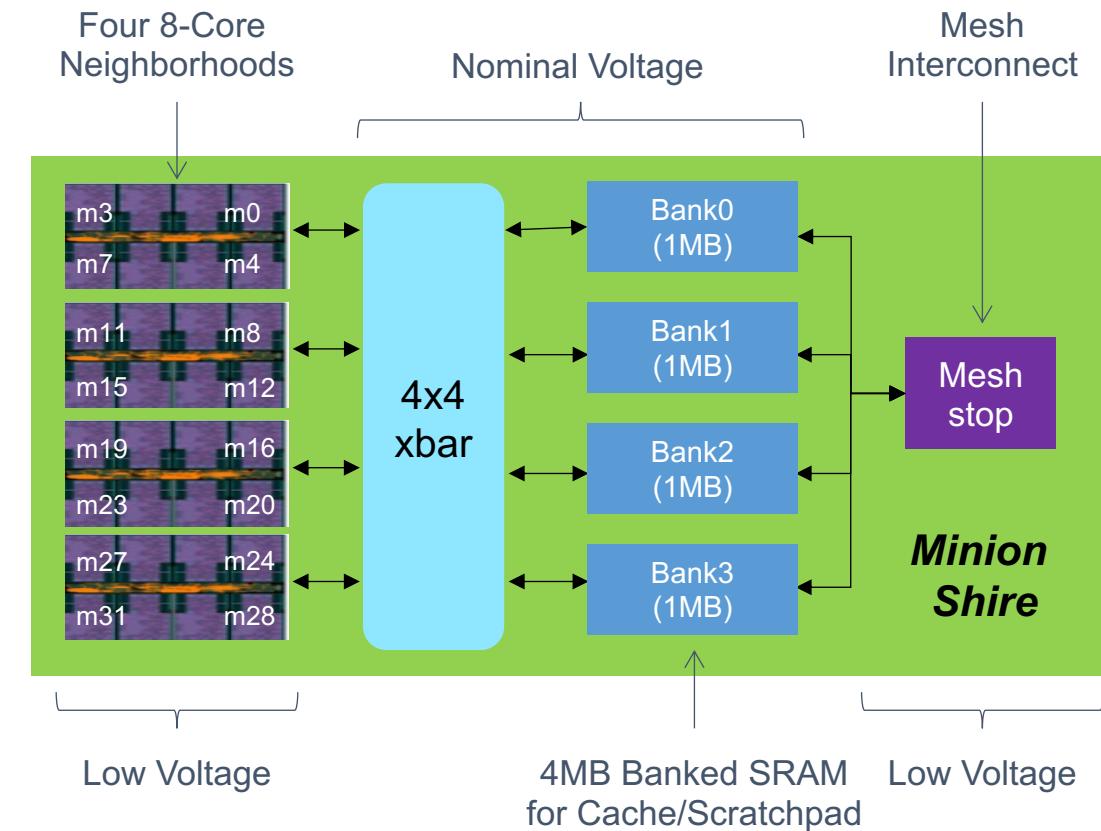
32 ET-Minion RISC-V cores per Minion Shire

- 64-bit RISC-V integer CPU cores each with 512-bit vector/tensor units
- Arranged in four 8-core neighborhoods

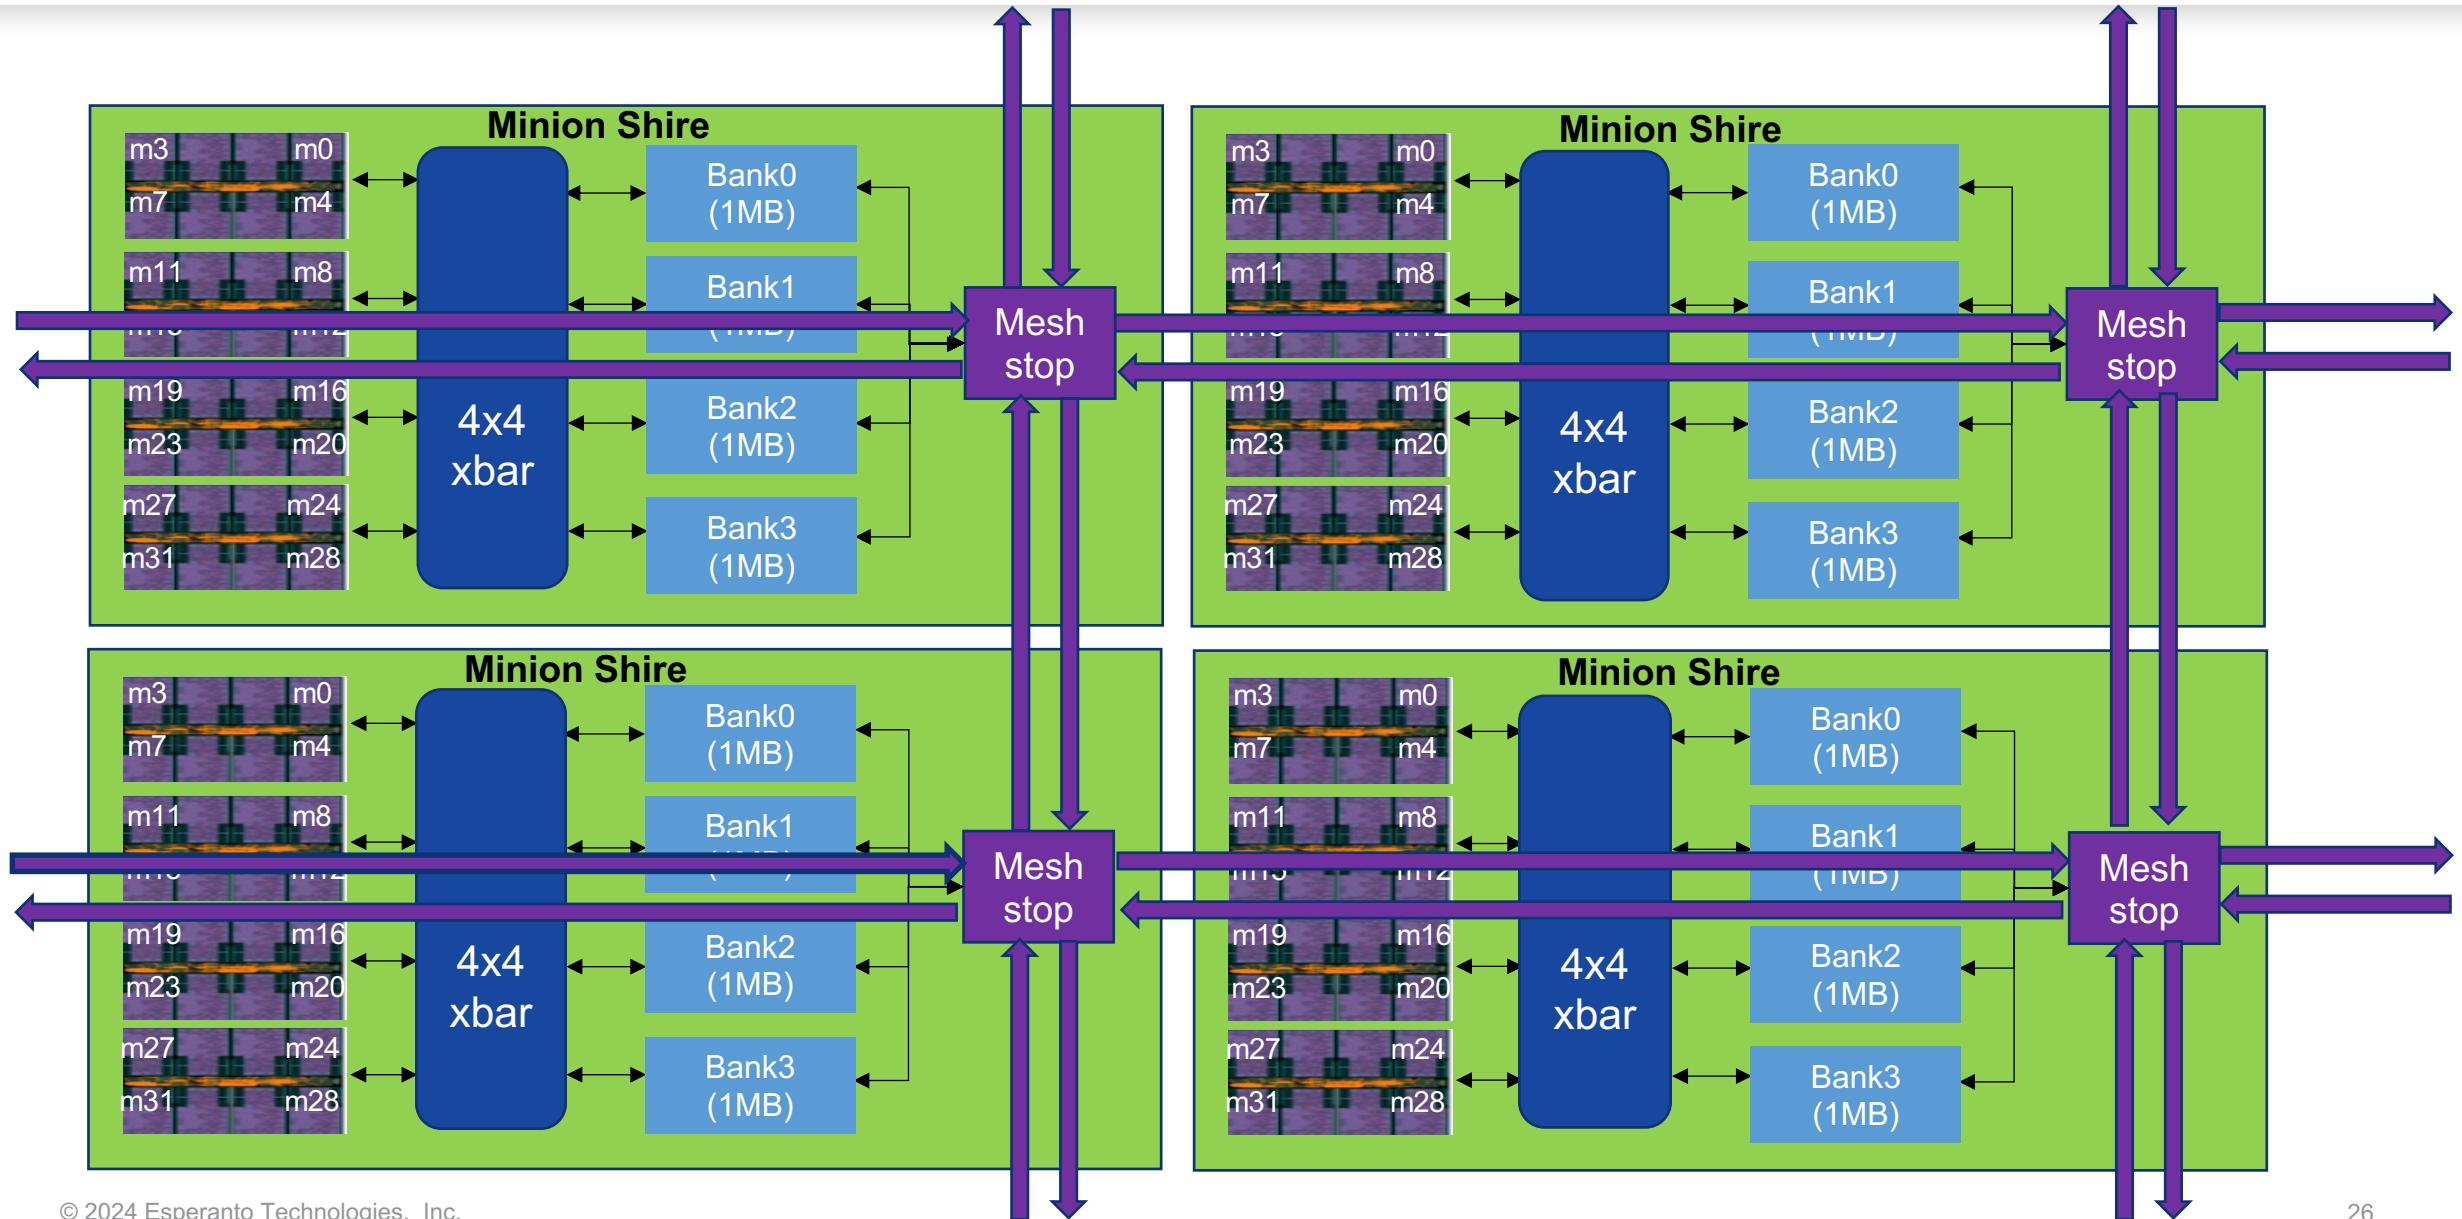
Software configurable memory hierarchy

- L1 data cache can also be configured as scratchpad
- Four 1MB SRAM banks can be partitioned as private L2, shared L3 and scratchpad

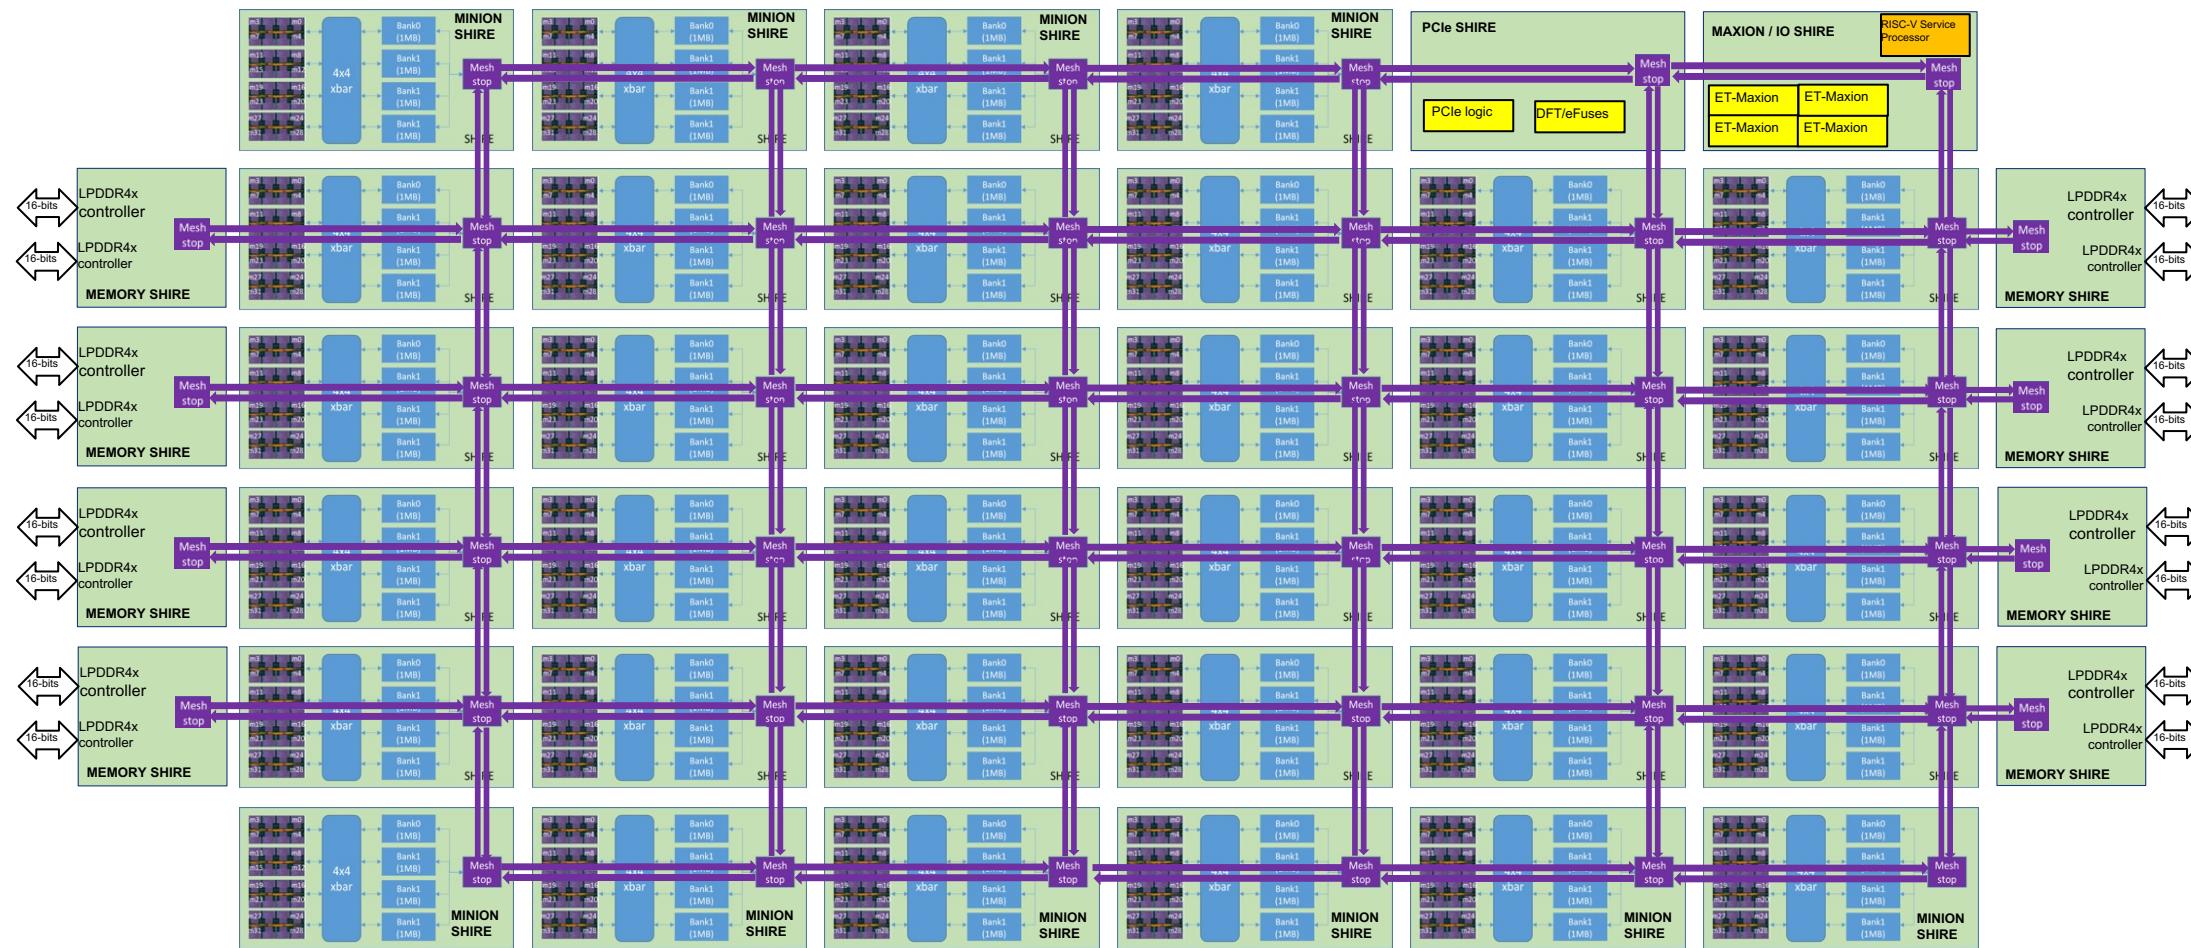
Shires connected with mesh network-on-chip



Shires are connected to each other and to external memory through Mesh Network



ET-SoC-1: Compute Fabric with Over a Thousand RISC-V CPU Cores



34 MINION SHIRES

- 1088 ET-Minions

8 MEMORY SHIRES

- LPDDR4x DRAM controllers
- up to 32 GB DRAM

1 MAXION / IO SHIRE

- 4 ET-Maxions
- 1 RISC-V Service Processor

PCIe SHIRE

160 million bytes of on-die SRAM

8 lanes of PCIe Gen4

Secure Root of Trust

The ET-SoC-1 is fabricated in TSMC 7nm

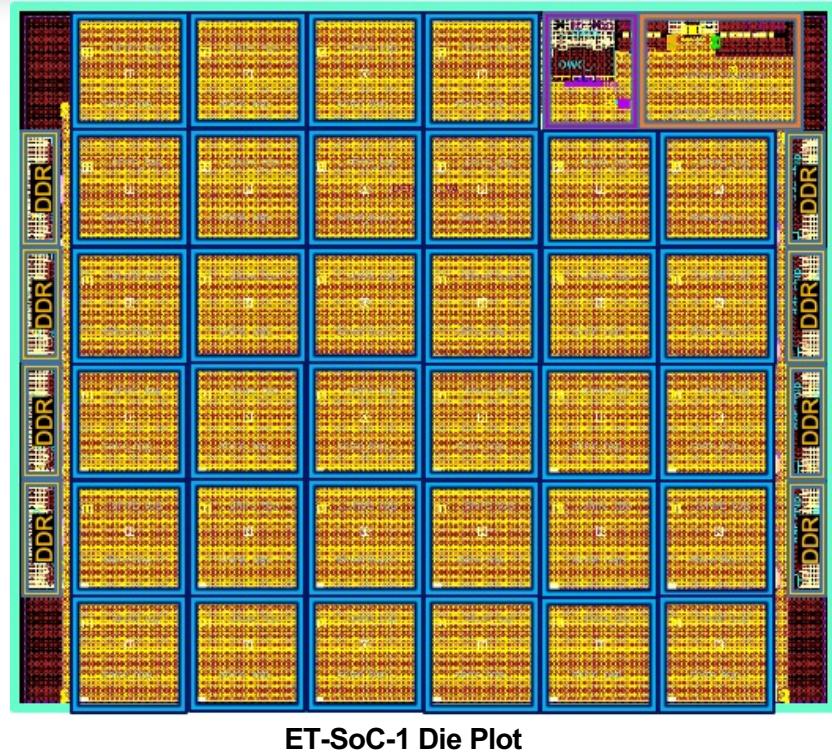
- 24 billion transistors
- Die-area: 570 mm²
- Low power: adjustable between 15 to 60 watts

Energy-efficient 64-bit RISC-V processors

- 1088 “ET-Minion” 64-bit RISC-V processors; vector/tensor unit
- 4 “ET-Maxion” 64-bit high-performance RISC-V out-of-order processors
- 1 64-bit RISC-V service processor

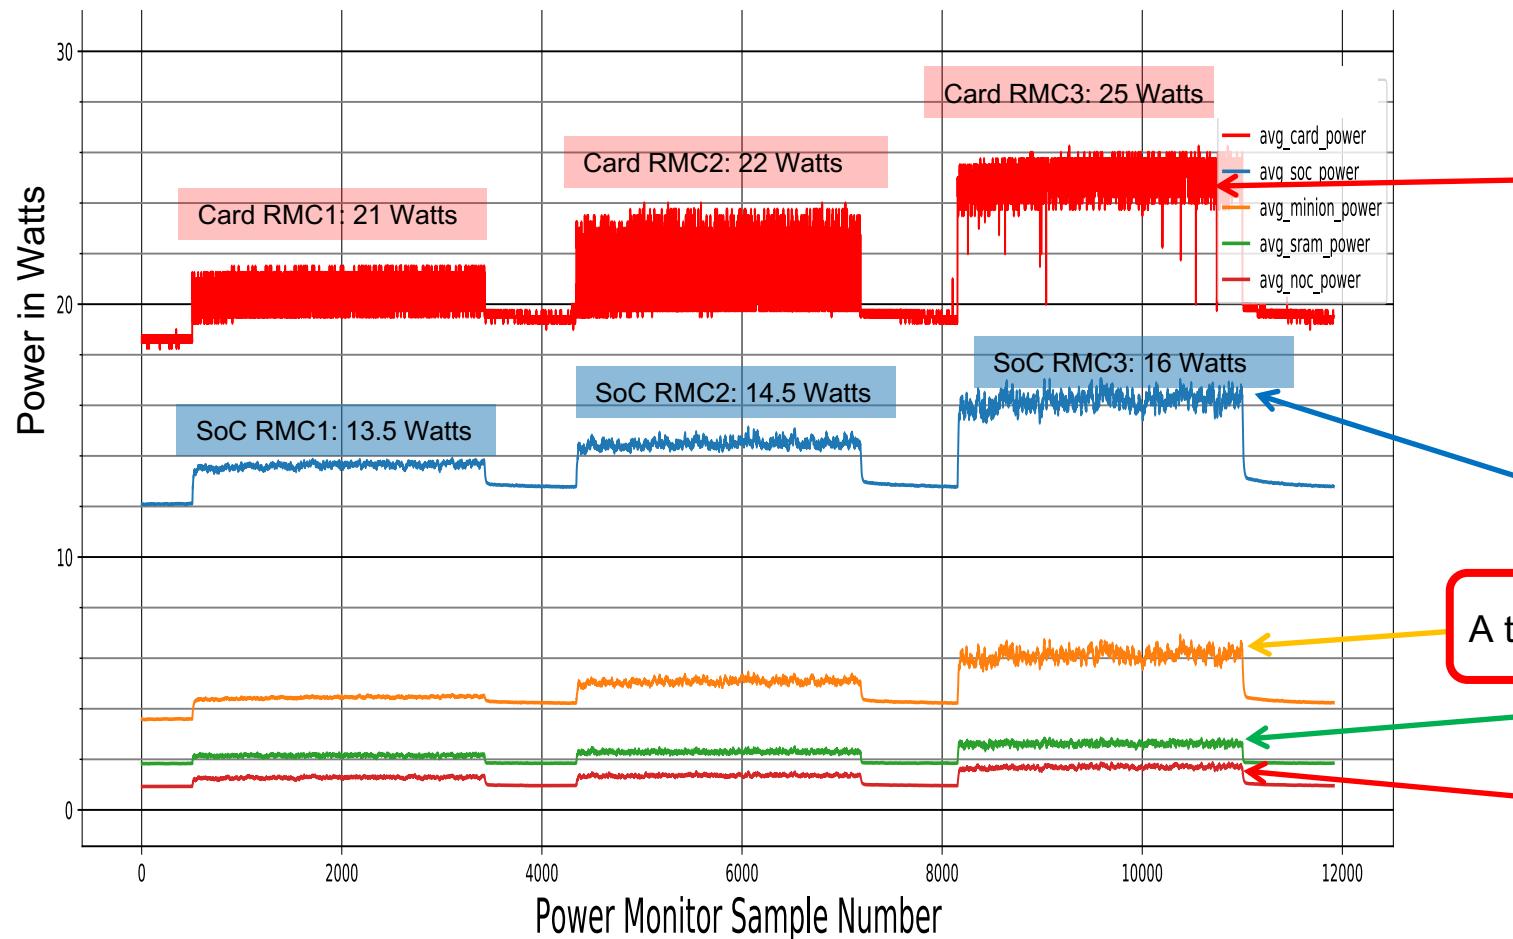
Other highlights

- 160MB of on-die SRAM used for caches and scratchpad memory
- Root of trust for secure boot
- Package: 45x45mm with 2494 balls to PCB, over 30,000 bumps to die



Example Card and ET-SoC-1 Power on ML Recommendation Benchmarks

DLRM FP16 RMC1, RMC2, RMC3



Power will vary depending on benchmark and operating conditions, but here is recently measured data.

Total PCIe card power under 25 watts for DLRM

- ET-SoC-1
- Eight LPDDR4x 32-bit wide DRAM die
- Flash memory and other misc components
- Power supplies

ET-SoC-1 power about 13.5 to 16 watts

A thousand 600 MHz ET-Minion RISC-V processors: 6 watts

128 MB of on-die SRAM for caches: 2.8 watts

44 Network-on-Chip connecting compute shires: 1.8 watts

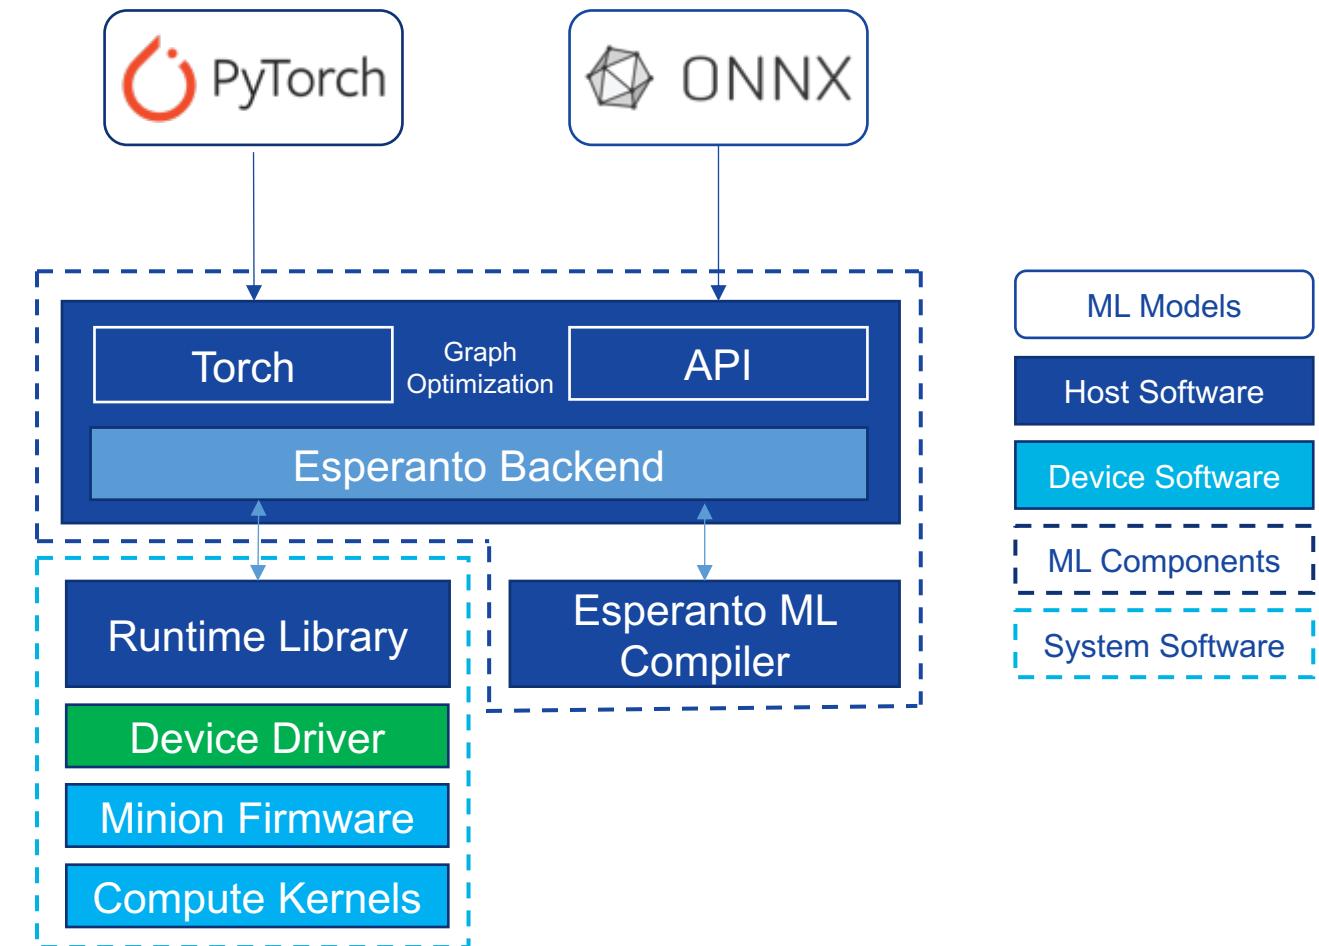
The Esperanto ML Software Stack

Delivered and integrated with each AI / HPC compute server

The ML framework compiler accepts a trained model expressed at a high level – PyTorch or TensorFlow (via ONNX) – and generates RISC-V executable code to run on ET-SoC-1(s)

Esperanto's ML stack performs device-independent graph optimization and then the Esperanto Backend perform device-specific optimization

Esperanto continues to adapt to market trends, new model directions and open-source software evolution



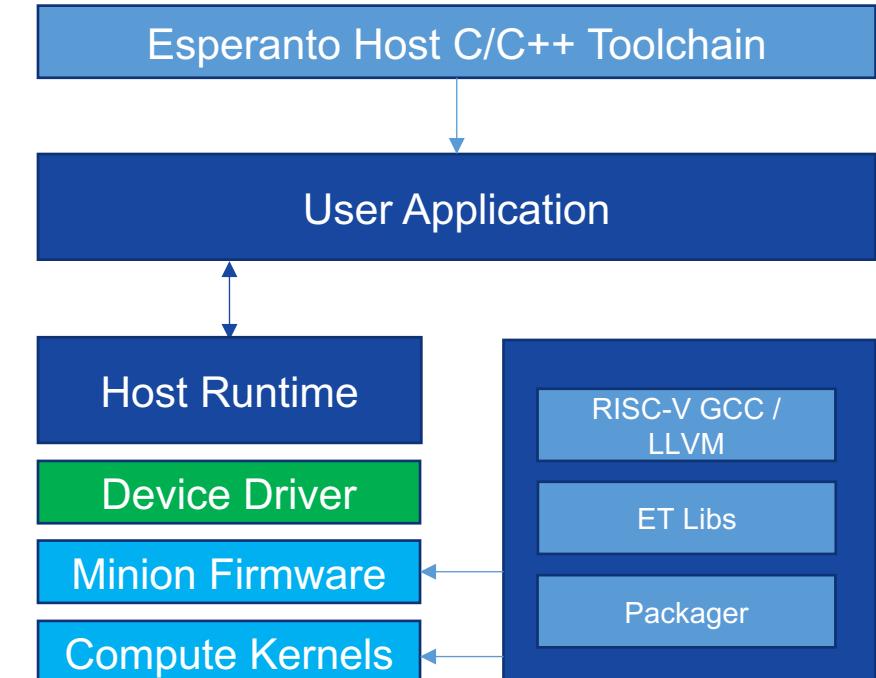
Comes integrated with each evaluation server

Enables general purpose parallel programming of 1024 64-bit ET-Minion cores

- All Minions can be individually programmed and can share common data address space
- Great for DSP and other highly parallel compute acceleration applications

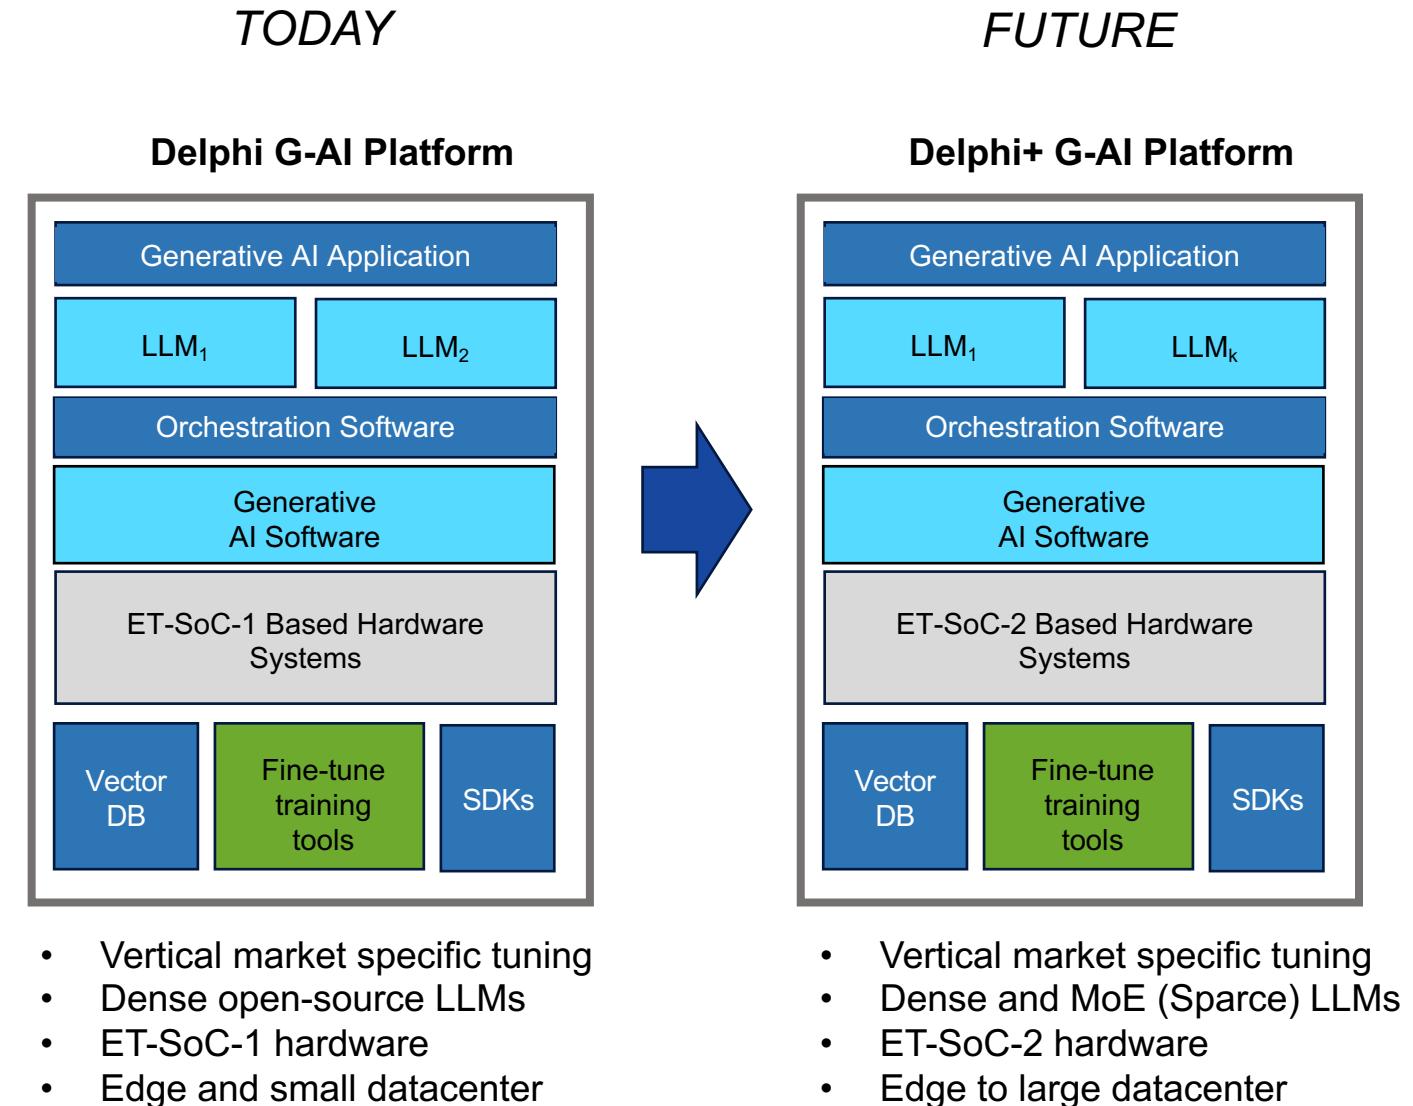
x86 host-based user application is compiled using Esperanto's Host C/C++ Toolchain

ET-Minion compute kernels and firmware are compiled using the RISC-V toolchain



- Esperanto strategy is focused on targeting ET-SoC-1 based systems with open-source models to bring differentiated value to vertical markets
 - Faster time to value and lower TCO
- Leverage Esperanto's technology for compute efficiency (low power) and scale-out computing
- Future support of MoE LLMs for inference as based model size is projected to increase beyond 1T parameters¹

¹Note: Nvidia's recent Blackwell based GPU systems are targeted at enabling training of 1T+ MoE LLMs



Generative AI Language Models Supported by Esperanto



Use Case	LLM Model	Parameters	Precision	Open-sourced since	ET Supported since
Text-to-text	Llama2, Vicuna	7B, 13B	fp16, int4	April 2023	September 2023
	Mistral	7B	fp16, int4	September 2023	November 2023
	Llama3.1	8B	fp16, int4	July 2024	August 2024
	Phi-3.5	Mini	fp16	August 2024	August 2024
	Gemma	2B, 7B	fp16	February 2024	June 2024
	Mixtral	8x7B	int4	January 2024	August 2024
Text-to-code	Starcoder	15B	fp16	May 2023	November 2023
Text-to-image	OpenJourney		fp16, fp32	November 2022	September 2023
	Stable Diffusion	1.5, 2.1	fp16	December 2022	November 2023
Image-to-text	LLAVA	1.5	fp16, int4	October 2023	July 2024
Speech-to-text	Whisper	V3 Medium, Large	fp16	November 2023	July 2024

PCIe Gen4 8-lane low-profile form factor

ET-SoC-1 and 32 GB LPDDR4x memory

Safety and emissions certification for multiple geographies

Developed, tested, certified and delivered by Esperanto's partner Penguin Solutions / Smart Global Holdings (SGH)

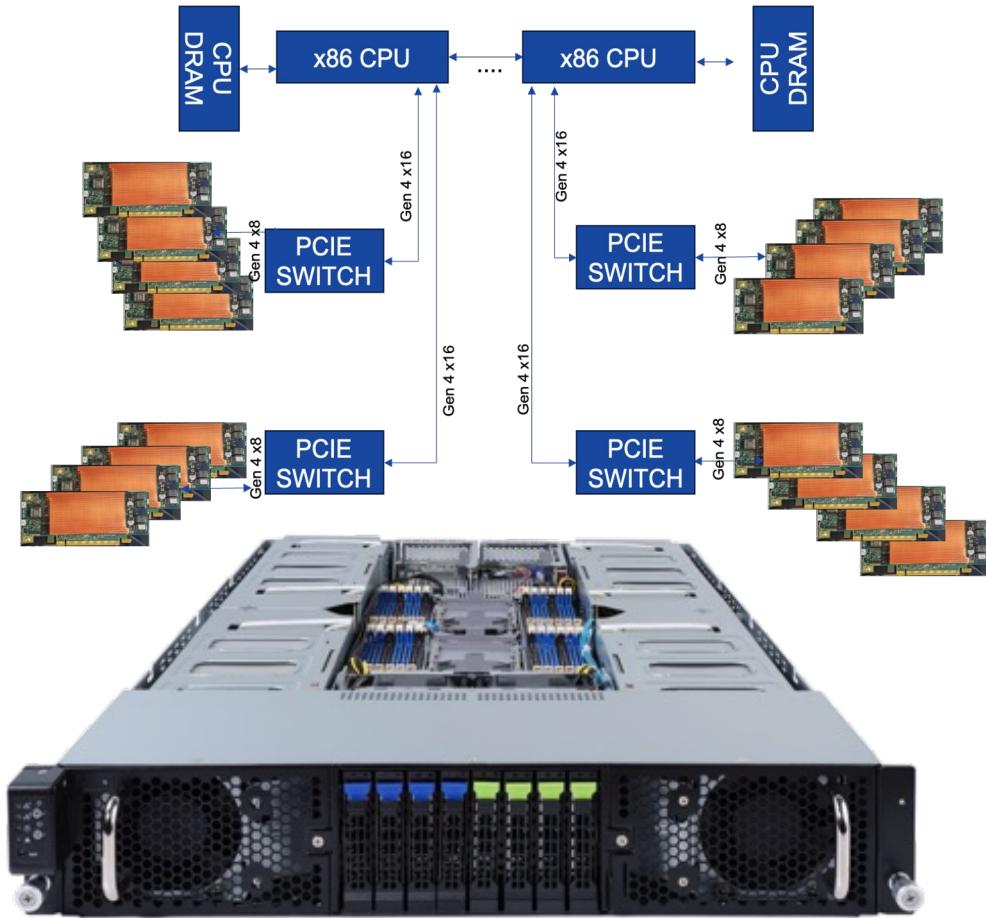
Implemented across various server form factors

- Single card (ruggedized edge server)
- Up to 6 cards (short-form 2U server for enterprise edge)
- Up to 16 cards (ultra-high density 2U server for data center)



Penguin Computing production ET-SOC-1 PCIe card with and without heatsink

Esperanto's ET-SoC-1 RISC-V Production System – Shipping Now



2U server with up to 16 Esperanto Accelerator Cards

Target environment	Data center	
Server configuration	Standard 2U 19" rack-mount chassis	
ET-SoC-1 PCIe cards (Gen4x8)	8 Cards	16 Cards
System host processor	2x Intel Xeon® Gold 6326 16-core	2x Intel Xeon Platinum 8358P 32-core
System memory	512GB DDR4-3200	1TB DDR4-3200
Storage	2x Samsung® PM9A3 3.84TB NVMe U.2 SSDs	
ET-SoC-1 frequency	600 MHz	
ET-SoC-1 power consumption	15W to 60W (workload dependent)	
ET-SoC-1 RISC-V CPUs	8,448 or 16,896 (high performance, low power)	
Pre-installed AI models	LLMs, Computer Vision, Recommendation, Transformer models	
ML SDK	Jupyter Notebook and command-line tools	
GP SDK	Included	
Performance, power, and trace analysis tools	Included (et-powertop, Perfetto)	
Training and documentation	Included	
Connectivity	2 x 10Gb/s LAN ports built in, 2 free PCIe Gen4 x8 slots	

Example ET-SoC-1 Rack Implementations

	Low Power	High Density	Ultra High Density	Notes
ET chip	ET-SoC-1	ET-SoC-1	ET-SoC-1	
Enclosure	GigaByte G292	GigaByte E252	GigaByte G292	US country of origin
Host	x86	x86	x86	"Self hosting" optional
PCIe	Gen 4	Gen 4	Gen 4	
Server form factor	2U	2U	2U	
# of cards / server	8	7	16	Customizable
# of server / rack	10	20	20	Customizable
# of 64bit RISC-V cores	80,000+	140,000+	320,000+	Low power cores
Vector / Tensor	Yes	Yes	Yes	
Rack power	8 kW*	16 kW**	35kW**	* measured average; DLRM @ FP16; ** estimated peak
Int8	6 PetaOps	10.5 PetaOps	24 PetaOps	Est. peak
FP16	1.5 PetaFlops	2.6 PetaFlops	6 PetaFlops	Est. peak
FP32	0.75 PetaFlops	1.3 PetaFlops	3 PetaFlops	Est. peak
Interconnect	Ethernet / InfiBand	Ethernet / InfiBand	Ethernet / InfiBand	Customizable
Software	Dual SDKs included	Dual SDKs included	Dual SDKs included	AI inference; single precision HPC
Pre-installed LLMs	Optional	Optional	Optional	Current: Vicuna, Llama, Mistral, StarCoder, OpenJourney, StableDiffusion, LLaVa
Analytics tools	Included	Included	Included	
Support / warranty	Included	Included	Included	Via Smart Global Holdings / Penguin Computing



80,000 low power RISC-V processors in Esperanto's Silicon Valley office

ESPERANTO'S NEXT GENERATIONS:

**Improvements for an energy-efficient Chiplet Compute Fabric
based on the latest RISC-V Vector Instruction Set
for both HPC and AI applications**



Starts with a

- Custom ET designed
- 64-bit RISC-V ISA compatible
- General-purpose
- FMA with 2 ops/cycle
- Simple in-order pipeline
- Low area
- Low power
- Dual hardware threads
- Supports a globally shared address space
- Circuits that work at low-Voltage
- Processor Core
- Esperanto calls this the ET-Minion2

With 1 GHz cores

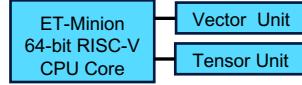
Int8: .002 TOps

FP8:

FP16:

FP32: .002 TFlops

FP64: .002 TFlops



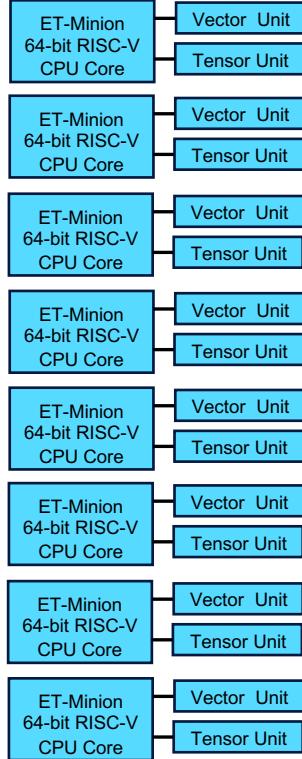
Add a

- RISC-V Vector Unit
 - 512 bits per cycle all RISC-V data types
- Tensor Unit
 - 1024 bits/cycle FP64/FP32
 - 2048 bits/cycle Int8/FP8/FP16

With 1 GHz cores
Int8: 0.512 TOps
FP8: 0.512 TFlops
FP16: 0.256 TFlops
FP32: 0.064 TFlops
FP64: 0.032 TFlops

Tensor FMA operations dominate power and area for ML,
meaning that providing the general-purpose RISC-V ISA has little penalty
if a simple, small and energy efficient design is used.

Compute Cluster for Parallel Programming (step 3 of 5)



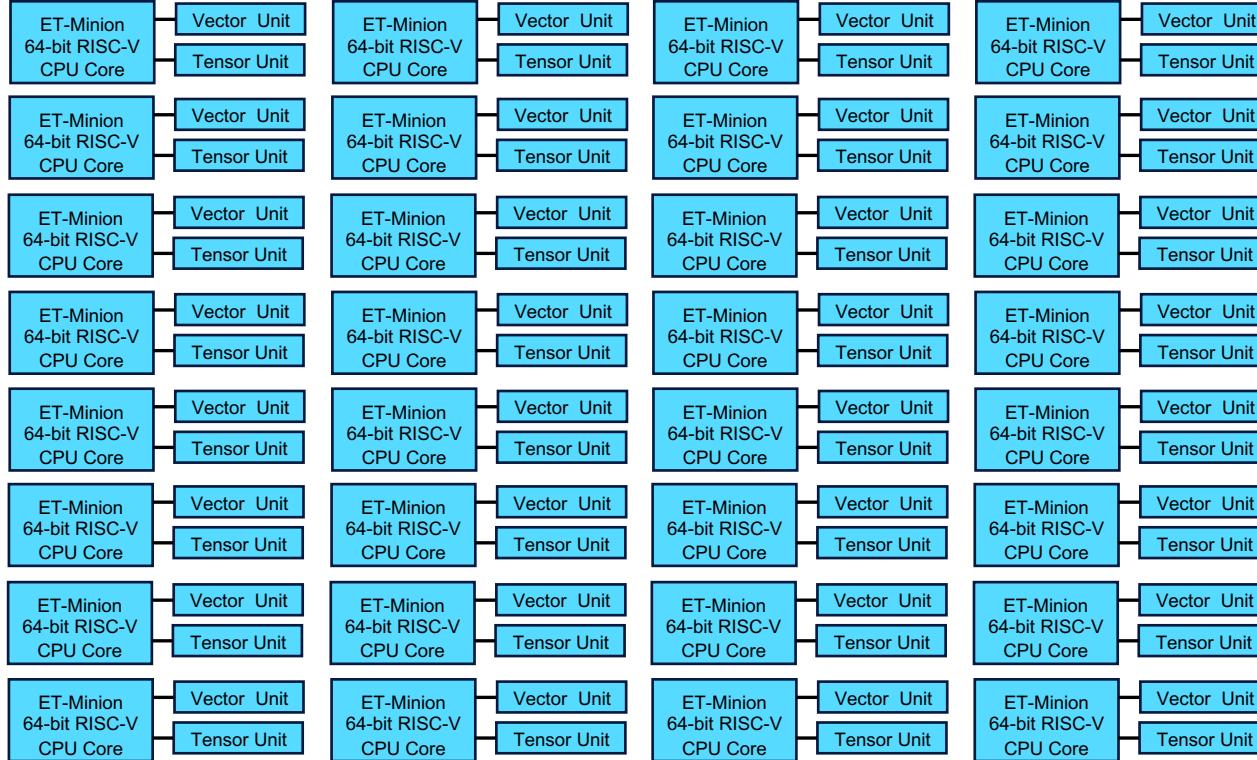
Add 8 ET-Minions each with

- Vector Unit
 - 512 bits per cycle all RISC-V data types
- Tensor Unit
 - 1024 bits/cycle FP64/FP32
 - 2048 bits/cycle Int8/FP8/FP16

With 1 GHz cores
Int8: 4.096 TOps
FP8: 4.096 TFlops
FP16: 2.048 TFlops
FP32: 0.512 TFlops
FP64: 0.256 TFlops

Compute Cluster for Parallel Programming (step 4 of 5)

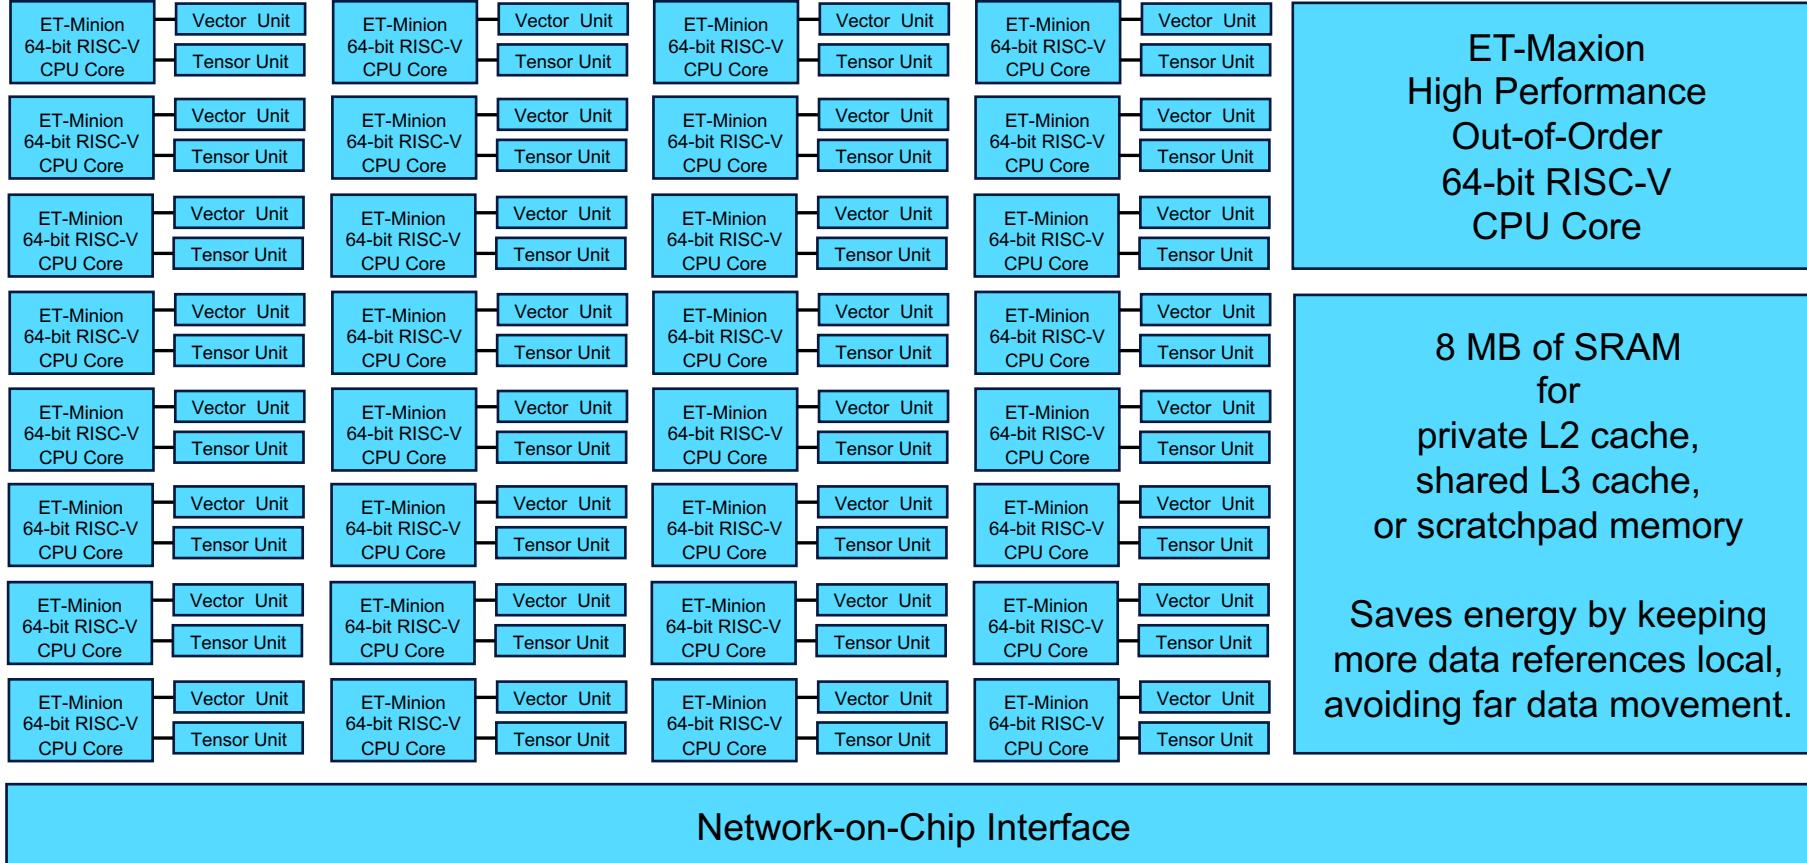
Build up to 32 ET-Minion Processors



With 1 GHz cores
Int8: 16.384 TOps
FP8: 16.384 TFlops
FP16: 8.192 TFlops
FP32: 2.048 TFlops
FP64: 1.024 TFlops

Compute Cluster for Parallel Programming (step 5 of 5)

Build up to 32 ET-Minion Processors



With 1 GHz cores
Int8: 16.384 TOPs
FP8: 16.384 TFlops
FP16: 8.192 TFlops
FP32: 2.048 TFlops
FP64: 1.024 TFlops

This is the cluster architecture for programmers/compiler
which is then replicated

Add compute clusters to form a compute fabric on a die



With 1 GHz cores

Int8: 131.1 TOPs

FP8: 131.1 TFlops

FP16: 65.5 TFlops

FP32: 16.4 TFlops

FP64: 8.2 TFlops

(note – only using
8 of the clusters,
1 is a spare)

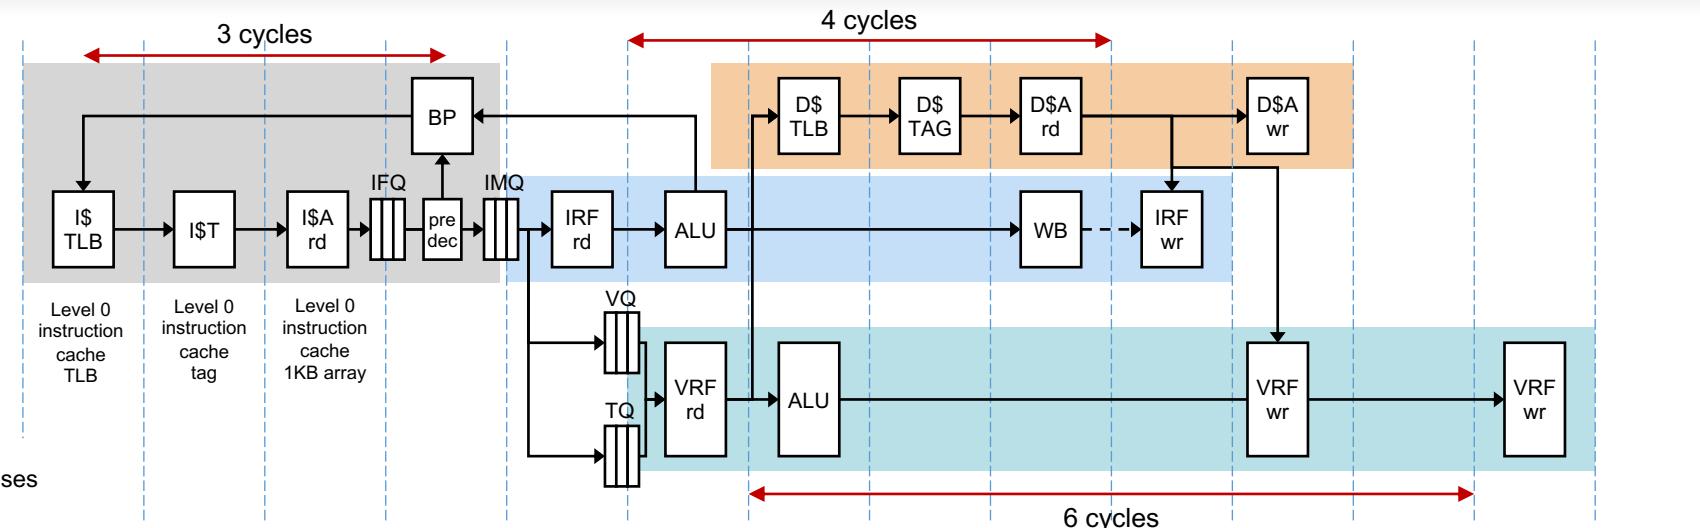
ET-SoC-2x
4nm

Add full support for HPC and larger AI tasks

New ET-Minion2 CPU: Small, low power and high vector/tensor throughput

Minion2 Scalar Pipe

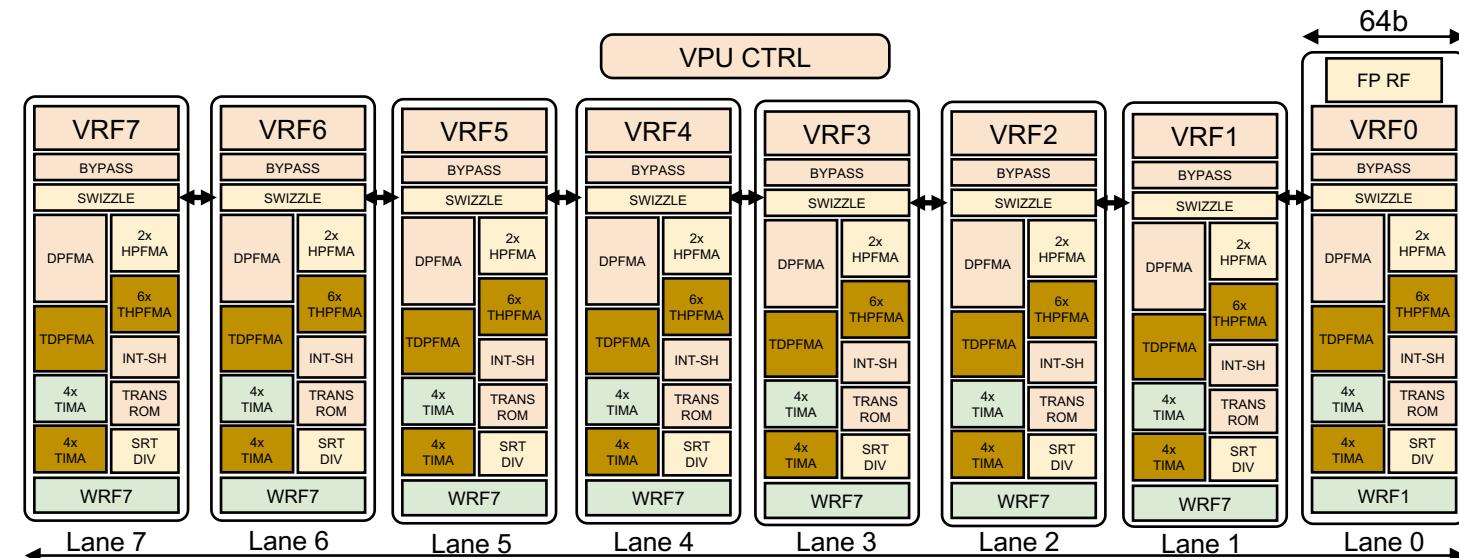
- 64-bit RISC-V compatible CPU up to 2 GHz
- 2 hardware threads
- Low-power in-order core
- Short 3 cycle branch bubble if unpredicted
- Decoupled Vector/Tensor unit
- In-order appearance of Tensor instructions
- 4 KB L1 Data Cache
- 4-way issue
 - Memory
 - Int
 - FP/Vector
 - Tensor
- 48-bit Virtual Addresses / 46-bit Physical Addresses



Minion2 Vector Unit

- RISC-V RVV Compatible
- Hand optimized for speed and low power
- 4 KB Vector Register File (32 64B reg x 2 threads)
- 16 KB Tensor Register File (16 1KB Tensor regs)

Data Type	Vector (RVV instructions) Ops/cycle	Tensor (Tensor instructions) Ops/cycle
FP64 DP	16	32
FP32 SP	32	64
FP16/BF16 HP	64	256
FP8 QP	128	512
INT8	64	512



New ET-SoC-2 SuperShire/Compute Cluster :

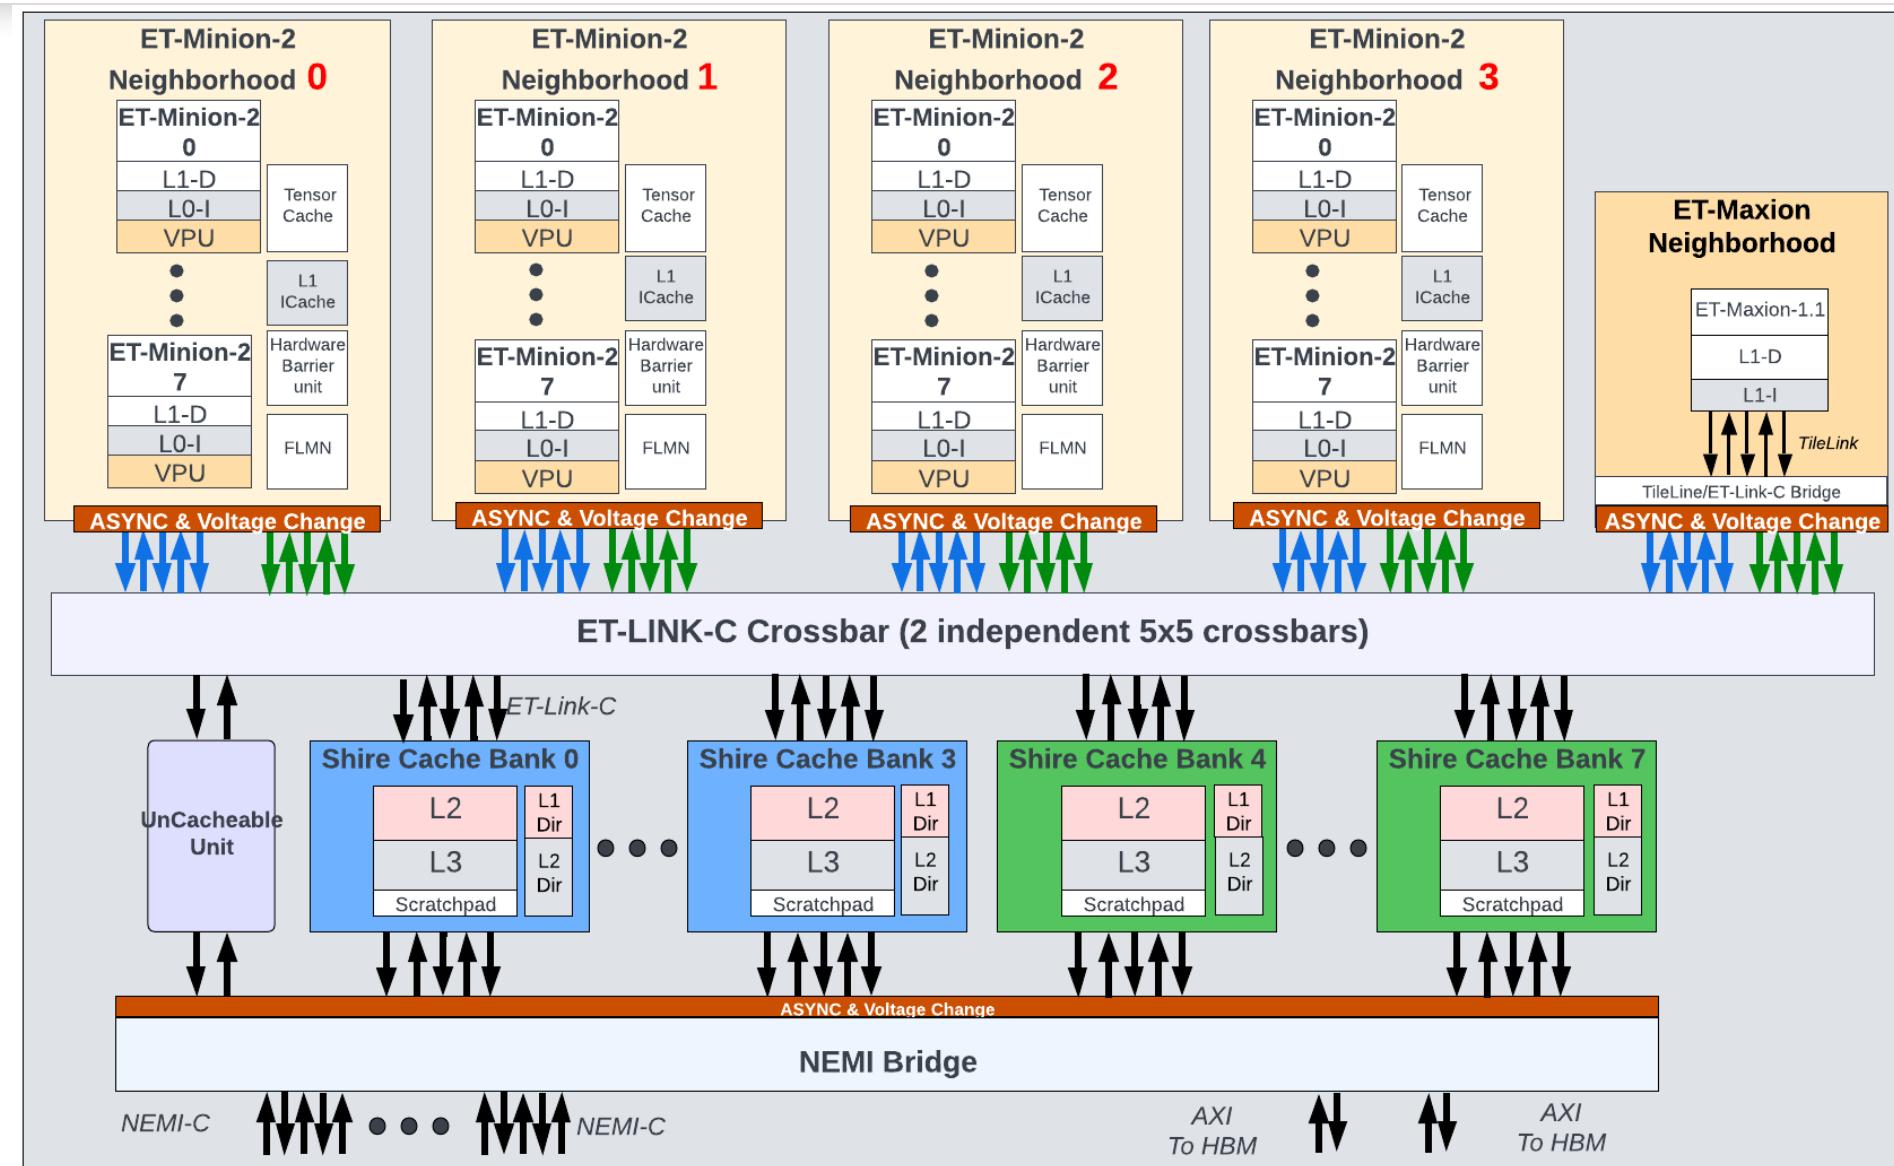
32 ET-Minion-2s, 1 ET-Maxion, 8 MB SRAM and NEMI NoC

Compute Fabric is built out of multiple **SuperShires**.

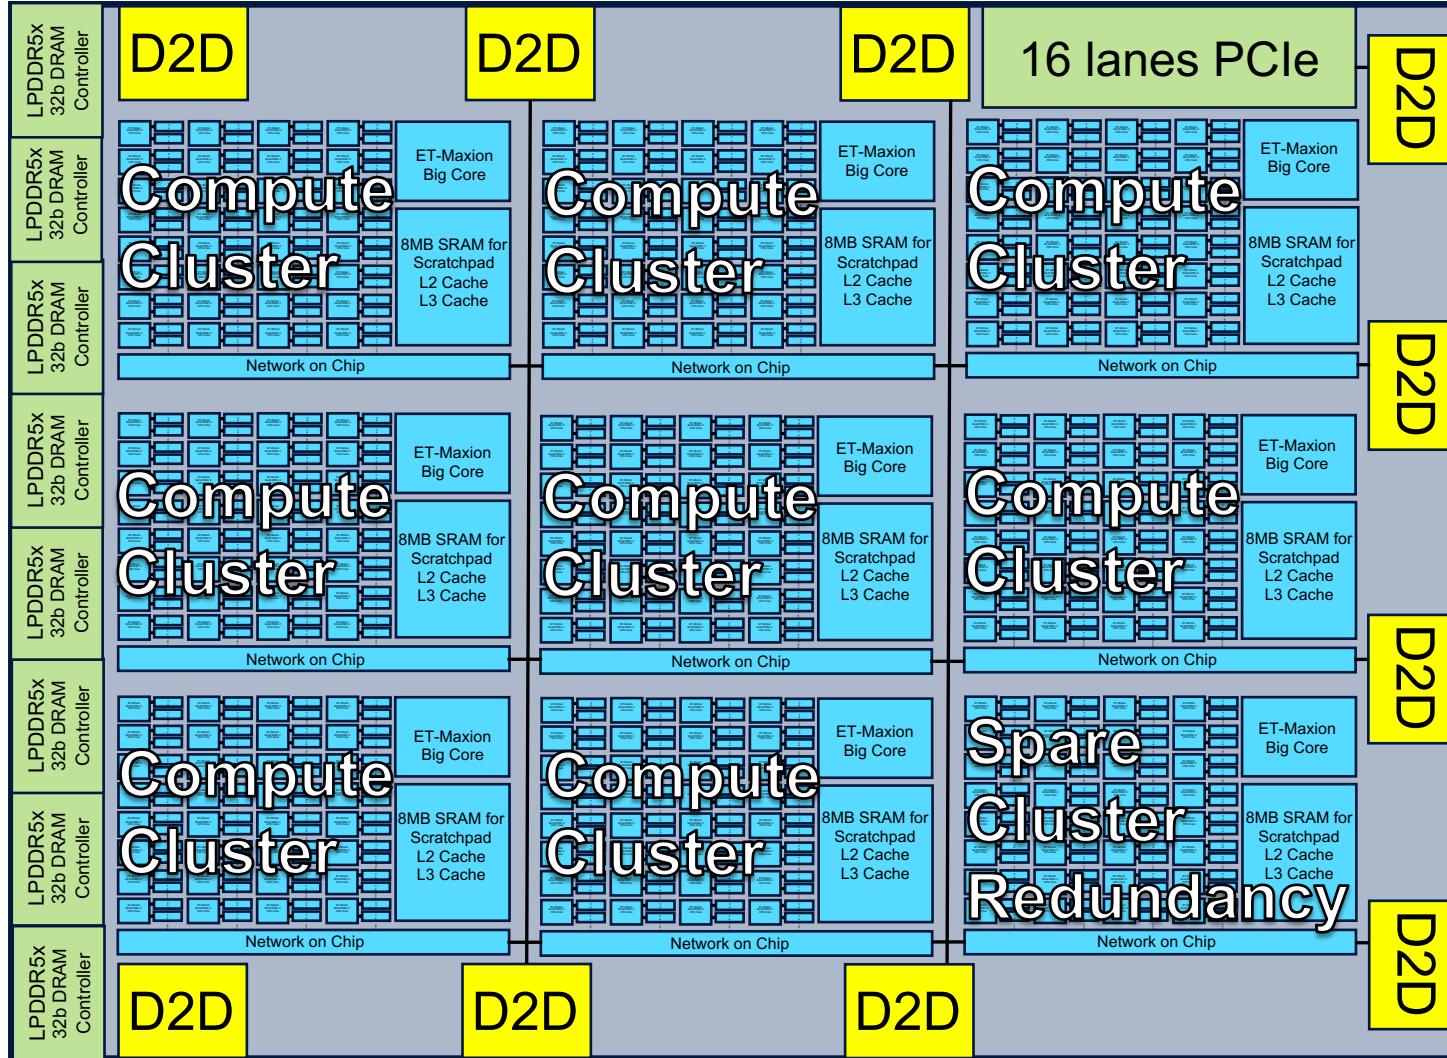
There are **9** SuperShires in the 4nm ET-SoC-2x.

Each SuperShire has:

- **1** ET-Maxion CPU
- **32** ET-Minion2 CPUs (Grouped in 4 Neighborhoods)
- Eight 1 MB SRAM Banks (software configurable as L2, L3 or scratchpad memory)
- Two 5x5 crossbar links
- NEMI NoC Bridge to other SuperShires or IO



ET-SoC-2x 4nm Compute Fabric Chiplet



ET-SoC-2x chiplet has:

- 8 Compute Clusters + 1 Spare
- 256 ET-Minion2 processors
- 8 ET-Maxion processors
- 64 MB on die SRAM
- 256 bit wide LPDDR5x-8500 controllers
- 16 lanes of PCIe Gen5

8 active compute clusters, each with

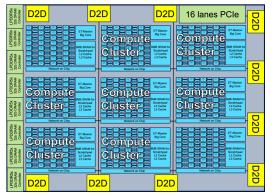
- One ET-Maxion “big core”
- 32 Minion2 “efficiency cores”
 - Each with its own Vector and Tensor units
- 8 MB of SRAM for L2, L3 and scratchpad
- Network on Chip to other compute clusters

1 compute cluster used for redundancy to increase yield and lower cost.

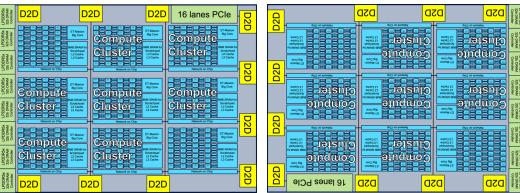
Die to die (D2D) interfaces allow fabric to grow each has similar bandwidth as on die NoC to NoC

ET-SoC-2x Allowed Compute Fabric Chiplet Configurations

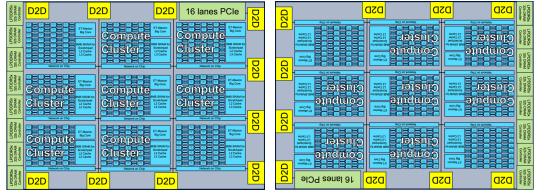
ET-SoC-2x1



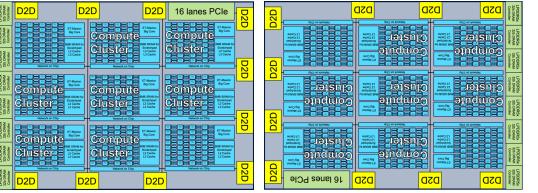
ET-SoC-2x2



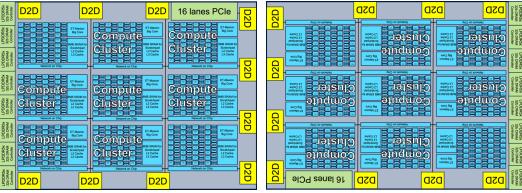
ET-SoC-2x4



ET-SoC-2x6



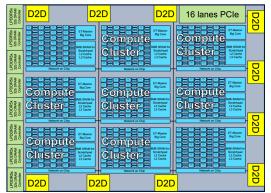
ET-SoC-2x8



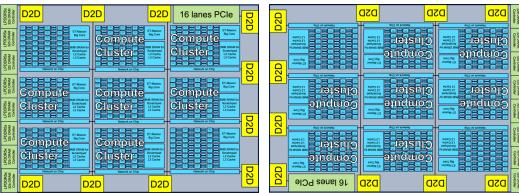
	ET-SoC-2x1	ET-SoC-2x2	ET-SoC-2x4	ET-SoC-2x6	ET-SoC-2x8
# Chiplets in package	1	2	4	6	8
LPDDR5x lanes	256	512	1024	1536	2048
Memory GB(Rank=4)	256	512	1024	1536	2048
Memory Bandwidth GB/s	272	544	1088	1632	2176
PCIe Gen 5 Lanes	16	32	64	96	128
PCIe Bandwidth GB/s	64	128	256	384	512
Peak TFlops	2				
FP8 or Int8	262	524	1048	1572	2096
FP16 / BF16	131	262	524	786	1048
FP32	32.8	65.6	131.2	196.8	262.4
FP64	16.4	32.8	65.6	98.4	131.2

ET-SoC-2x Allowed Compute Fabric Chiplet Configurations

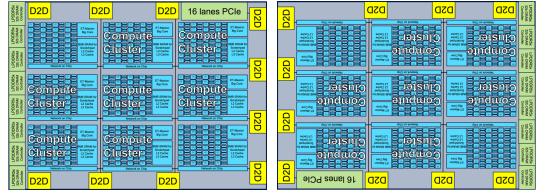
ET-SoC-2x1



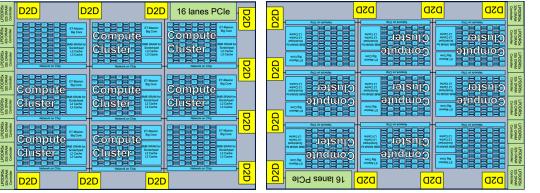
ET-SoC-2x2



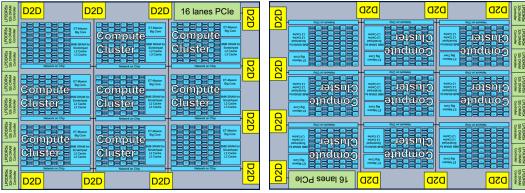
ET-SoC-2x4



ET-SoC-2x6

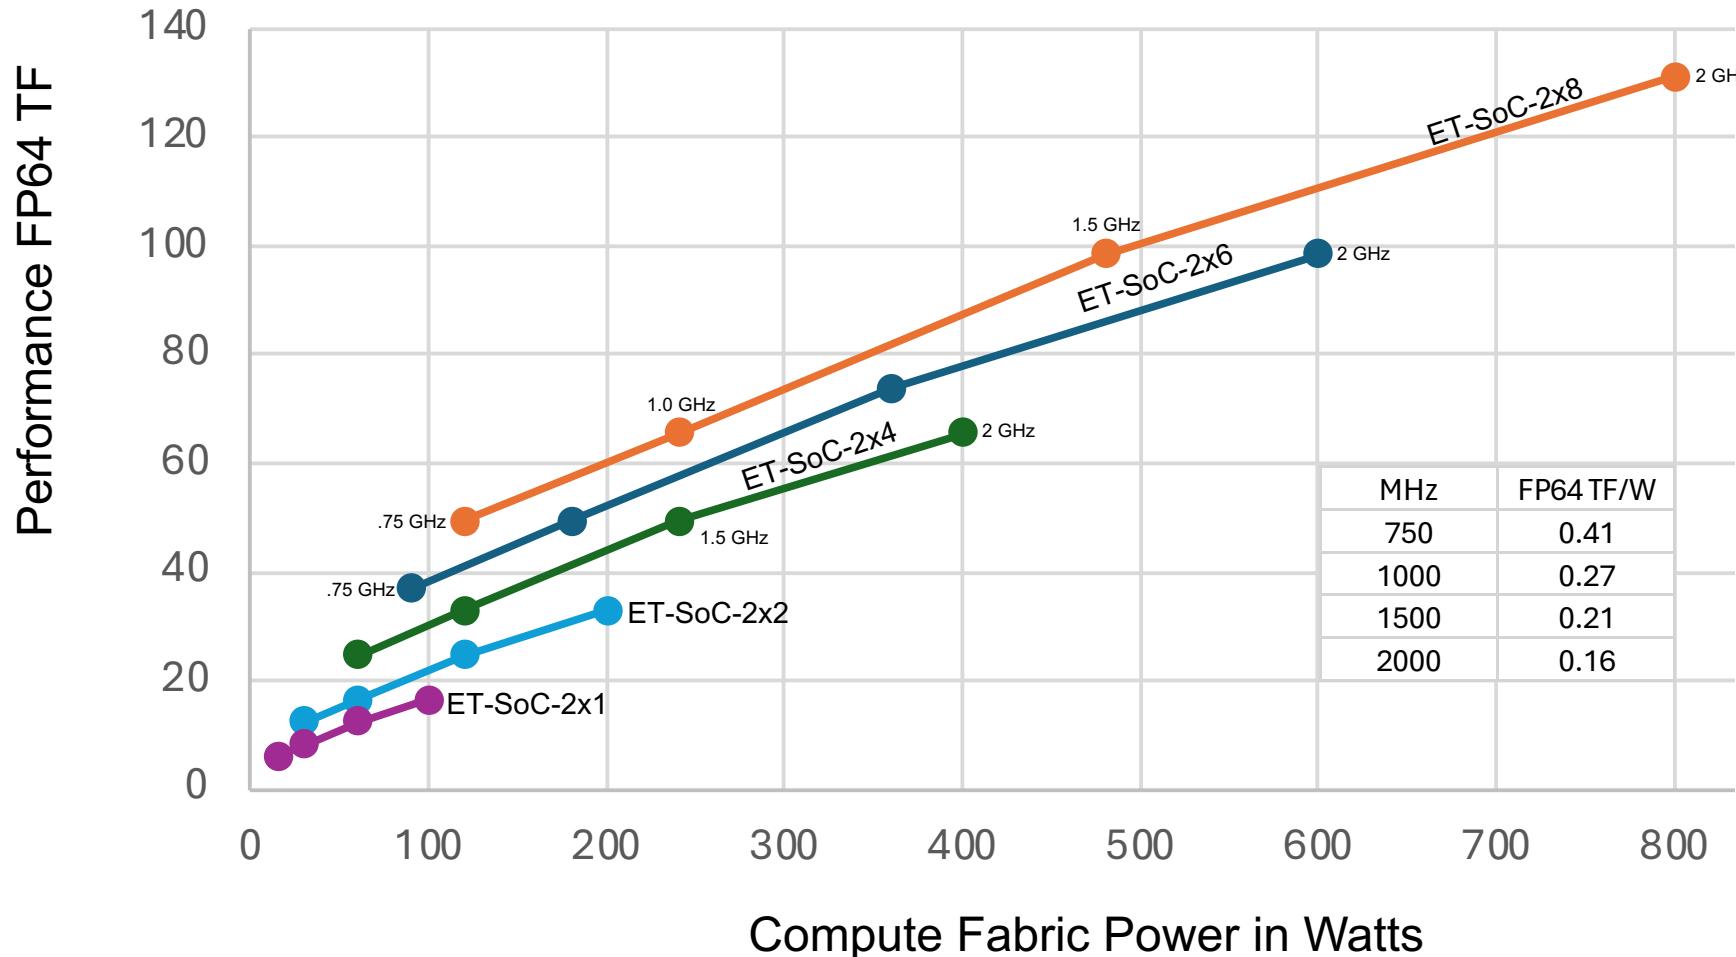


ET-SoC-2x8

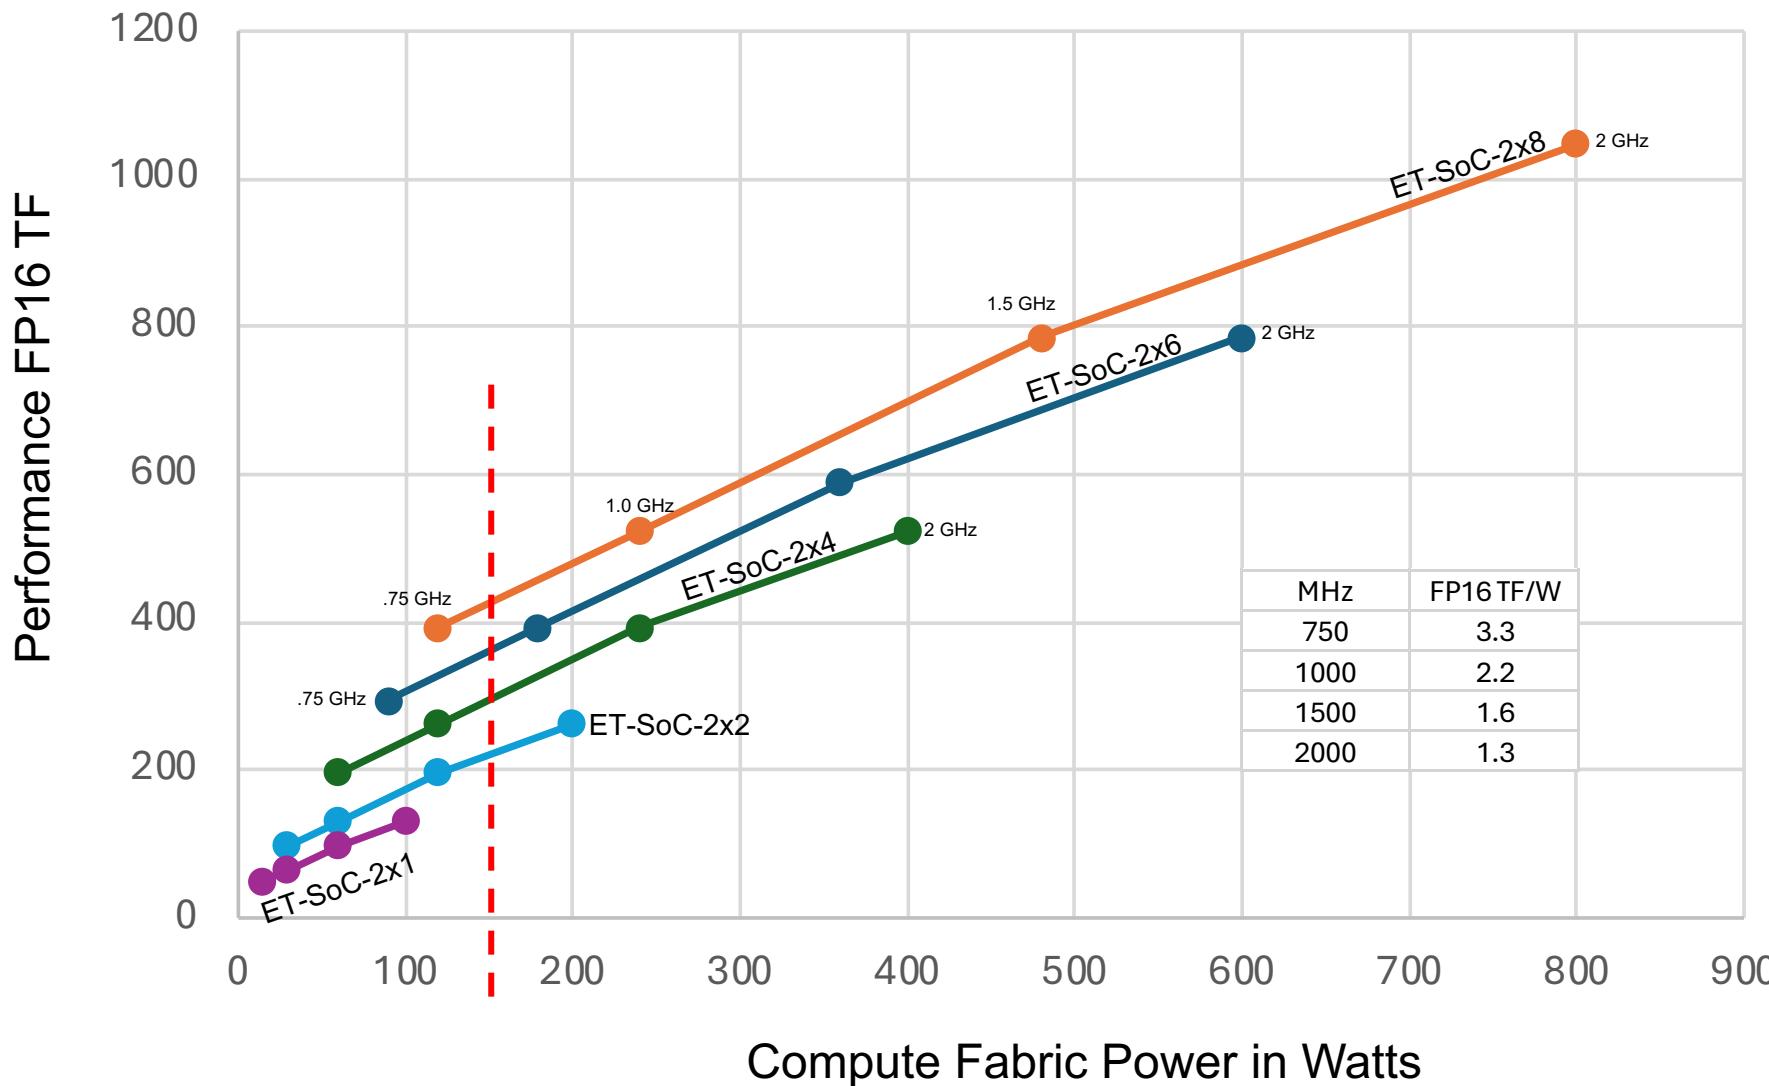


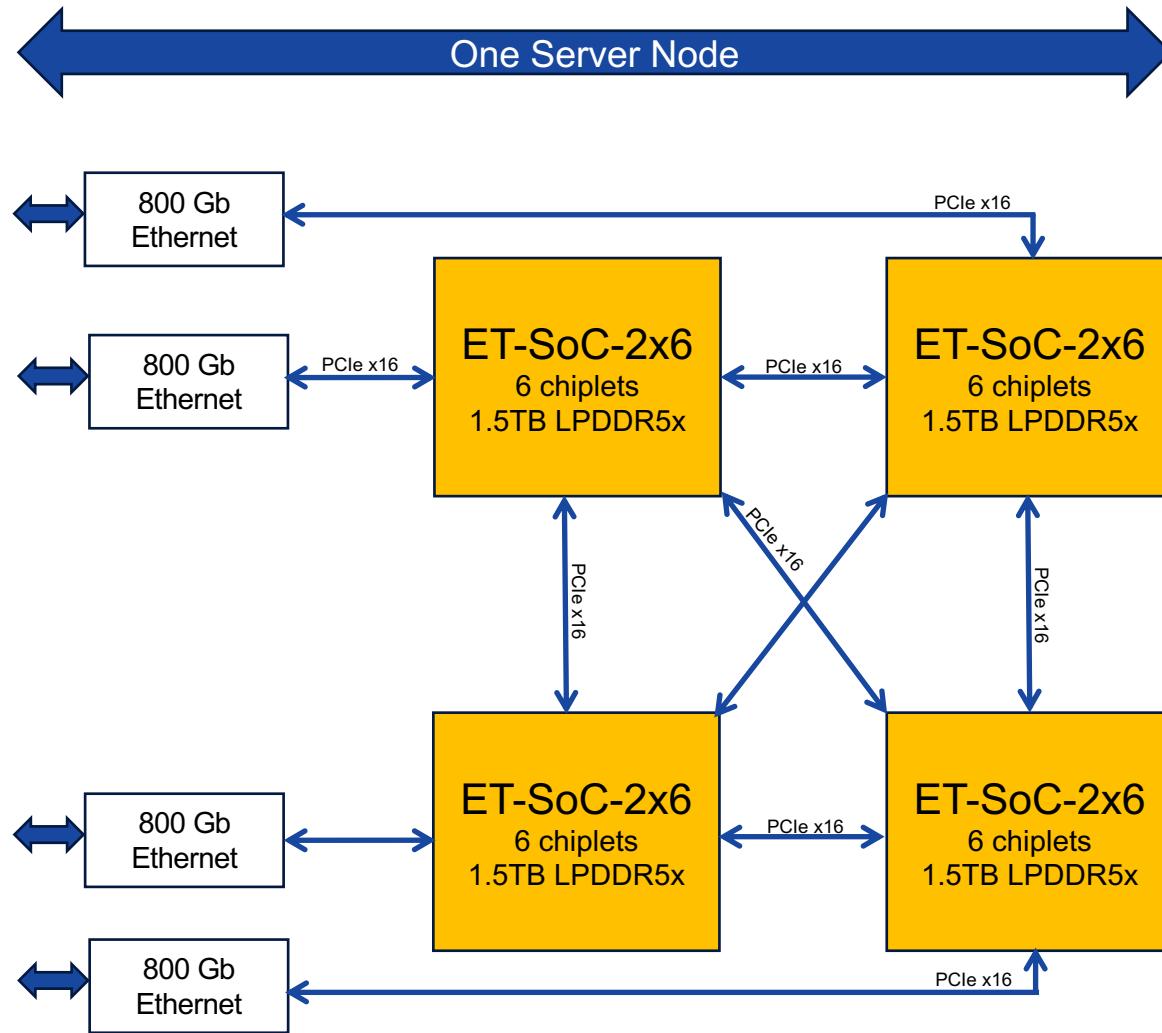
	ET-SoC-2x1	ET-SoC-2x2	ET-SoC-2x4	ET-SoC-2x6	ET-SoC-2x8
# Chiplets in package	1	2	4	6	8
LPDDR5x lanes	256	512	1024	1536	2048
Memory GB(Rank=4)	256	512	1024	1536	2048
Memory Bandwidth GB/s	272	544	1088	1632	2176
PCIe Gen 5 Lanes	16	32	64	96	128
PCIe Bandwidth GB/s	64	128	256	384	512
Peak TFlops	2				
FP8 or Int8	262	524	1048	1572	2096
FP16 / BF16	131	262	524	786	1048
FP32	32.8	65.6	131.2	196.8	262.4
FP64	16.4	32.8	65.6	98.4	131.2

FP64 Performance vs Power



FP16 Performance vs Power for different 4nm ET-SoC-2 Packages





Four CPU packages on a motherboard per node

Each package has six ET-SoC-2x chiplets

All CPUs share single address space and memory

6144 RISC-V vector/tensor processors

6 TB of shared DRAM

Peak performance per server:

- 384 TF FP64
- 768 TF FP32
- 3072 TF FP16
- 6144 TF FP8

Performance per chiplet doubles with 2nm ET-SoC-3

About 36 racks to provide 1 ExaFlop DP peak

ET-SoC-3x

2nm

Main changes to ET-SoC-3x Chiplet



Similar software interface between 4nm ET-SoC-2x and 2nm ET-SoC-3x

2nm technology allows 512 ET-Minion-3 RISC-V processors, double that of 4nm ET-SoC-2x

Improved memory system to support up to 2 TB of useable bandwidth per chiplet

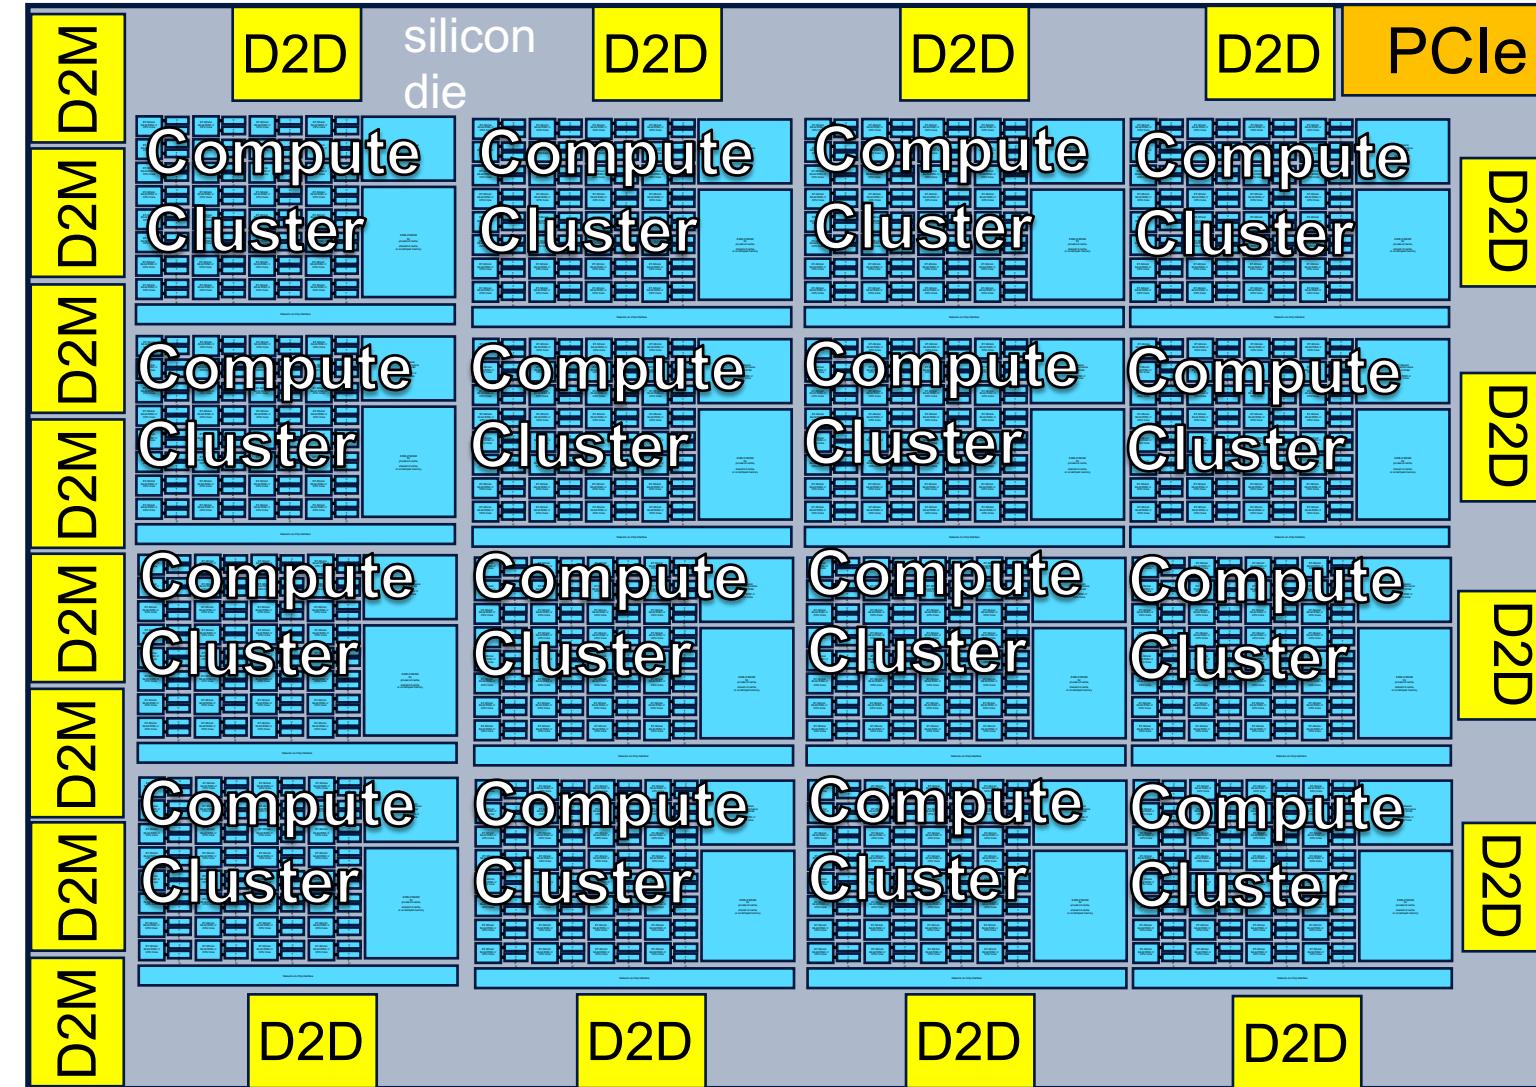
Higher speed/bandwidth for chiplet die to die interconnect

Expanded physical and virtual address space (57 bits)

PCIe Gen6 or Gen7 IO

Other chiplet support, eg optical IO, TBD, depending on customer interest

ET-SoC-3x 2nm Compute Fabric Chiplet: Twice the compute vs 4nm



ET-SoC-3x chiplet has:

- 16 Compute Clusters with
- 512 ET-Minion2 processors
- 16 ET-Maxion processors
- 128 MB on die SRAM

16 compute clusters, each with

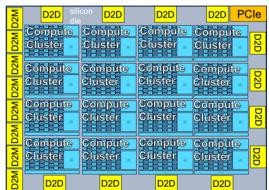
- One ET-Maxion “big core”
- 32 Minion “efficiency cores” in each cluster
- 8 MB of SRAM for L2, L3 and scratchpad
- Network on Chip to other Compute Clusters

Die to die (D2D) interfaces allow fabric to grow.

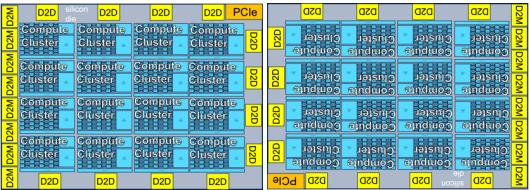
Die to Memory allows up to 2 TB/s bandwidth
direct to memory side cache on die

ET-SoC-3x Allowed Compute Fabric Configurations

ET-SoC-3x1



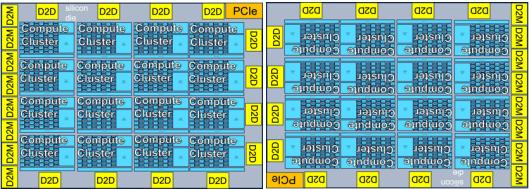
ET-SoC-3x2



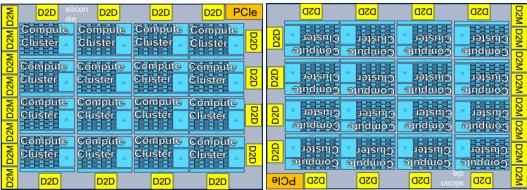
ET-x4



ET-SoC-3x6

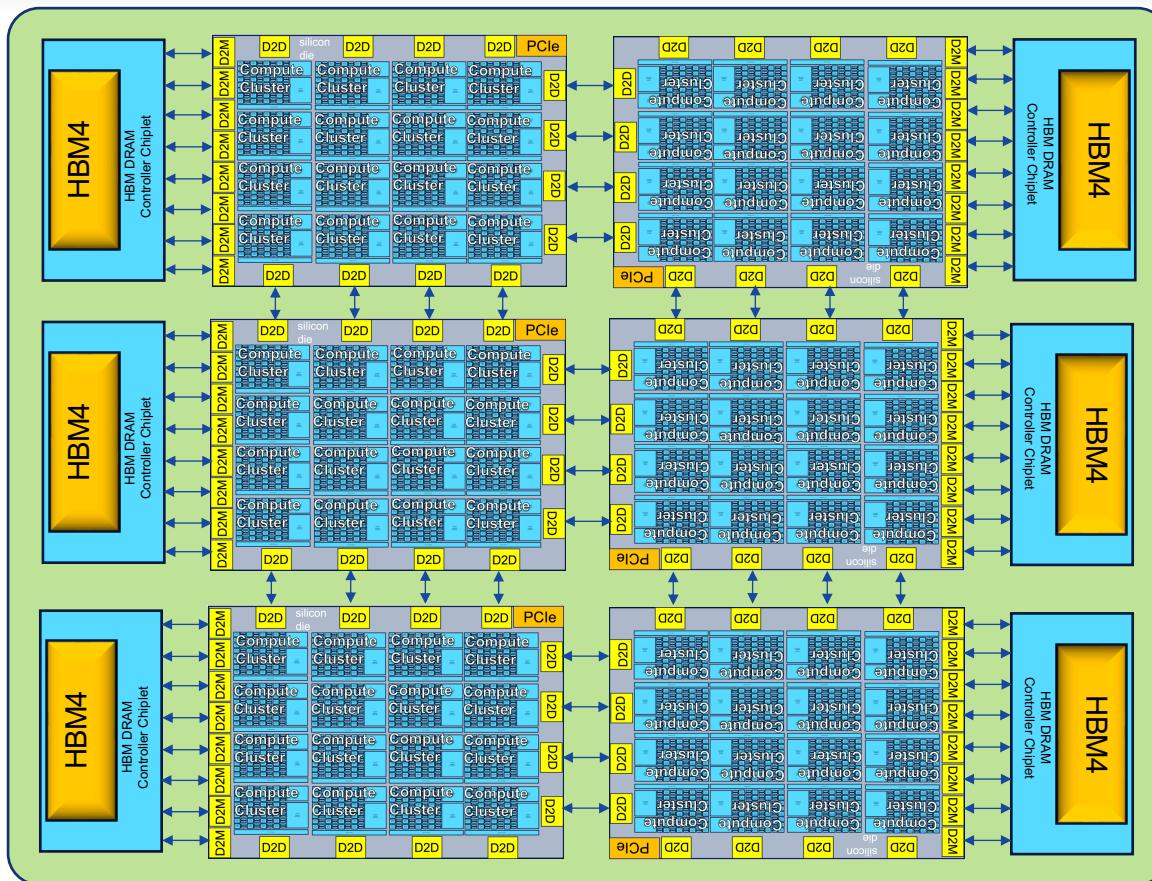


ET-SoC-3x8



	ET-SoC-3x1	ET-SoC-3x2	ET-SoC-3x4	ET-SoC-3x6	ET-SoC-3x8
# Chiplets in package	1	2	4	6	8
# of HBM4e supported	1	2	4	6	8
HBM4e Memory GB	64	128	256	384	512
Memory Bandwidth TB/s	2	4	8	12	16
PCIe Gen 6/7 Lanes	16	32	64	96	128
PCIe Bandwidth GB/s	128	256	384	512	1024
Peak TFlops	GHz->	2			
FP8 or Int8	524	1048	2096	3144	4192
FP16 / BF16	262	524	1048	1572	2096
FP32	65.6	131.2	262.4	394	524
FP64	32.8	65.6	131.2	196	262

2nm ET-SoC-3 (x6) package configuration example product



Six 2 GHz chiplets have 200 TFlops FP64 / 460 watts = 434 GFlops/watt

Six 2nm ET-SoC-3 Chiplets

- Max perf example @2 GHz
- 77 Watts per Compute Chiplet
- 460 Watts for 6 Compute Chiplets

Peak Performance:

- 200 TFlops FP64
- 400 TFlops FP32
- 1600 TFlops FP16/BF16
- 3200 TFlops FP8
- 3200 TOps Int8

Six HBM4e on module (2TB/s 64 GB)

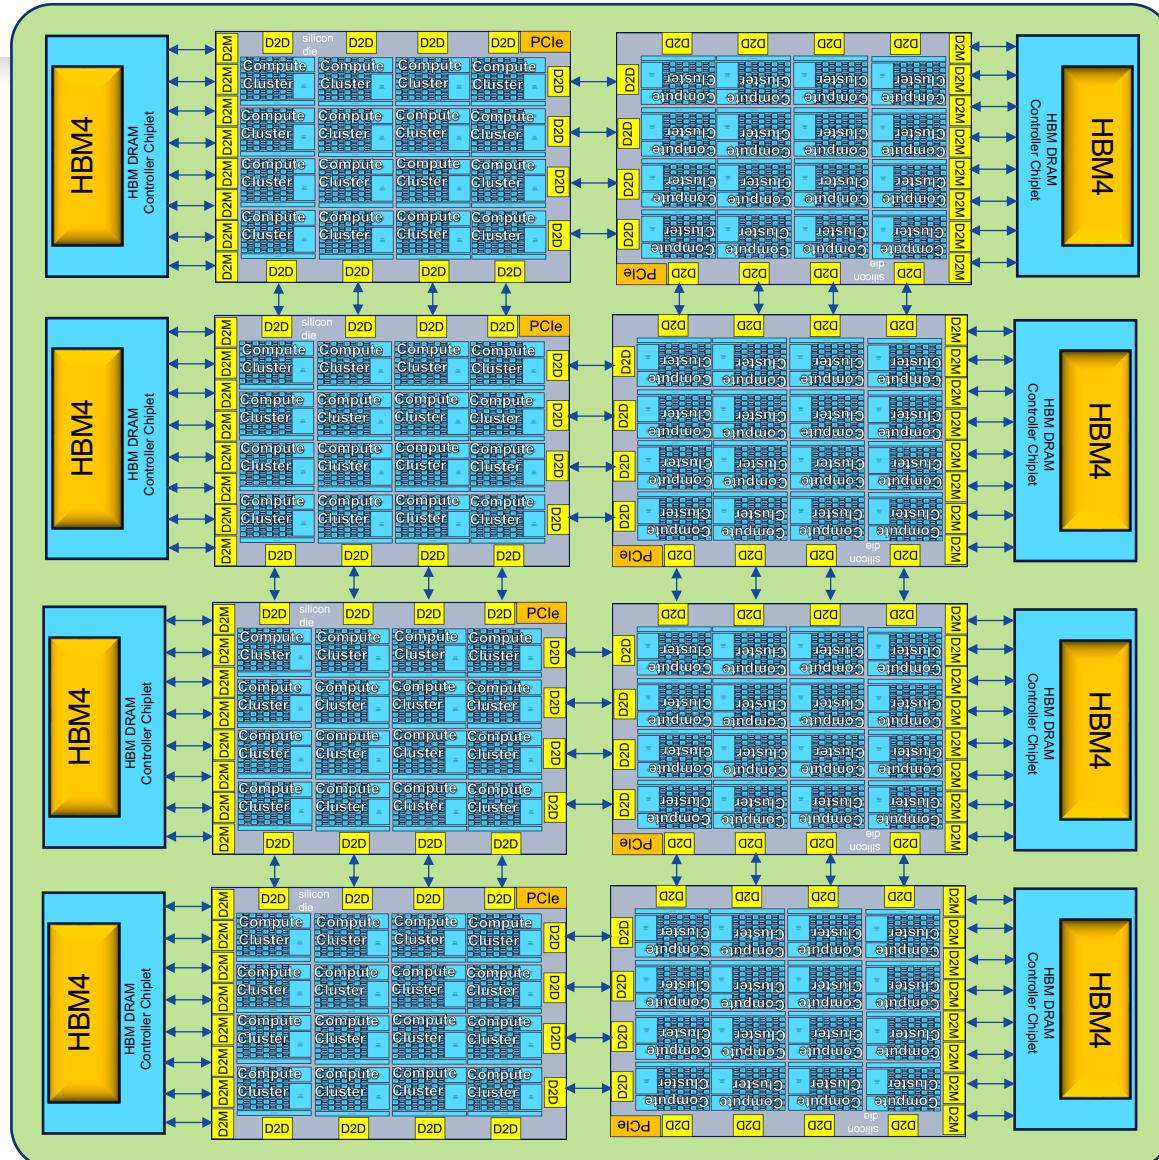
- 384 GB capacity
- 12 TB/sec bandwidth

96 lanes PCIe Gen6/7

Expansion DRAM off module options

Better FP64 performance than new NVIDIA Blackwell 1KW B200 GPU which is only 40 TFlops FP64

2nm ET-SoC-3 (x8) package configuration example product



Eight 1.5 GHz chiplets have 200 TFlops FP64 / 350 watts = 571 GFlops/watt

Eight 2nm ET-SoC-3 Chiplets

- Efficient example @1.5 GHz
- 44 Watts per Compute Chiplet
- 350 Watts for compute chiplets

Peak Performance:

- 200 TFlops FP64
- 400 TFlops FP32
- 1600 TFlops FP16/BF16
- 3200 TFlops FP8
- 3200 TOps Int8

Eight HBM4e on module (2TB/s 64 GB)

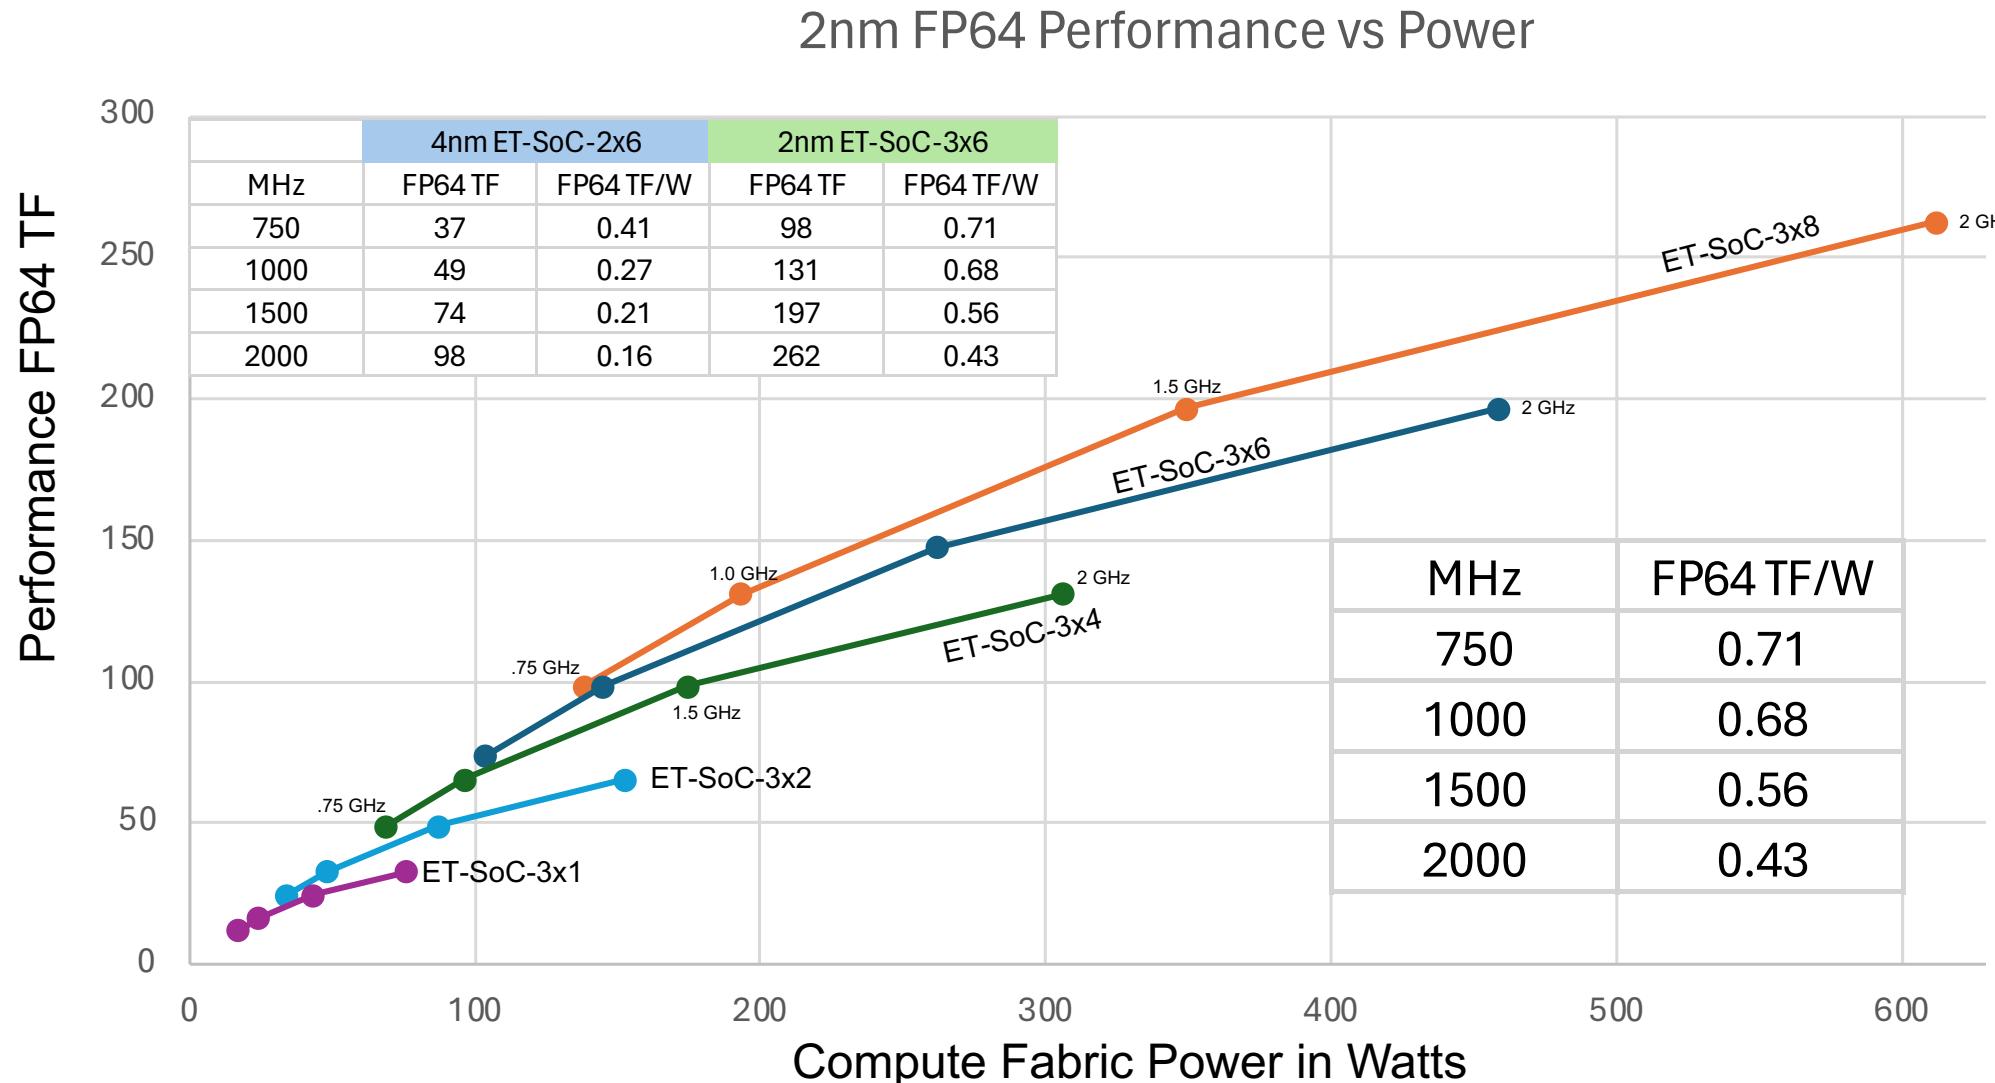
- 512 GB capacity
- 16 TB/sec bandwidth

128 lanes PCIe Gen6/7

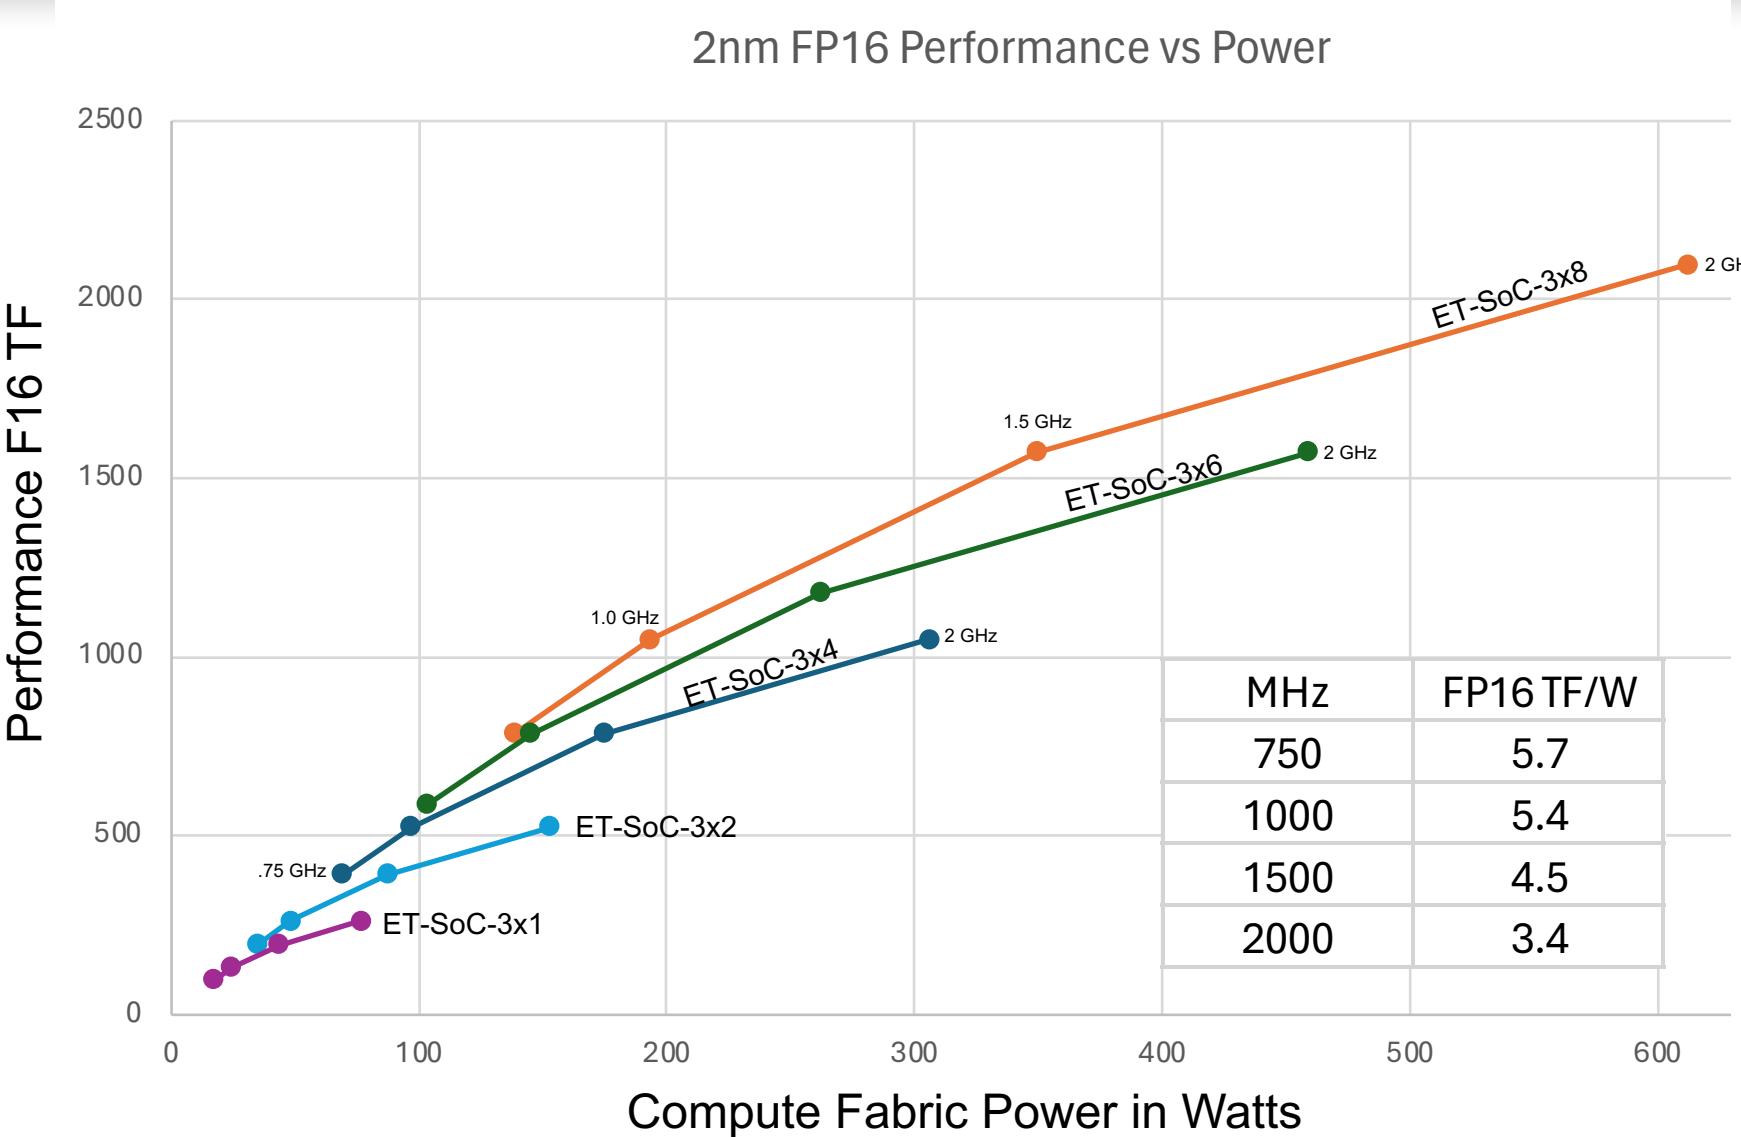
Expansion DRAM off module options

Better FP64 performance than new NVIDIA Blackwell 1KW B200 GPU which is only 40 TFlops FP64

FP64 Performance vs Power for different 2nm ET-SoC-3 Packages



FP16 Performance vs Power for different 2nm ET-SoC-3 Packages

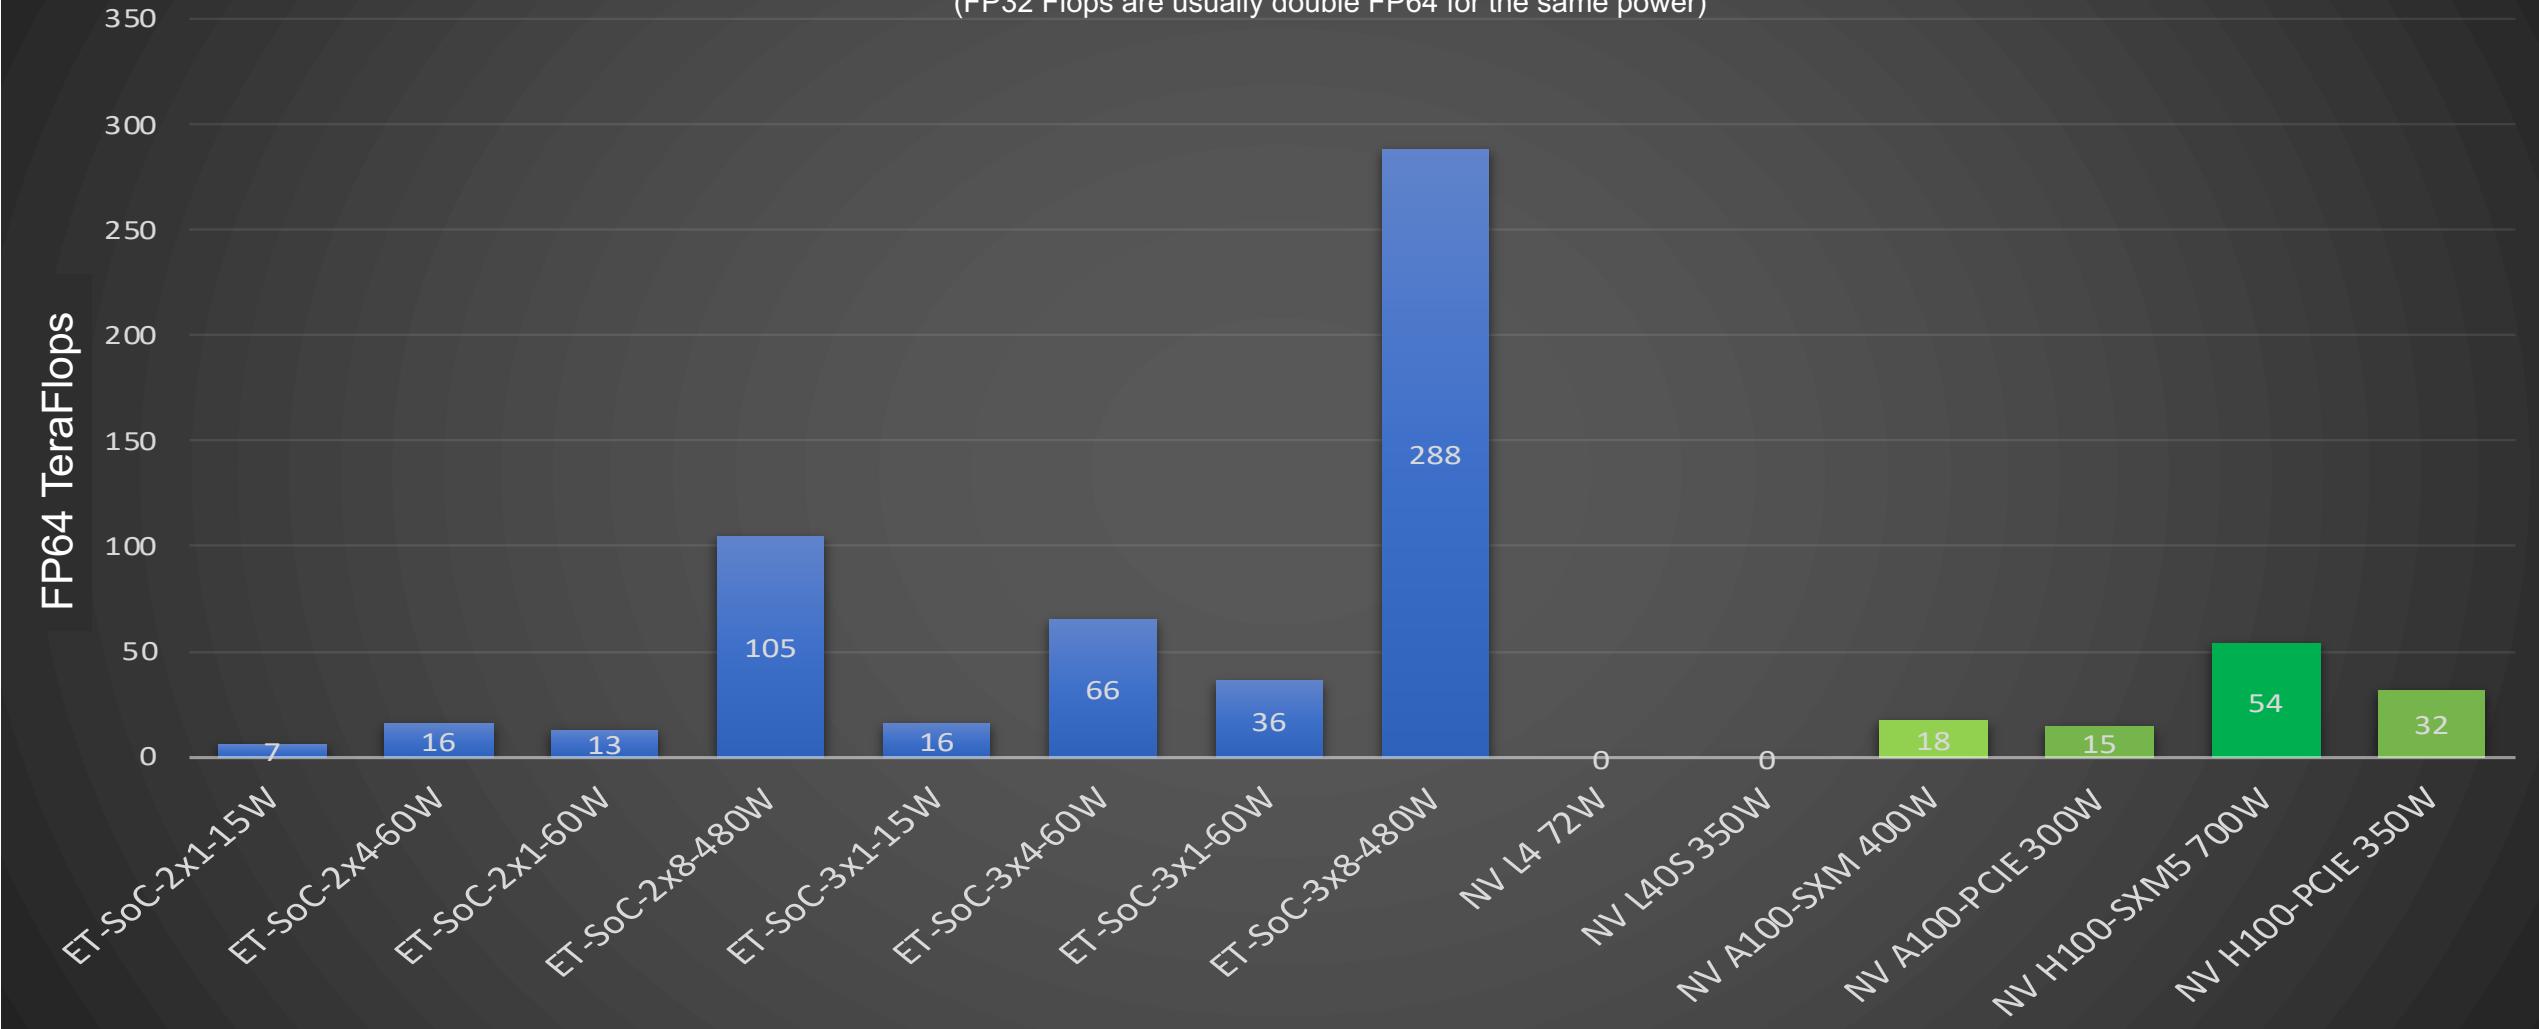


Performance, Power and Performance per Watt Comparisons

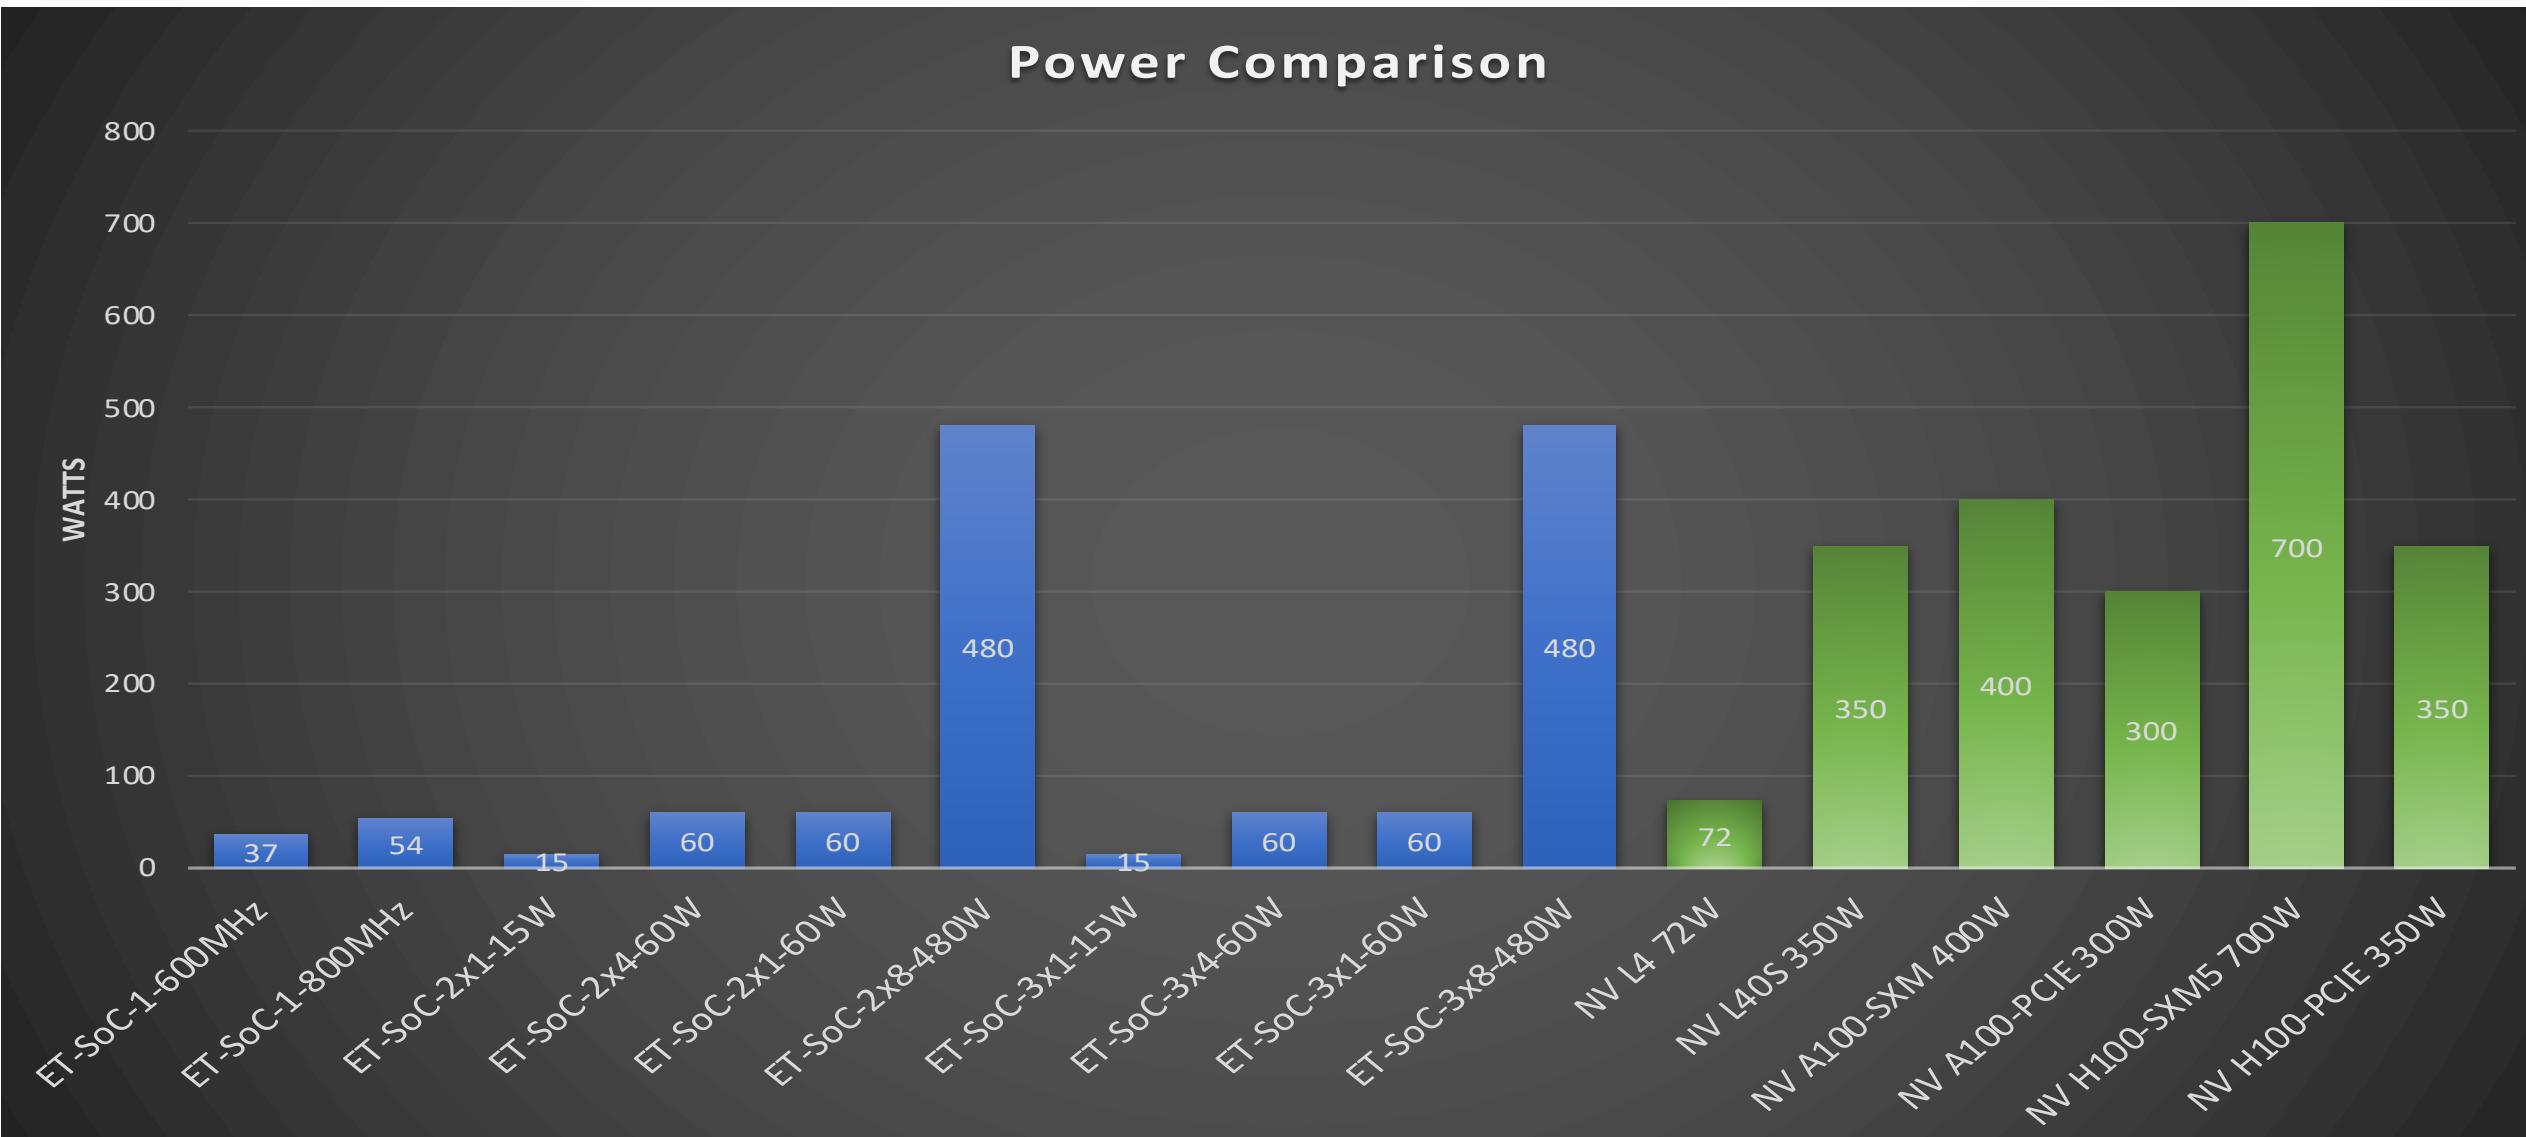
Note – to provide more apples to apples comparisons, the following graphs compare computation rates assuming no sparsity and sustainable base clock rates that do not exceed TDP.

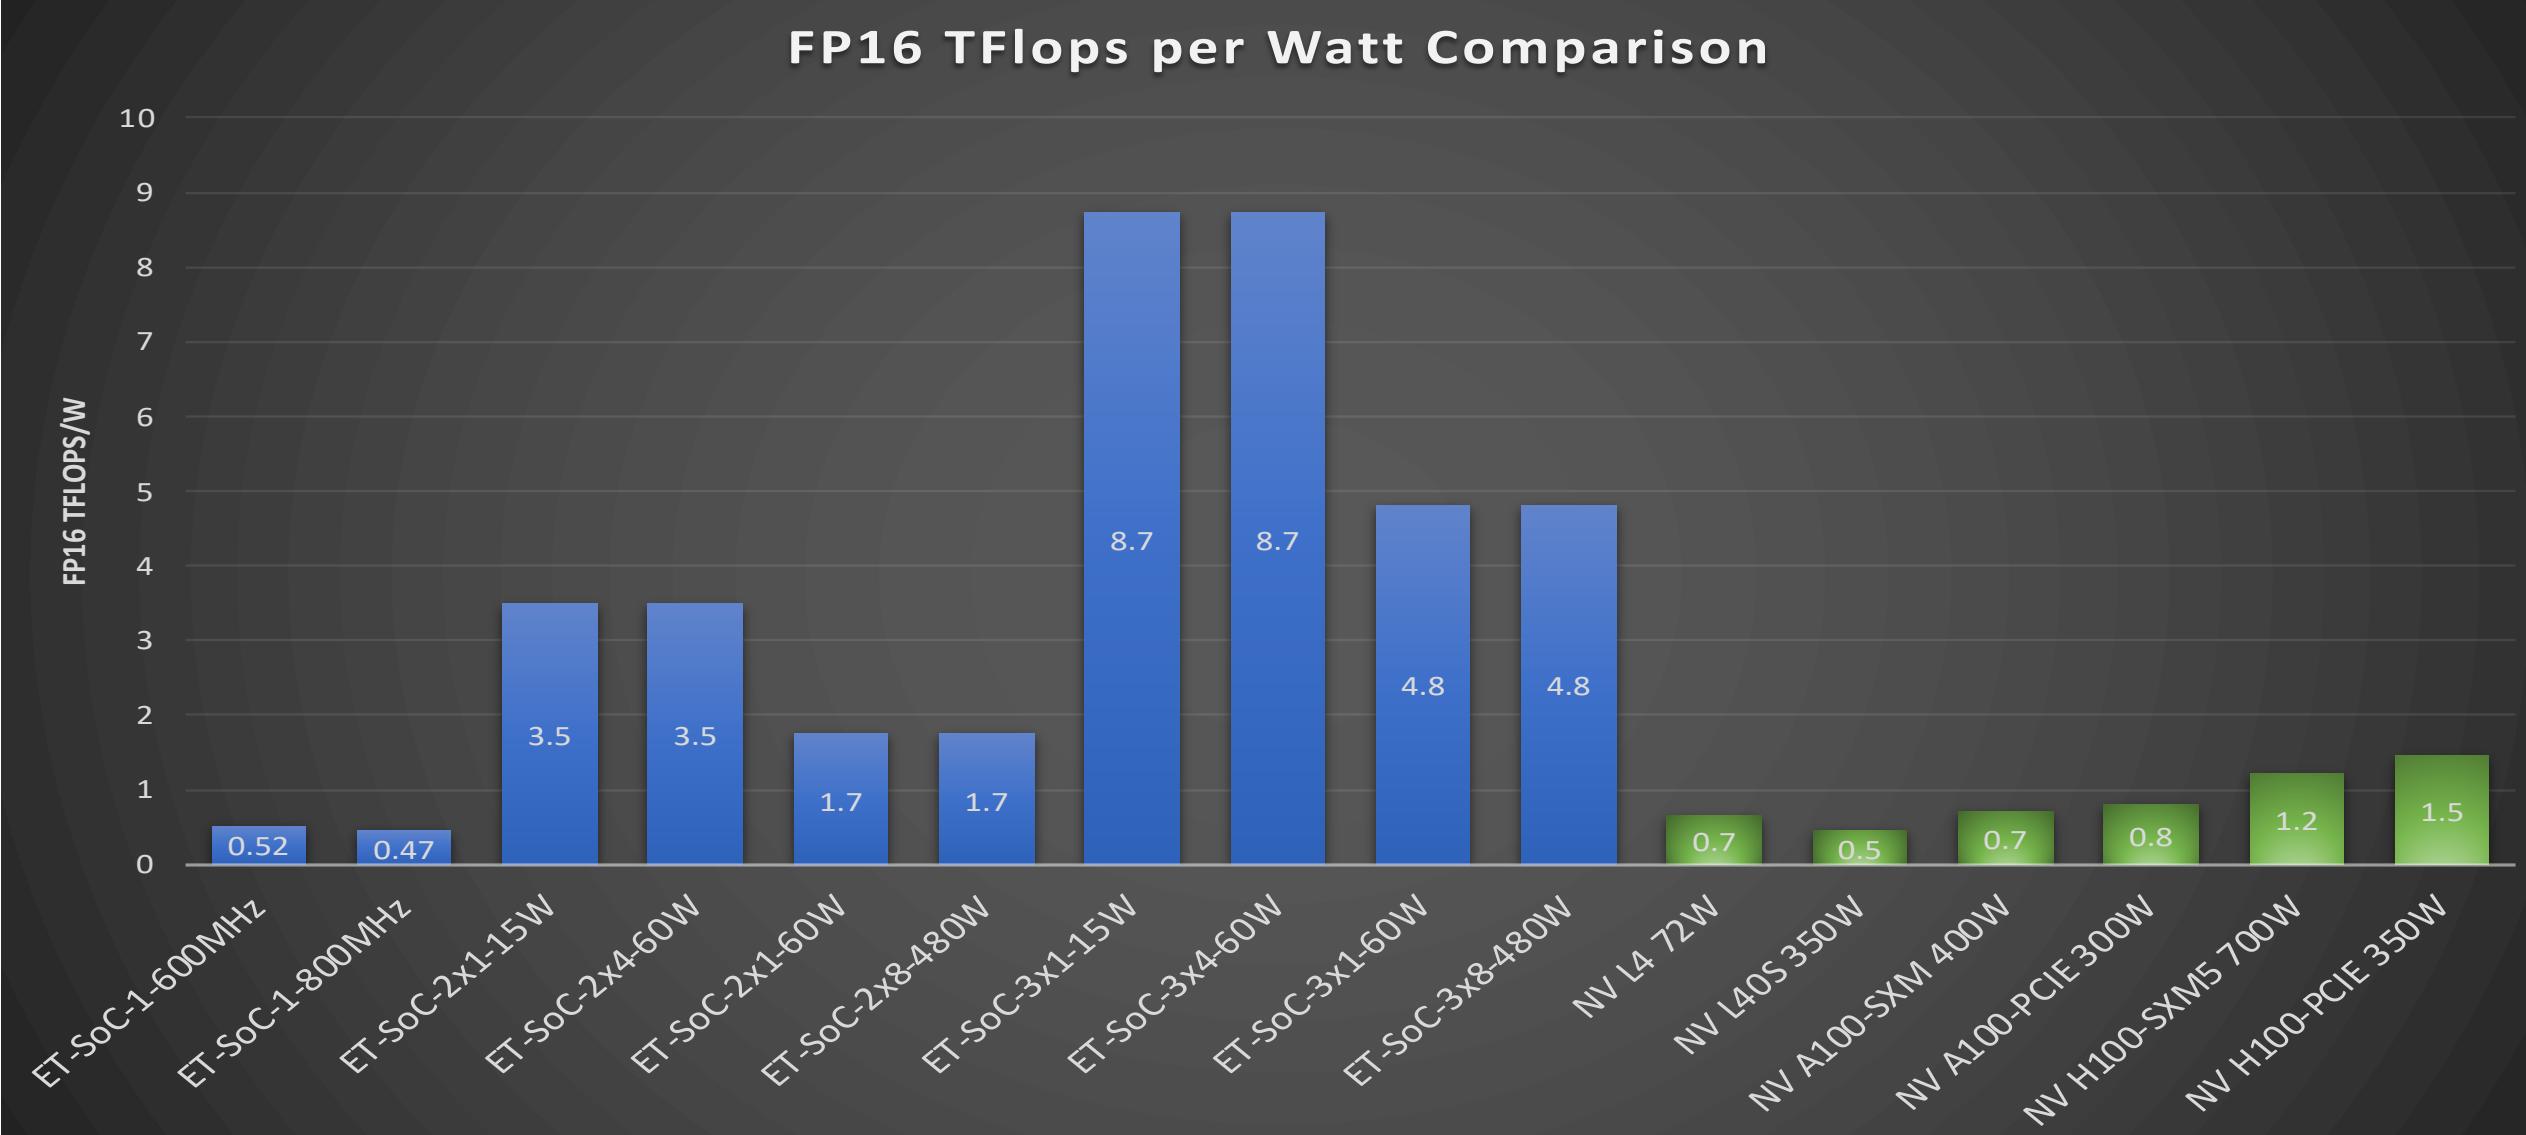
FP64 TFlops Comparison

(FP32 Flops are usually double FP64 for the same power)

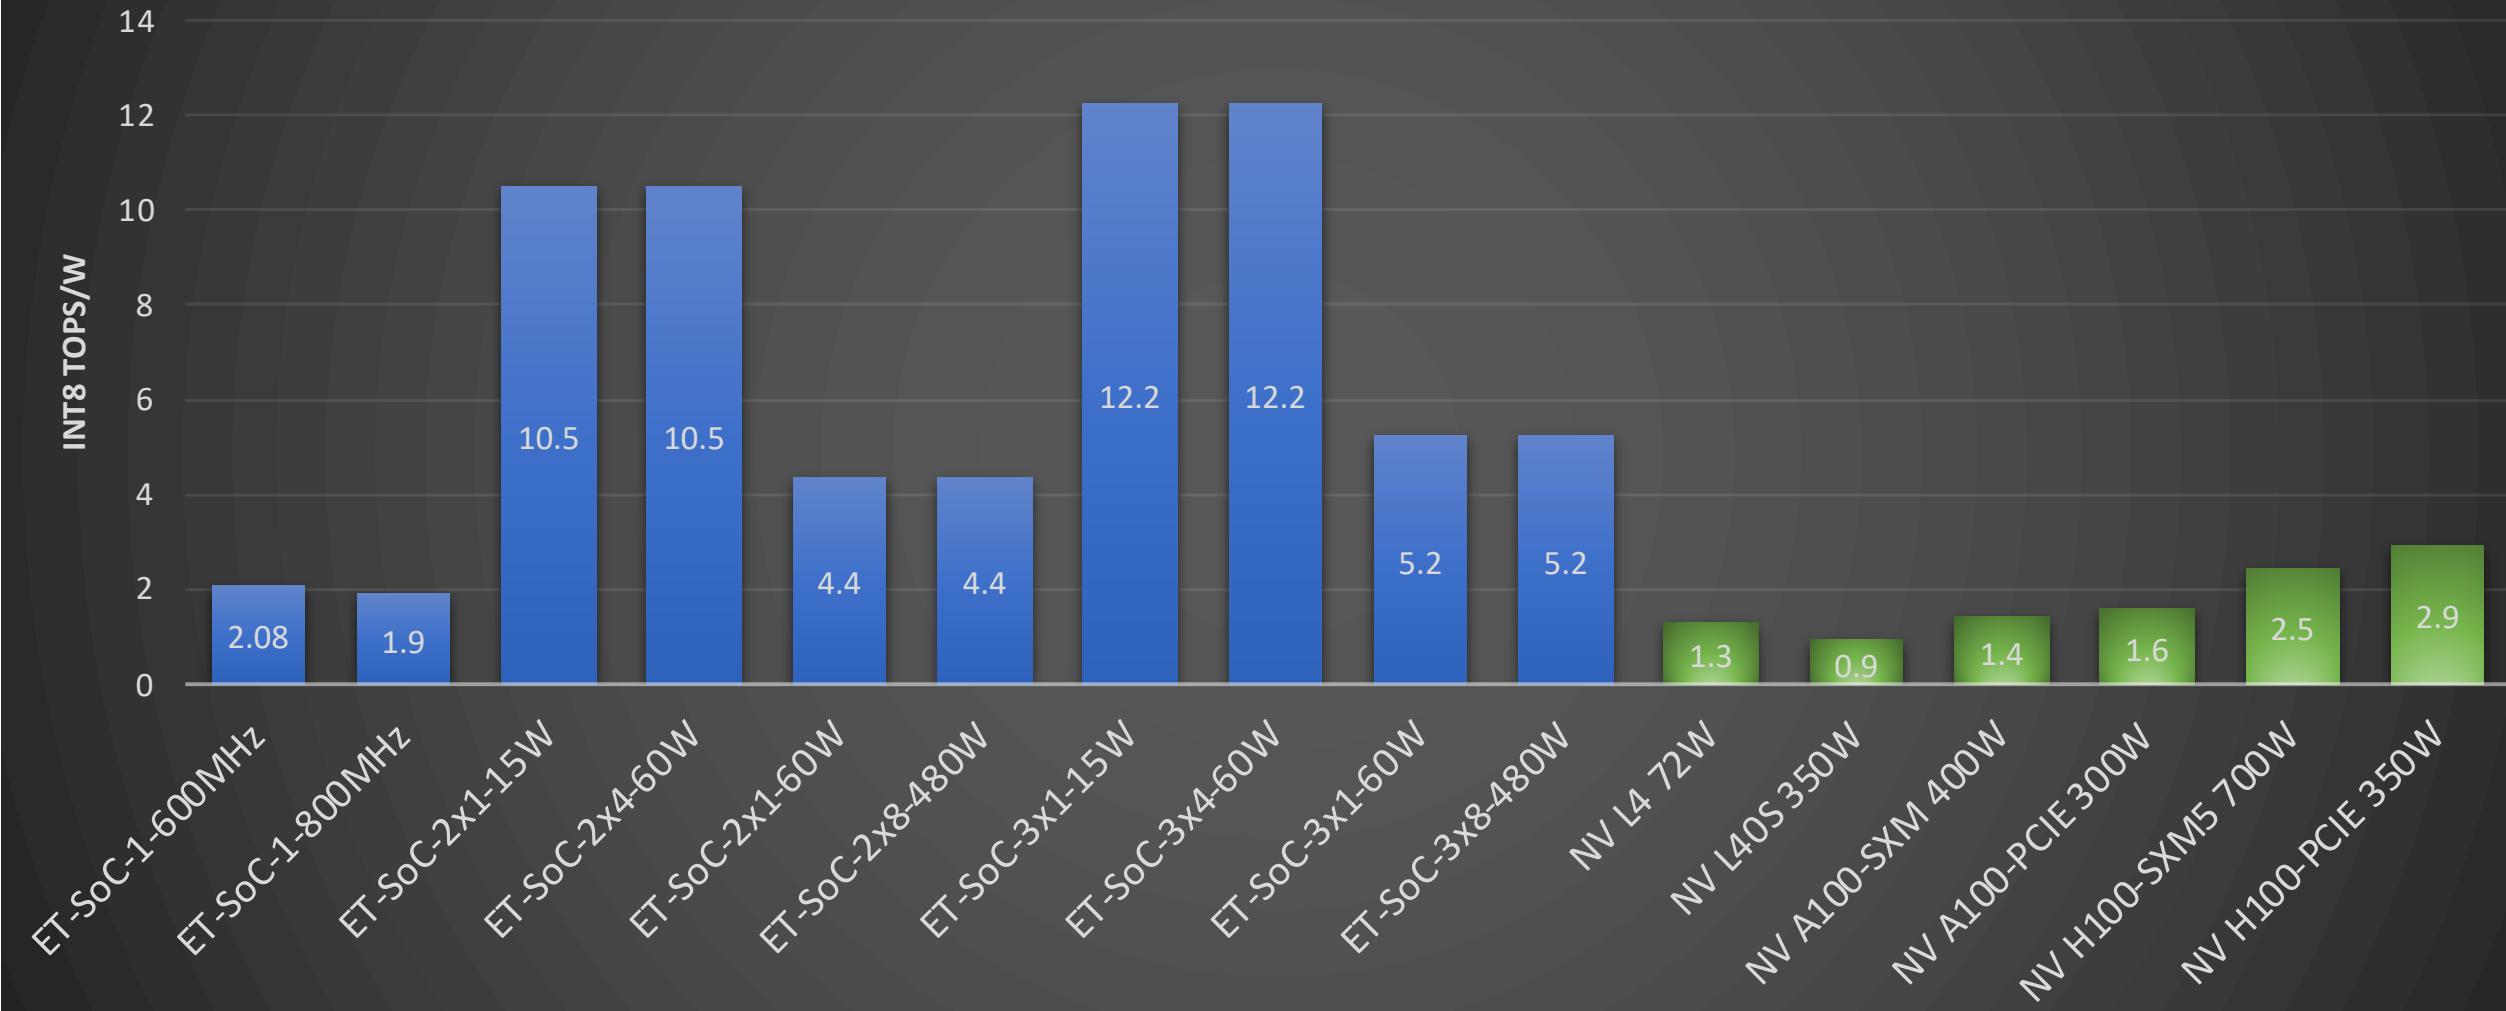


Thermal Design Power (TDP) comparisons

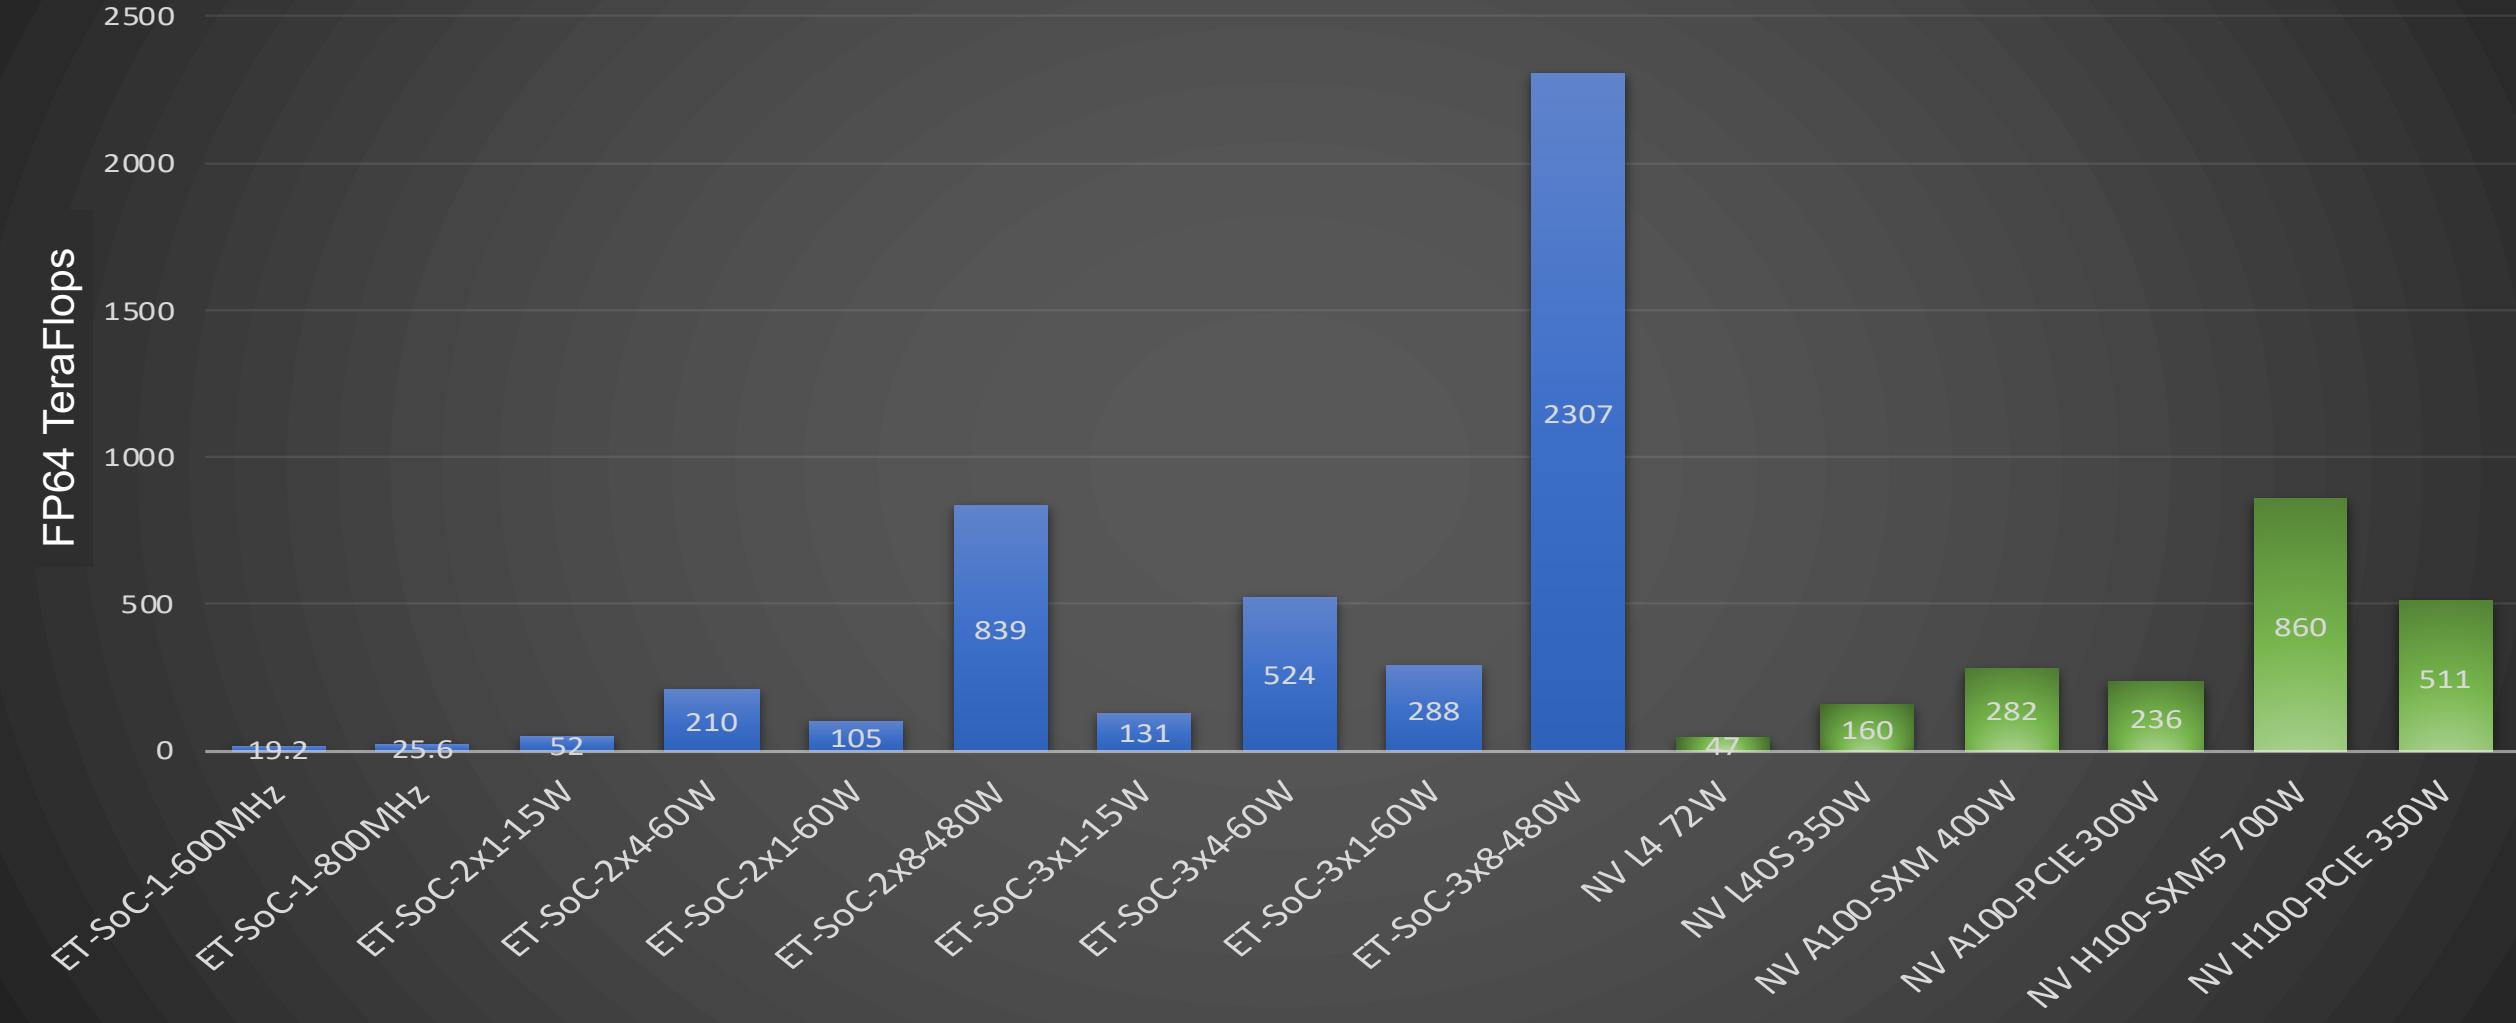




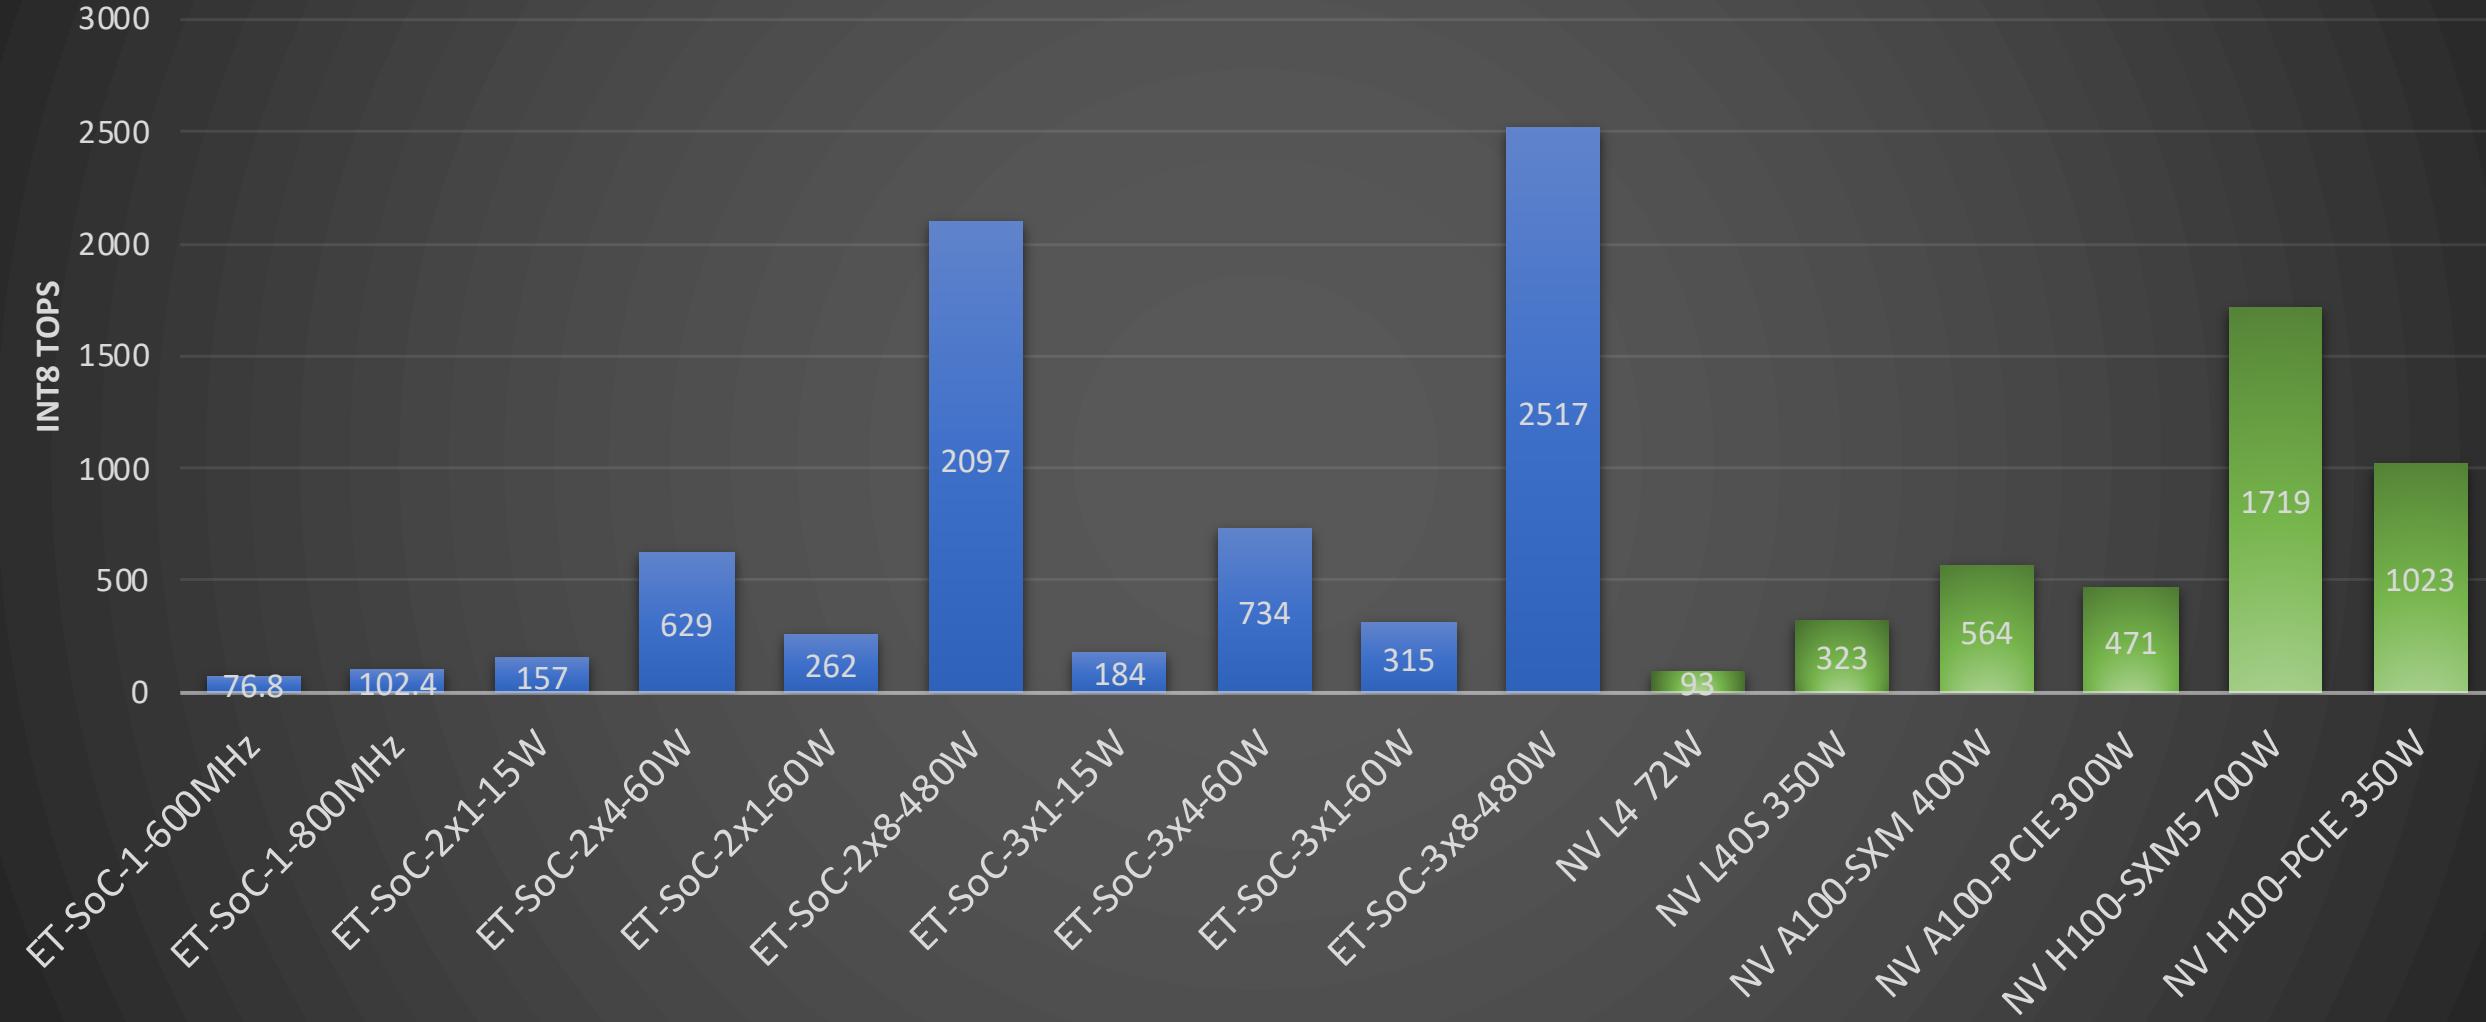
INT8 TOPs per Watt Comparison



FP16 TFlops Comparison



INT8 TOps Comparison



Have worked with many instruction sets: CISC, VLIW, x86, ARM, and RISC-V

RISC-V is the best of the alternatives

- RISC-V is better technically
- RISC-V is better legally ("Free and Open", i.e. no license needed)
- RISC-V is still maturing and has a lot of room to grow

Join us in making RISC-V successful in the HPC space!



www.esperanto.ai