

Simulating Hybrid Analog + RISC-V Systems for HPC Applications

SuperComputing '25 International Workshop on HPC for RISC-V

November 17, 2025

Cameron Durbin¹, Jacob Flores², Thomas Weatherly³, Ben Feinberg²

¹ University of Oregon, ² Sandia National Laboratories, ³ Georgia Institute of Technology



Why Consider Analog?

Scaling Reality

Dennard scaling is over – more performance costs power. Most of that energy is burned moving data.

Analog Promise

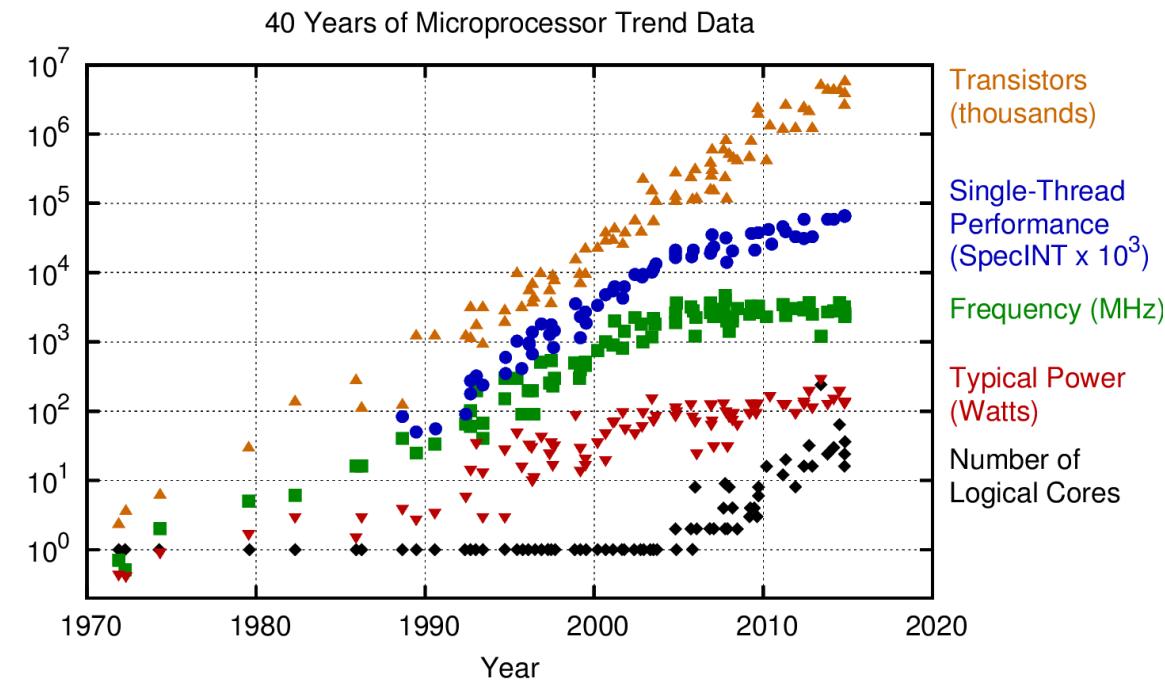
Crossbar MVMs can deliver 10-100x energy/throughput gains *per MVM*.

Integration Challenge

Crossbar MVMs only pays off when it's coupled with a programmable digital system.

Core Question:

How do we make analog acceleration usable and efficient inside a full RISC-V system?



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2015 by K. Rupp

Why RISC-V?

Open and Extensible ISA

Allows us to add our analog MVM instructions without proprietary toolchains.

Mature Software Stack

LLVM and OpenMP support fast, flexible HPC development on RISC-V.

Tight Integration Path

The RoCC interface enables low-overhead attachment of custom accelerators to each core.

Full-System Simulation

SST models RISC-V with cycle granularity, enabling detailed studies of analog-digital interactions.

Research Agility

RISC-V's transparent ecosystem and reproducible infrastructure make it ideal for exploring hybrid architectures.

Bringing It Together

RISC-V offers the flexibility and maturity needed to explore how analog acceleration behaves within a fully programmable system.

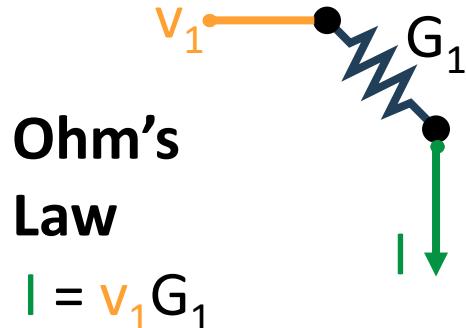
Analog MVMs

How it works

- Program the conductances of \mathbf{A}
- Apply voltages of vector \mathbf{x} to produce vector \mathbf{y}

Program-once-ish

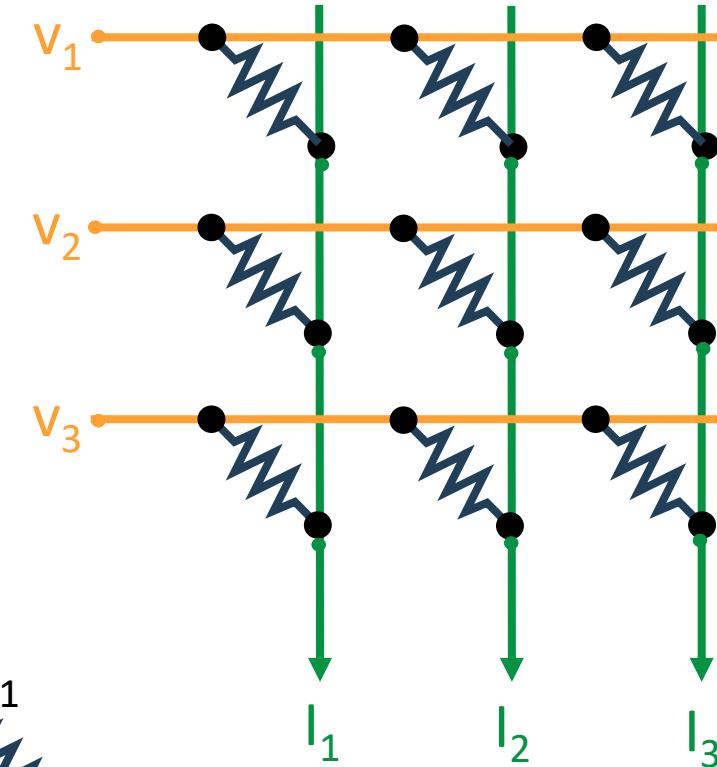
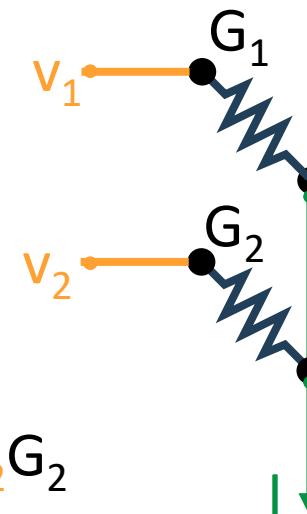
Writing resistive states is slow and consumes substantial energy.



Kirchhoff's Law

$$I = V_1 G_1 + V_2 G_2$$

A circuit diagram showing two voltage sources V_1 and V_2 connected in series. They are connected to two resistors G_1 and G_2 , respectively. A green arrow labeled I indicates the direction of current flow through the resistors.



Matrix-Vector Multiplication

$$\mathbf{I} = \mathbf{V}\mathbf{G}$$

System Architecture

Analog Arrays

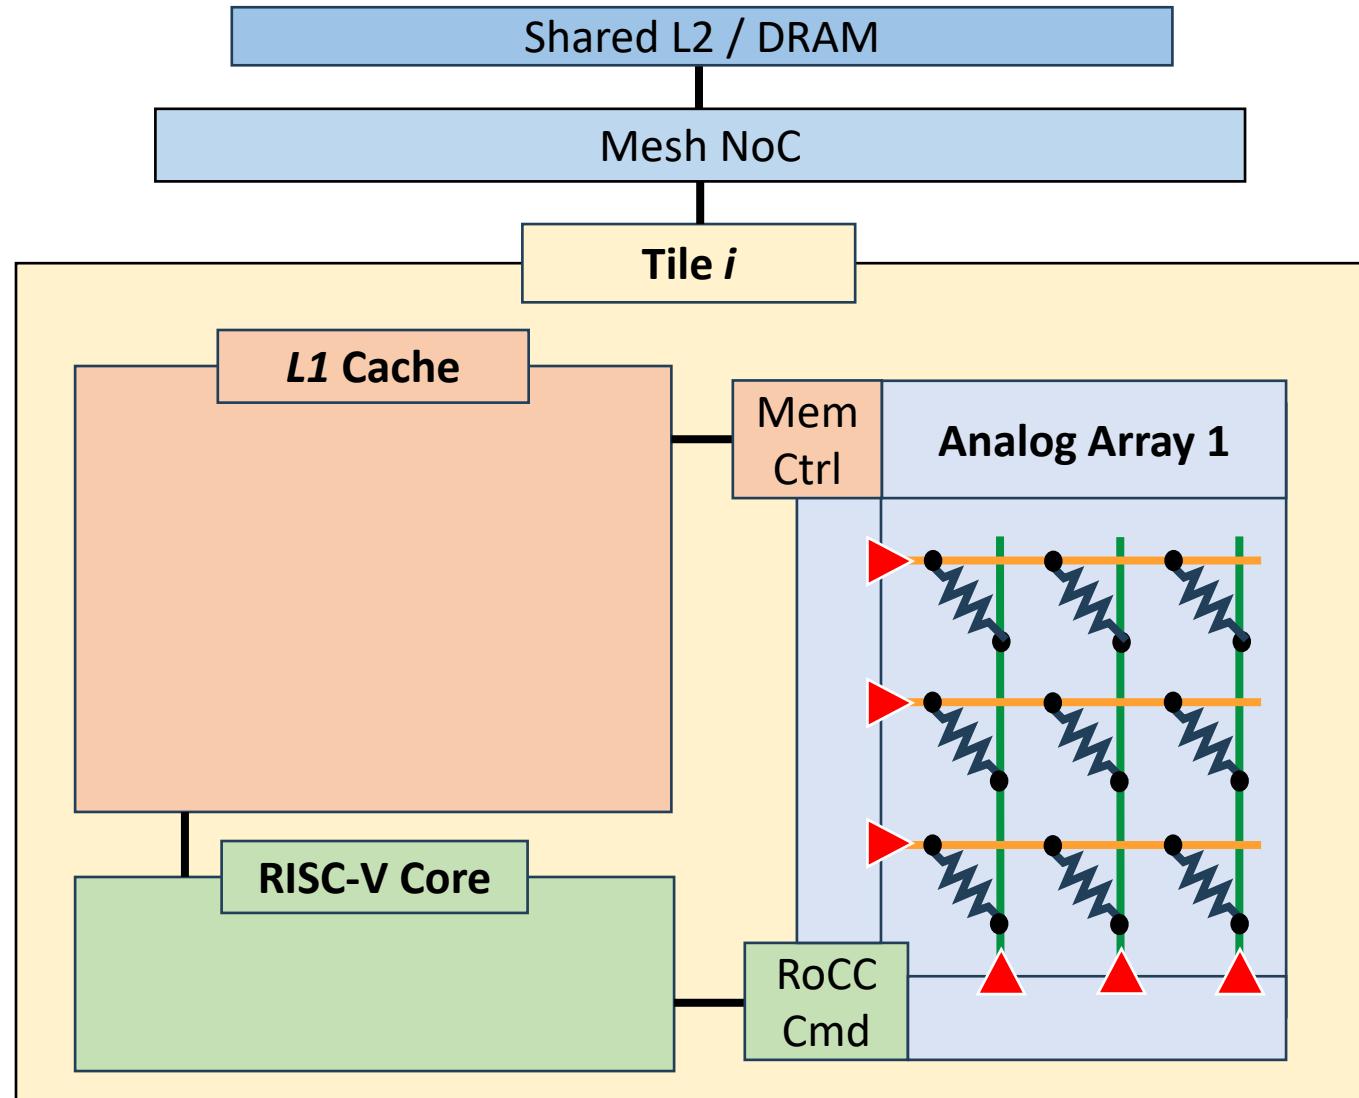
Architectural visible input/output buffers, DMA, RoCC Interface

Tile

Contains RISC-V Core, n analog arrays, L1 Cache

NoC

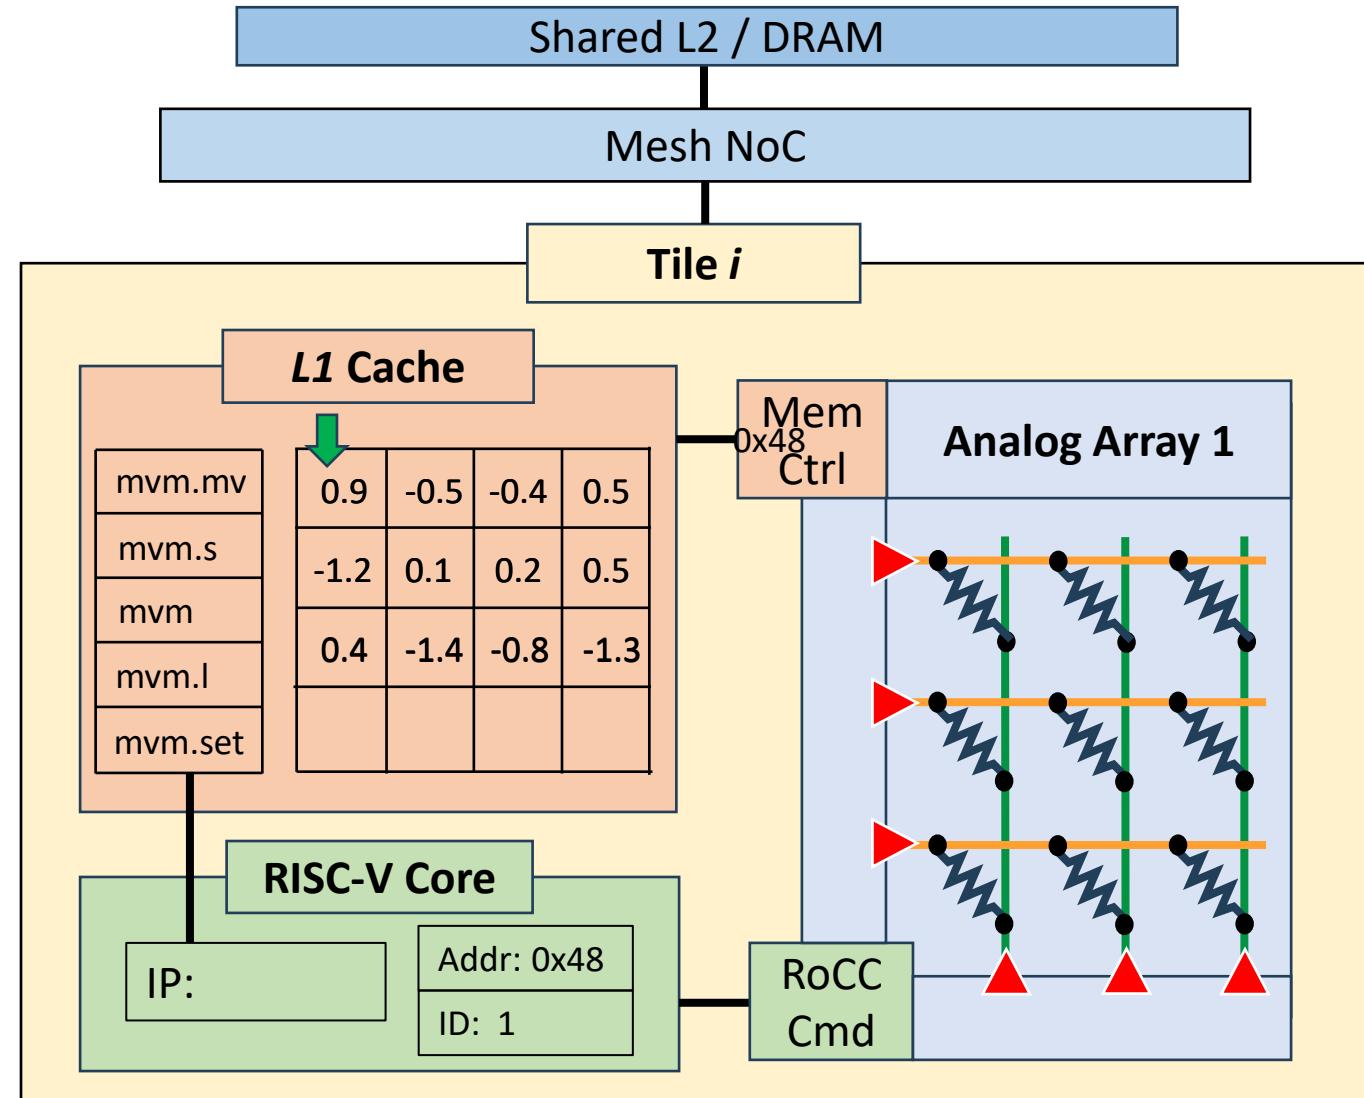
Connects m tiles in a mesh, shared L2 cache, directory-based MESI coherence



Minimal Analog ISA

mvm.set %addr, id

- Program the conductances of analog array *id* using matrix A found at *%addr*

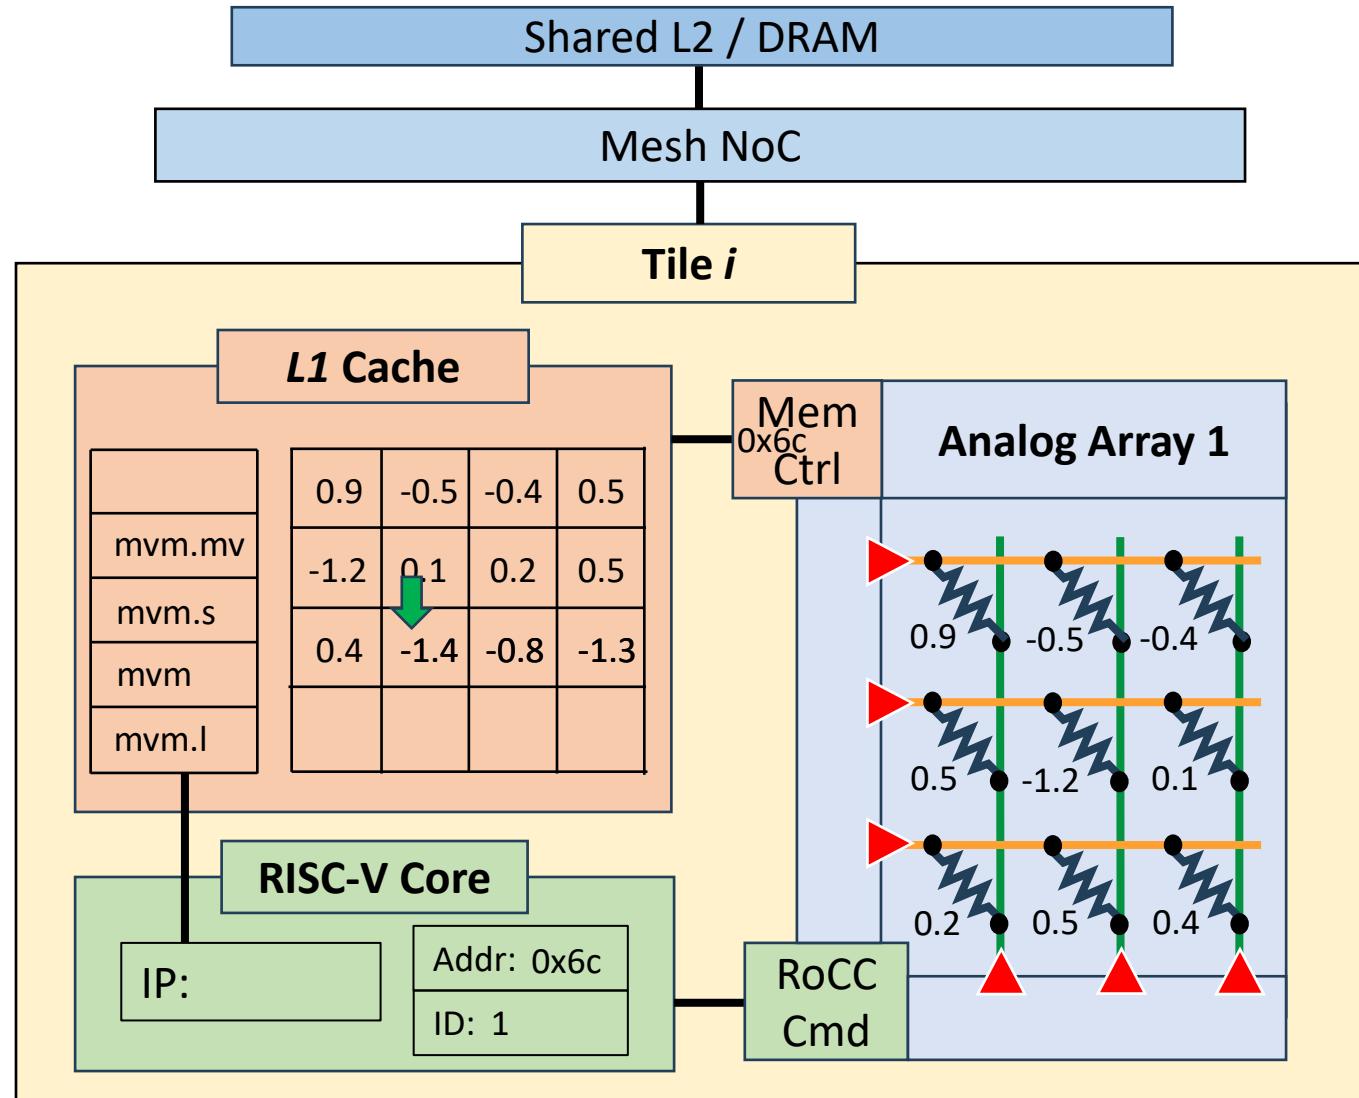


Minimal Analog ISA

mvm.set %addr, id

mvm.l %addr, id

- Load an input vector into analog array **id** using vector x found at **%addr**



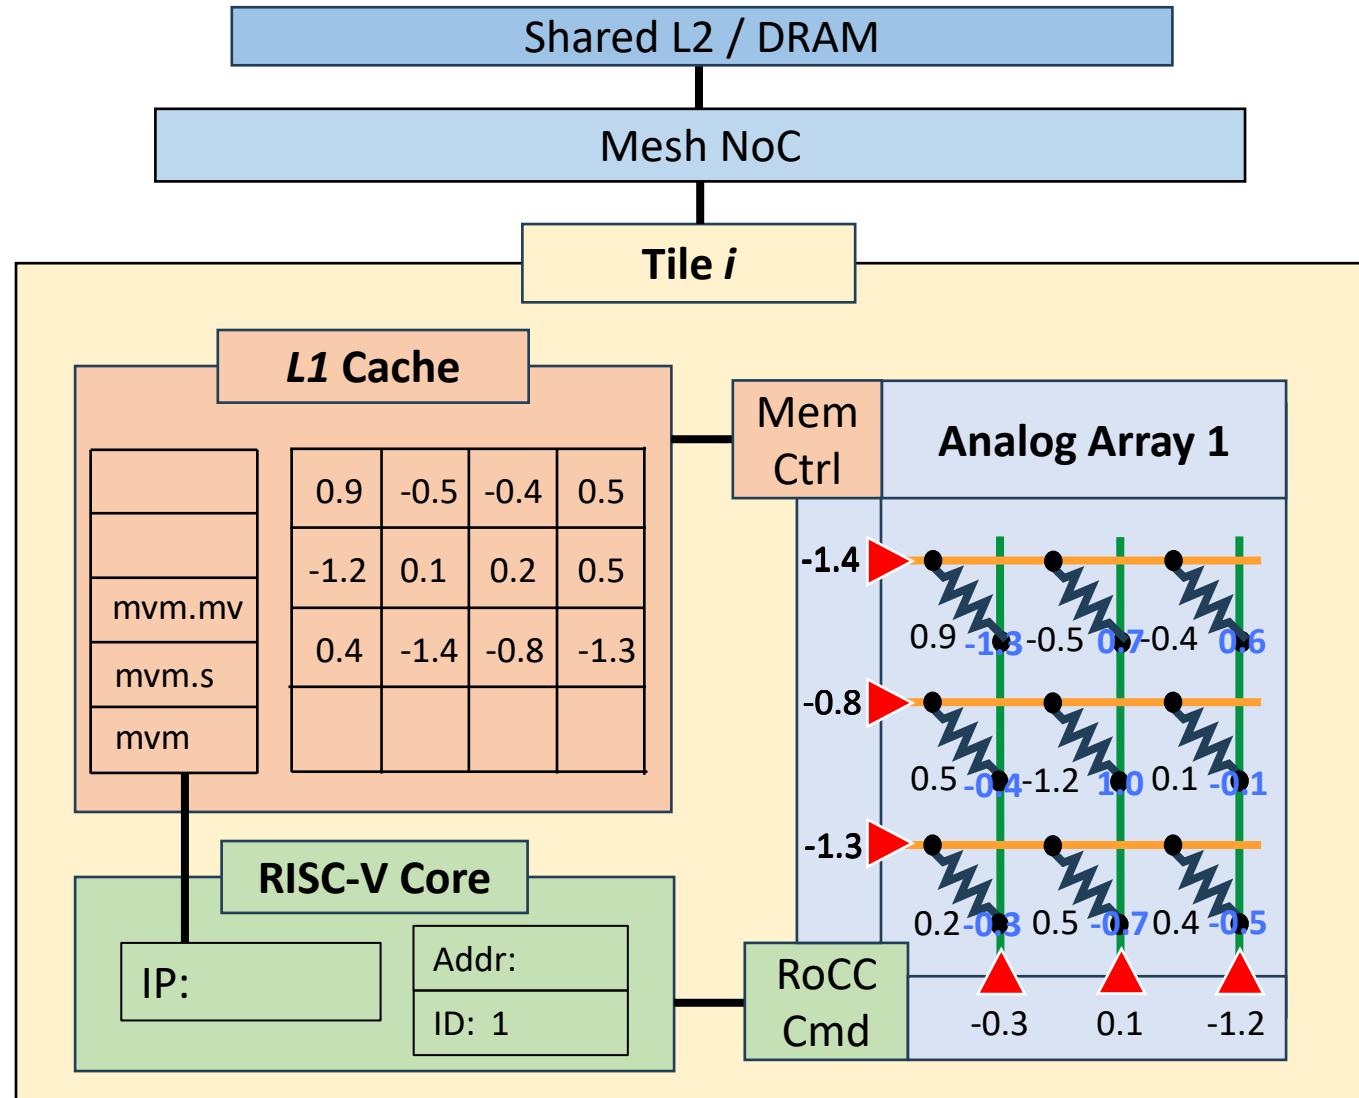
Minimal Analog ISA

mvm.set %addr, id

mvm.I %addr, id

mvm id

- Perform an MVM using the crossbar and the input buffer on array **id**



Minimal Analog ISA

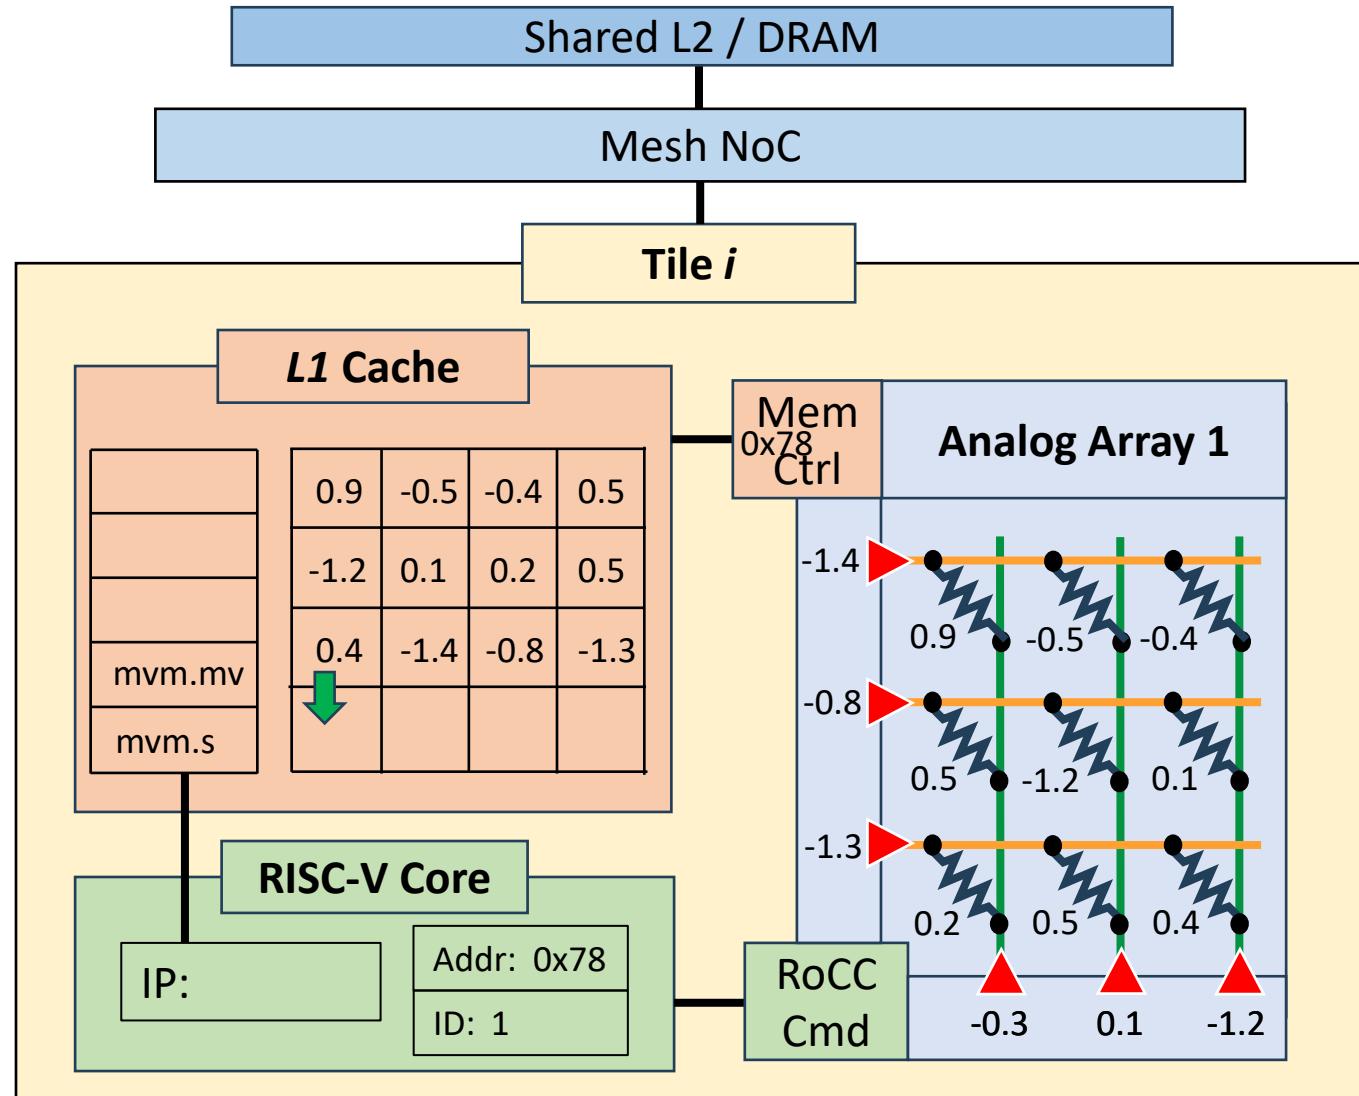
mvm.set %addr, id

mvm.l %addr, id

mvm id

mvm.s %addr, id

- Store the contents of the output buffer of array *id* at *%addr*



Minimal Analog ISA

mvm.set %addr, id

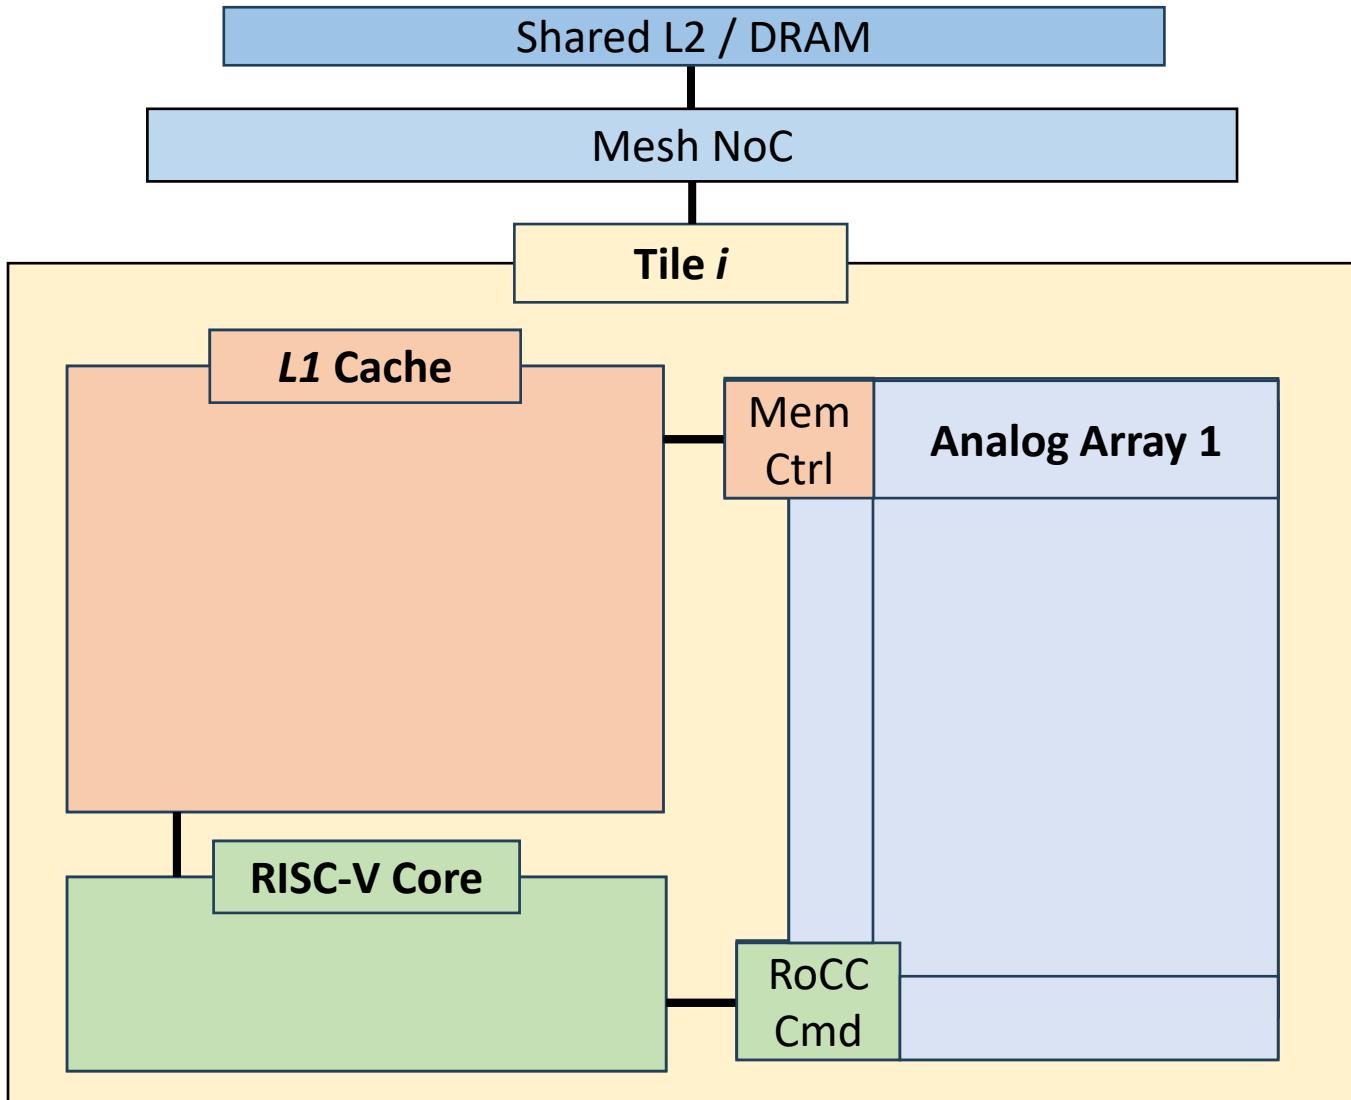
mvm.l %addr, id

mvm id

mvm.s %addr, id

mvm.mv src_id, dst_id

- Move the contents of the array *src_id*'s output buffer to array *dst_id*'s input buffer

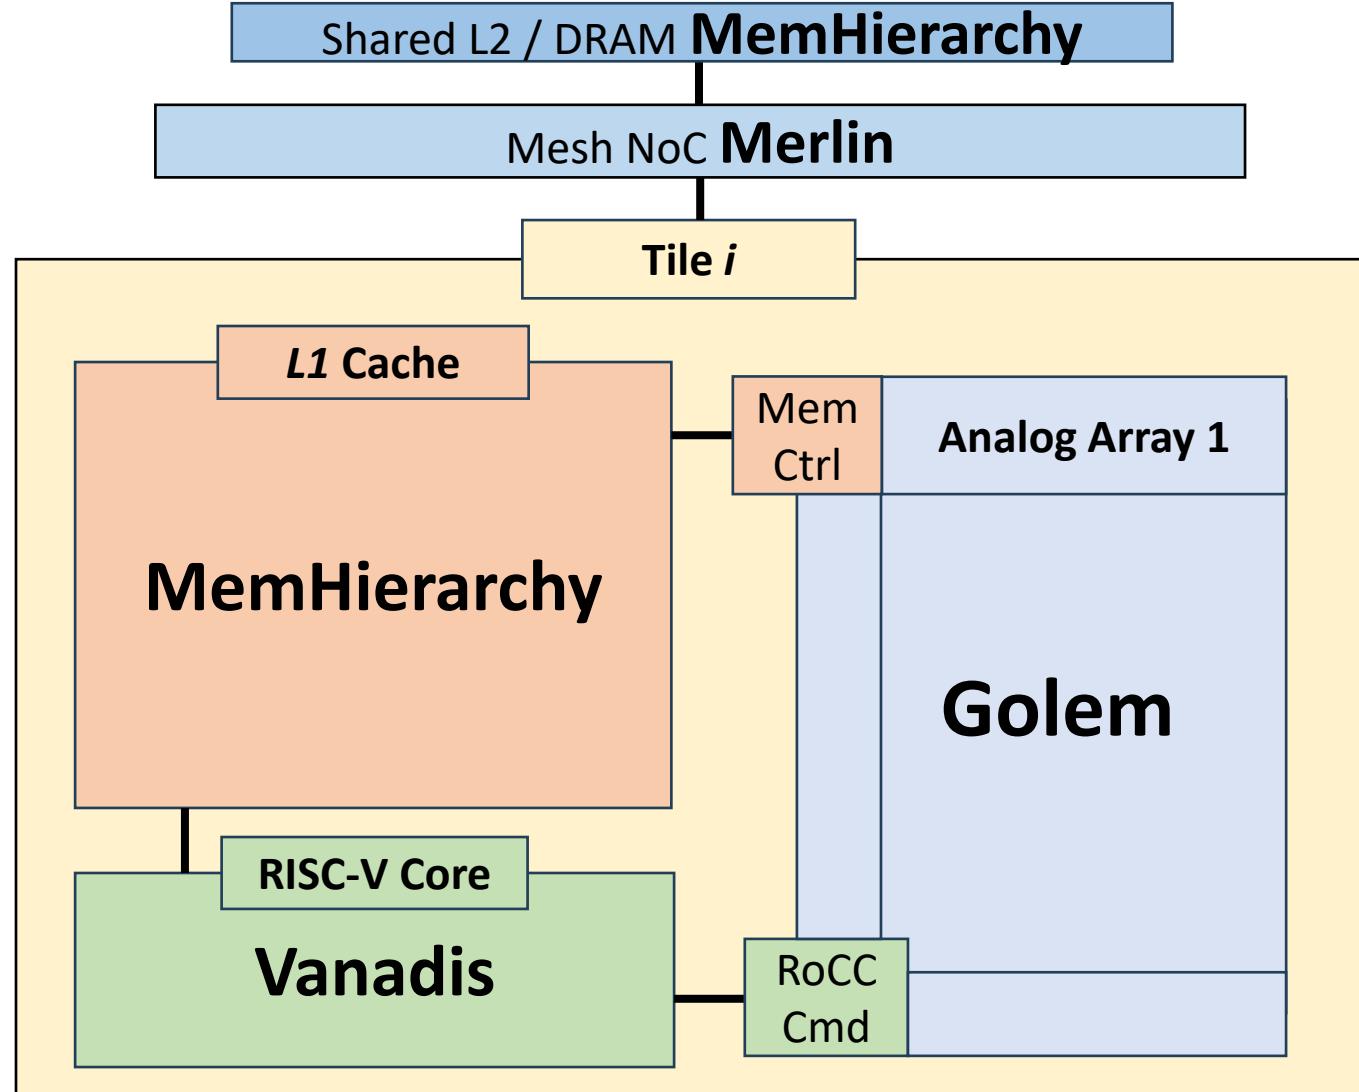


Simulation Stack

Structural Simulation Toolkit (SST)

SST simulates our infrastructure using the following components:

- MemHierarchy
- Vanadis
- Golem (Optional: CrossSim)
- Merlin



Software Stack

```
int mvm_set(int A, int t) {
    int s;
    asm volatile (
        "mvm.set %0, %1, %2"
        : "=r"(s)
        : "r"(A), "r"(t)
        : "memory"
    );
    return s;
}
```

```
int mvm_load(int x, int t) {
    int s;
    asm volatile (
        "mvm.l %0, %1, %2"
        : "=r"(s)
        : "r"(x), "r"(t)
        : "memory"
    );
    return s;
}
```

```
int mvm_exec(int t) {
    int s;
    asm volatile (
        "mvm %0, %1, x0"
        : "=r"(s)
        : "r"(t)
    );
    return s;
}
```

```
int mvm_store(int y, int t) {
    int s;
    asm volatile (
        "mvm.s %0, %1, %2"
        : "=r"(s)
        : "r"(y), "r"(t)
        : "memory"
    );
    return s;
}
```

AnalogM

```
#pragma
for (in
    int
    int
    int
    float
    ana
    ana
    } )
```

```
ma
for(in
    int
    int
    float
    float
    ana
    ) ;
```

);

Standard BLAS SGEMV

```
void sgemv(
```

```
    char trans, _____ → char trans,
    int m, int n, _____ → int m, int n,
    float alpha, _____ → float alpha,
    const float *A, _____ → const float *A,
    int lda, _____ → int lda,
    } AnalogMatrix;
```

ma

for

```
    const float *x, _____ → const float *x,
    int incx, _____ → int incx,
    float beta, _____ → float beta,
    float *y, _____ → float *y,
    int incy _____ → int incy,
    );
```

AnalogBLAS SGEMV

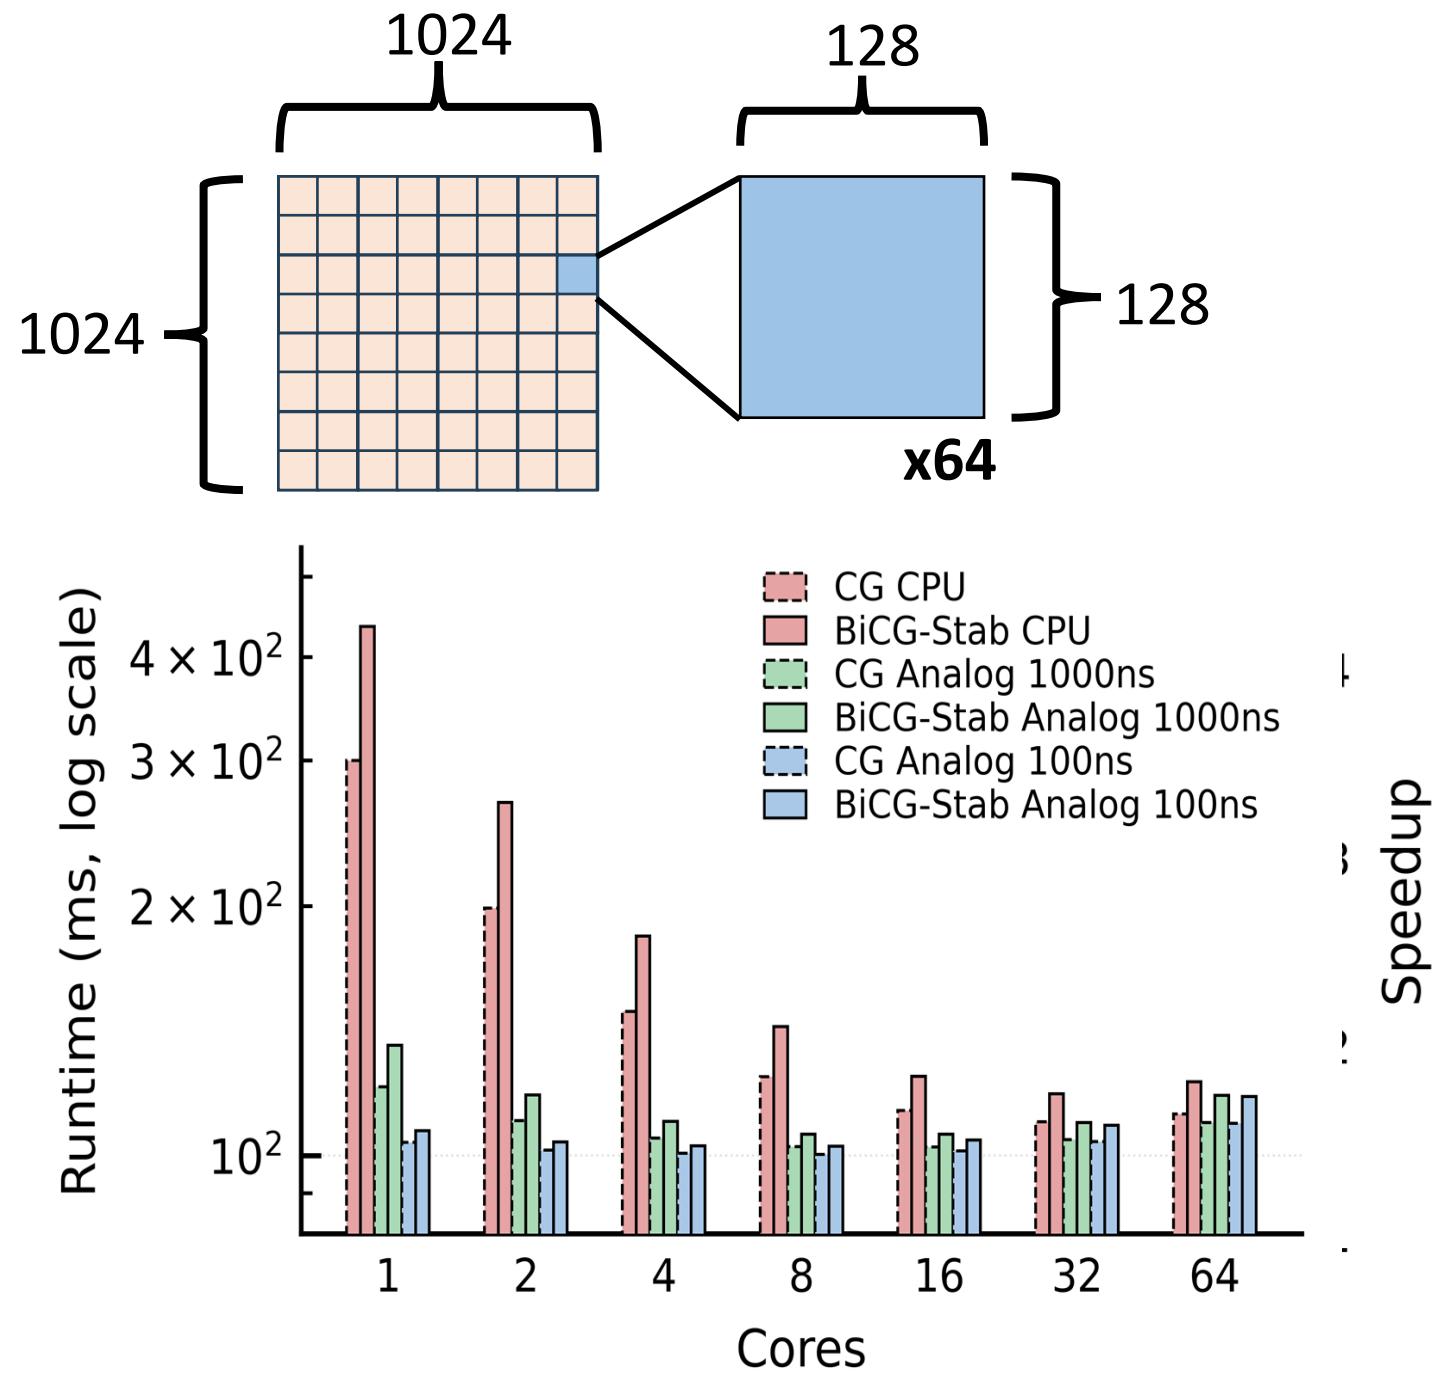
Evaluation

Algorithms

- Conjugate Gradient
- Stabilized Biconjugate Gradient

Setup

- **1024x1024 matrix** broken up into **64 128x128 submatrices**.
- Distribute 64 submatrices evenly across **1-64 cores**.
- Use **100ns** and **1000ns** latency for MVM operations.
- Comparing CPU runtime performance vs CPU + analog MVM runtime performance.



Conclusion + Future Work

Contributions

- First full system simulation of a tightly coupled **analog MVM + RISC-V architecture** for HPC-style workloads.
- A **minimal, composable ISA** enabling on-tile analog pipelines.
- **Reusable golem SST component** to explore analog/digital trade-offs and non-idealities.

Future Work

- Create an MLIR dialect to automatically partition workloads across multiple arrays. Lower PyTorch models to analog compatible binaries.
- Implement co-processor scratchpads and workflows.
- Experiment with precision of analog crossbar MVMs.
- Implement the RISC-V tile on FPGA hardware and couple it with an analog coprocessor for in-silicon validation.
- Measure energy usage and performance.

Questions?