

# RISC-V Vectorization Coverage for HPC: A TSVC-Based Analysis

RISCV-HPC @ SC25

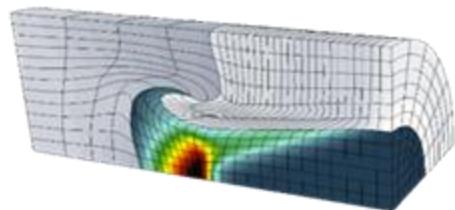
11/17/2025

Hung-Ming Lai<sup>1,4</sup>, Pei-Hung Lin<sup>1</sup>, Maya B. Gokhale<sup>1</sup>, Ivy Peng<sup>2</sup>, Hiren Patel<sup>3</sup>, Jenq-Kuen Lee<sup>4</sup>

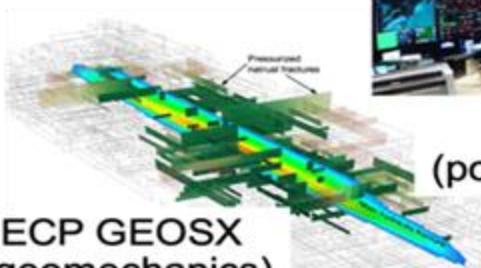
Center for Applied Scientific Computing LLNL<sup>1</sup>

KTH<sup>2</sup>, University of Waterloo<sup>3</sup>, NTHU<sup>4</sup>





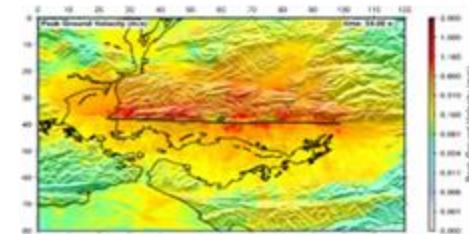
LLNL ECP/ATDM  
(high-order ALE hydro)



ECP GEOSX  
(geomechanics)



ECP ExaSGD  
(power grid optimization)



ECP SW4  
(earthquake modeling)

plus  
others...



Perlmutter  
Perlmutter (LBL)  
AMD Milan CPUs +  
NVIDIA Ampere GPUs



VANGUARD  
Astra  
Astra (SNL)  
ARM architecture



Sierra (LLNL)  
IBM P9 CPUs + NVIDIA Volta GPUs

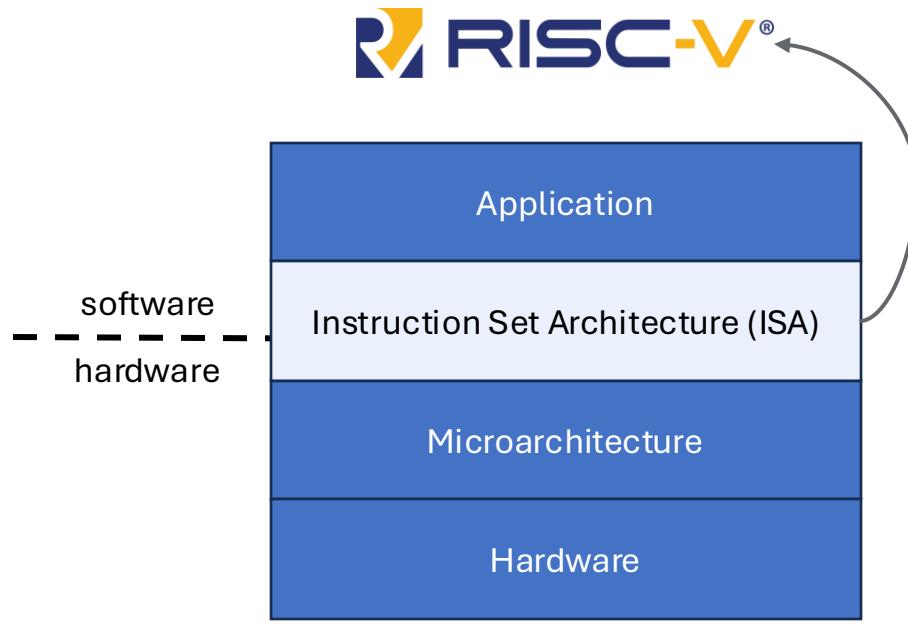


Aurora (ANL)  
Intel Xeon CPUs + Xe GPUs



Frontier (ORNL) &  
El Capitan (LLNL)  
AMD CPUs + GPUs

# The potential of RISC-V's open, modular design in next-gen HPC

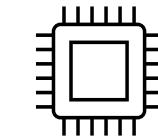


- Proprietary vs Open-standard ISA

	Proprietary (x86, Arm, ...)	Open-standard (RISC-V)
Licensing Cost	Expensive, restricted	Free and royalty-free
Customizability	Limited or restricted	Fully customizable
Vendor Lock-in	High	Low — broad vendor ecosystem
Innovation Speed	Slower, centralized	Fast, community-driven

- Modular ISA Extensions

RV32IMC



MCUs

RV64IMAFDC



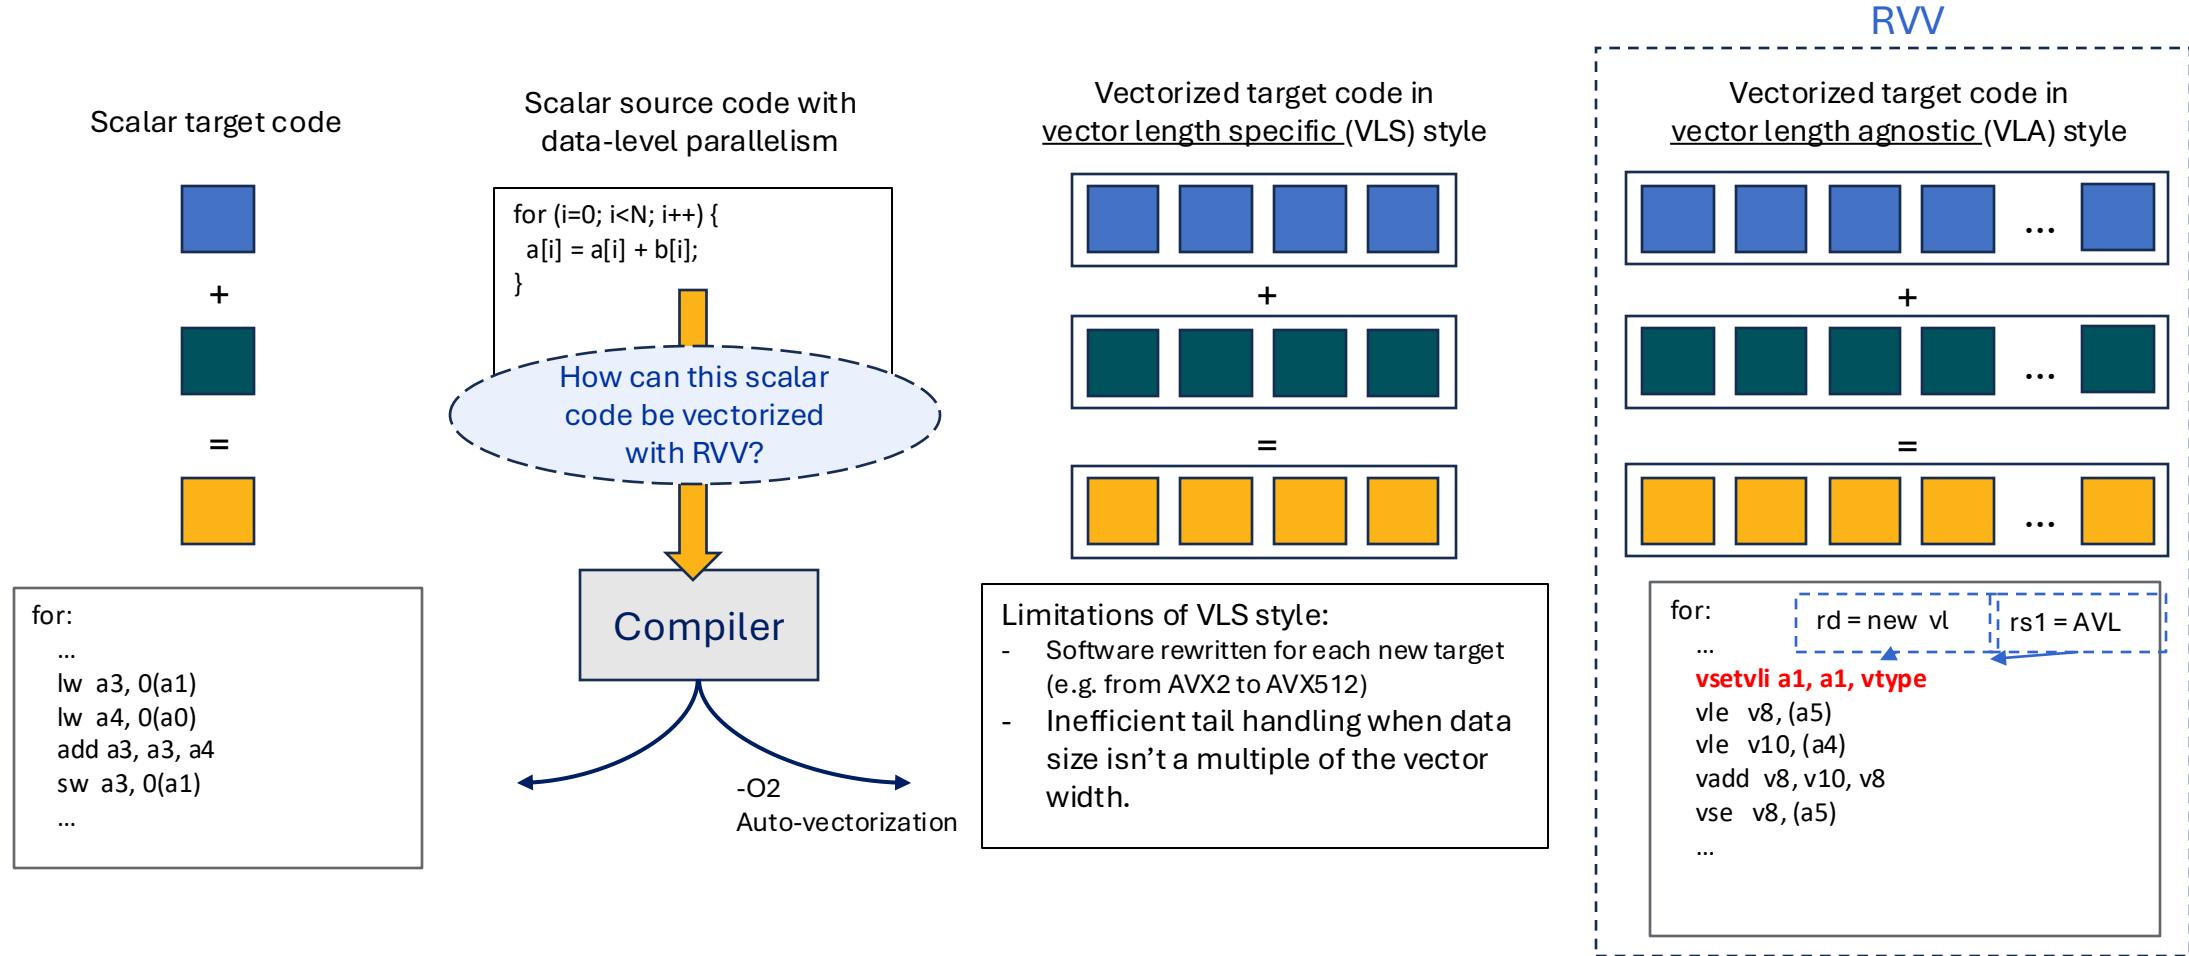
CPUs

RV64IMAFDCV



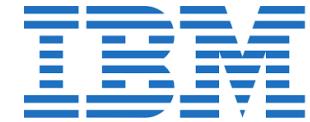
Accelerators

# Scalable vectorization with RISC-V vector extension and compiler auto-vectorization



# TSVC: Test Suite for Vectorizing Compilers

- A benchmark designed to evaluate compiler auto-vectorization capabilities.
  - Originated in the 80's
  - loops derived from real-world and synthetic codes
  - grouped by dependency and control-flow characteristics
  - Part of LLVM test suite
- Statistics
  - ~151 loop kernels across control flow, memory access, and arithmetic patterns
  - Covers scalar, array, pointer-based, and conditional loops
  - Used by research and industrial compiler teams to assess vectorization performance
- Applied in studies for various CPUs



# Methodology

## RVV Instruction Coverage Analysis

### Missed Instruction Identification

- Compiler backend or pattern-matching limitations
- Lack of representative use cases
- Non-trivial vectorization requirements or unsupported cases

### Future Benchmark Enhancements

# Key Results - RVV Coverage Statistics

Instruction group	config-setting	Load & Store	Integer	Fixed-Point	Floating-Point	Reduction	Mask	Permutation	Sum of all groups
RVV Instruction count	3	21	55	13	41	16	15	15	179
GNU-VLA	2 (66%)	10 (47.6%)	9 (16.4%)	0 (0%)	10 (24.4%)	1 (6.3%)	1 (6.7%)	8 (53.3%)	41 (22.9%)
GNU-VLS	2 (66%)	11 (52.4%)	9 (16.4%)	0 (0%)	12 (29.3%)	1 (6.3%)	1 (6.7%)	8 (53.3%)	44 (24.6%)
LLVM-VLA	2 (66%)	10 (47.6%)	10 (18.2%)	0 (0%)	17 (41.4%)	3 (18.8%)	1 (6.7%)	9 (60%)	52 (29.1%)
LLVM-VLS	2 (66%)	10 (47.6%)	10 (18.2%)	0 (0%)	17 (41.4%)	3 (18.8%)	1 (6.7%)	8 (53.3%)	51 (28.5%)

# Key Results - Detailed Observations

---

- Config-setting: Only vsetvli/vsetivli observed, not vsetvl
- Load/Store: Missed mask register, segmented, and advanced indexed instructions
- Integer: No widening instructions; limited by TSVC's floating-point focus
- Fixed-point: No coverage due to lack of kernels
- Floating-point: Missing widening and FMA instructions; compilers can generate with suitable patterns
- Reduction: Basic reductions covered, but widening and advanced forms missing
- Mask: Only vid observed; advanced mask ops absent
- Permutation: vslideup/vslidedown present; single-element and some merge instructions missing

# Key Results - VLA vs. VLS

---

- Minor differences between modes
  - GNU VLS: Uses `vlm` for mask loads, FMA instructions
  - LLVM: `vrgatherei16` in VLA, `vrgather` with 32-bit indices in VLS
- 
- Differences:
    - Load & Store: **`vlm`** employed in VLS by GNU
    - Floating-Point
      - GNU VLS: `vfmadd`, `vfmacc`
      - GNU VLA: `vfmul` and `vfadd`
    - Permutation
      - GNU VLA: `vrgatherei16`
      - GNU VLS: `vrgather`

# Key Results - Vector Length (VL) in VLS

---

- Coverage stable across vector lengths except for wide vectors (GNU emits `vmv.s.x` for  $VL \geq 512$ )
  - S4116: `sum += a[off] * aa[j -1][ ip[i]];`
  - Code generation for reductions under VLS differs based on vector length (VL):
    - For  $VL < 512$ : Uses scalar floating-point register for accumulation, seeded with `vfmv.s.f` and updated via `vfmv.f.s.`
    - For  $VL \geq 512$ : Keeps accumulator in vector registers, initialized with `vmv.s.x` and reduced with `vfredosum.vs`, avoiding scalar register traffic.

# Classification of Auto-Vectorization Gaps

---

## Key Failure Causes:

- Unidentified PHI
  - cyclic data dependencies between instructions
  - non-predictable control-flow variables
- Memory Conflict
  - conservatively assumed aliasing
  - Overlapping memory access patterns
- Switch Statement
- Non-intrinsic Call
- Early Exit

# Failure Causes

Cause (# of loops)	loop
Unidentified PHI (37)	s116, s1213, s123, s126, s13110, s141, s161, s211, s2111, s212, s221, s222, s2233, s2251, s231, s232, s233, s235, s242, s256, s258, s261, s275, s277, s291, s292, s3110, s3112, s315, s318, s321, s322, s323, s331, s341, s342, s343
Memory conflict (7)	s1113, s114, s1161, s1244, s241, s244, s281
Switch statement (1)	s442
Non-intrinsic call (1)	va
Early exit (3)	S332, s481, s482

**First failure cause encountered by the LLVM LoopVectorize pass for each non-vectorized loop in TSVC**

# Failure Cause: Unidentified PHI (cont'd)

- cyclic data dependencies between instructions
- non-predictable control-flow variables

```
// s221
for (int i = 1; i < LEN_1D; i++) {
    a[i] += c[i] * d[i];
    b[i] = b[i - 1] + a[i] + d[i];
}
```

LV: Not vectorizing: Found an unidentified PHI  
%0 = phi float [ %.pre, %for.cond2.preheader ], [ %6,  
%for.body5 ]

control-flow pattern  
not recognized

- non-predictable control-flow variables

```
// s161
for (int i = 0; i < LEN_1D-1; ++i) {
    if (b[i] < (real_t)0.) {
        goto L20;
    }
    a[i] = c[i] + d[i] * e[i];
    goto L10;
L20:
    c[i+1] = a[i] + d[i] * d[i];
L10:
    ;
}
```

LV: Not vectorizing: Found an unidentified PHI  
%indvars.iv = phi i64 [ 0, %for.cond2.preheader ], [  
%indvars.iv.next.pre-phi, %for.inc ]

# Examples of Other Failure Causes

- Memory conflict

```
// s1113
for (int i = 0; i < LEN_1D; i++) {
    a[i] = a[LEN_1D/2] + b[i];
}
```

- Switch statement

```
// s442
for (int nl = 0; nl < iterations/2; nl++) {
    for (int i = 0; i < LEN_1D; i++) {
        switch (indx[i]) {
            case 1: goto L15;
            case 2: goto L20;
            case 3: goto L30;
            case 4: goto L40;
        }
    L15:
    ...
}
```

- Non-intrinsic call

```
// va
for (int i = 0; i < LEN_1D; i++) {
    a[i] = b[i];
}
```

- Early exit

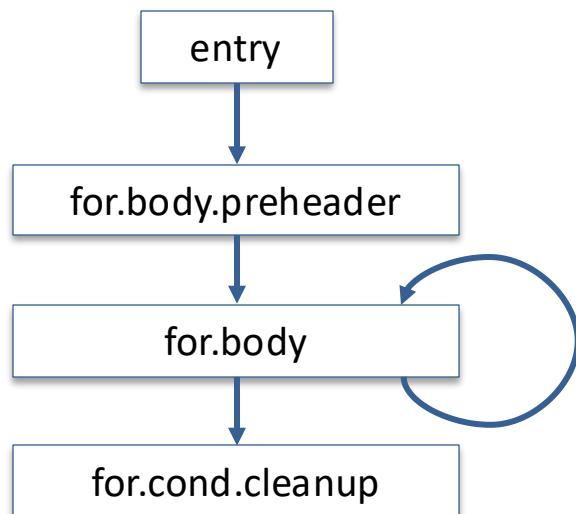
```
// s482
for (int i = 0; i < LEN_1D; i++) {
    a[i] += b[i] * c[i];
    if (c[i] > b[i]) break;
}
```

whole loop optimized to  
@llvm.memcpy by previous passes

# Failure Cause: Unidentified PHI

- SSA PHI nodes from a loop

```
void foo(int *a, int *b, int n) {  
    for (int i = 1; i < n; i++) {  
        a[i] = a[i - 1] + b[i];  
    }  
}
```



```
entry:  
    %cmp10 = icmp sgt i32 %n, 1  
    br i1 %cmp10, label %for.body.preheader, label %for.cond.cleanup  
  
for.body.preheader:  
    %wide.trip.count = zext nneg i32 %n to i64  
    %load_initial = load i32, ptr %a, align 4  
    br label %for.body  
  
for.cond.cleanup:  
    ret void  
  
for.body:  
    %store_forwarded = phi i32 [ %load_initial, %for.body.preheader ], [ %add, %for.body ]  
    %indvars.iv = phi i64 [ 1, %for.body.preheader ], [ %indvars.iv.next, %for.body ]  
    %0 = getelementptr i32, ptr %a, i64 %indvars.iv  
    %arrayidx2 = getelementptr inbounds i32, ptr %b, i64 %indvars.iv  
    %1 = load i32, ptr %arrayidx2, align 4  
    %add = add nsw i32 %1, %store_forwarded  
    store i32 %add, ptr %0, align 4  
    %indvars.iv.next = add nuw nsw i64 %indvars.iv, 1  
    %exitcond.not = icmp eq i64 %indvars.iv.next, %wide.trip.count  
    br i1 %exitcond.not, label %for.cond.cleanup, label %for.body
```

compiled LLVM IR

# Failure Cause: Unidentified PHI (cont'd)

- “Identified” PHIs
  - induction PHI
    - induction scalar → value of  $i$
    - induction pointer → pointer of  $a[i]$  and  $b[i]$
  - reduction PHI → value of sum
  - first-order recurrence (PHI from optimized load with no actual data dependence) → value of  $a[i - 1]$

```
for (int i = 1; i < n; ++i) {  
    sum += a[i];  
    b[i] = a[i] - a[i - 1];  
}
```

# Auto-Vectorization Gaps: Root Causes and Implications

---

## 1st Core Barrier:

- Control Flow Complexity
  - Divergent paths (e.g., branches, switches, early exits) prevent lockstep parallel execution.
  - Modern ISAs (e.g., RVV) mitigate via masking/predication.
  - Compilers can apply if-conversion, but success depends on control simplicity.
  - Complex cases may require speculation or explicit restructuring.

# Auto-Vectorization Gaps: Root Causes and Implications

---

## 2nd Core Barrier:

- Data Dependence Ambiguity
  - Must ensure vectorization doesn't violate loop-carried dependencies.
  - Failures like Unidentified PHI and Memory Conflict stem from:
    - Unclear SSA dependencies
    - Insufficient aliasing analysis
  - Often not real barriers, but compiler conservatism.
  - Mitigation techniques:
    - Loop interchange – reorder iterations
    - Loop distribution – isolate vectorizable sections
    - Prefix-sum transformations – handle recurrences

# Missed Instruction Analysis in TSVC

- “not vectorizable” due to a loop-carried dependency on  $j$
  - $j$  depends on how many earlier iterations satisfy  $c[i] > 0$
  - This creates a scalar dependency chain  
→ sequential updates only
  - Conventional SIMD auto-vectorizers cannot break this dependency
  - TSVC loop s123
- ```
int j = -1;
for (int i = 0; i < (LEN_1D/2); i++) {
    j++;
    a[j] = b[i] + d[i] * e[i];
    if (c[i] > (real_t)0.) {
        j++;
        a[j] = c[i] + d[i] * e[i];
    }
}
```

# RVV Solution: Masked Prefix-Sum Instructions

- RISC-V Vector Extension (RVV) introduces hardware support for per-lane index computation:

| Instruction | Function                    | Purpose                               |
|-------------|-----------------------------|---------------------------------------|
| viota.m     | masked exclusive prefix sum | gives each active lane its own offset |
| vcpop.m     | population count of mask    | advances base index for next vector   |

- Mask:  $m = (c > 0)$
- Offsets:  $vIO = \text{viota}.m(m)$
- Base advance:  $j\_base += vl + \text{vcpop}.m(m)$

# Vectorized Form and Compiler Insights

- Recognize prefix-sum recurrence on induction variables.
- Apply if-conversion + masked index lowering.
- Extend LLVM/GCC to map these patterns to `viota.m` / `vcpop.m`.

- Vectorized pseudocode (simplified):

```
while (p < N/2) {  
    vl = vsetvl(N/2 - p);  
    vB = vle(b+p); vC = vle(c+p);  
    m = vgt(vC, 0);  
    vIO = viota.m(m);  
    idx_b = j_base + vid() + vIO;  
    vsuxei(a, idx_b, vB + d*e);  
    idx_c = idx_b + 1;  
    vsuxei.m(a, idx_c, vC + d*e, m);  
    j_base += vl + vcpop.m(m);  
    p += vl;  
}
```

# Conclusions & Future Directions

---

- TSVC covers basic RVV instructions, but many advanced features missed
- Benchmark ambiguity complicates evaluation
- Recommend:
  - Extending TSVC for modern vector ISAs
  - Clearer vectorizability criteria
  - Benchmarks inspired by RAJA, Kokkos
- Systematic assessment framework needed for compiler and architecture co-design



Lawrence Livermore  
National Laboratory