



Cuzco: From Open-Source to High-Performance RISC-V CPU IP

International Workshop on RISC-V for HPC (RISCVHPC)
SC25, St. Louis, November 17, 2025

Shashank Nemawarkar, Senior Director of Architecture

Intro to Condor Computing

- Condor Computing, a wholly owned USA subsidiary of Andes Technology, was founded in 2023 with the goal of creating the highest performance, licensable RISC-V CPU IP in the industry
- We are a tight knit team (~50 engineers), with very light management and overhead, entirely focused on bringing an innovative new micro-architecture to the RISC-V CPU market
- We intend to demonstrate that RISC-V can be competitive in any high-performance computing application, from datacenters to handsets, and up to automotive

Cuzco Processor Summary

- Licensable O-O-O CPU core IP designed for highest end performance application processors
- Innovative, time-based scheduling to eliminate Tomasulo algorithm to achieve similar performance much less power than comparable cores
- Support for up to 8 cores with private L2\$ in a coherent cluster with a shared L3\$
- Latest RISC-V profile support (RVA23) for maximum software compatibility
- Full support for ISA customization

RISC-V Ecosystem: SIG Performance Modeling


- Charter:
“...address our community's need for cycle-accurate simulation and aims to establish a common workflow and toolscape to use through the product lifecycle...”
- Interface with other SIGs for common topics
- Some participants: Condor, MIPS, Ventana, SiFive, TensTorrent, Imagination, and individuals
- Confluence: <https://riscv.atlassian.net/browse/RVG-9>
- Github: <https://github.com/riscv-software-src/riscv-perf-model>
- Webpage: <https://lists.riscv.org/g/sig-perf-modeling>

Platform and Modeling Framework

- Provides a basic processor model
 - Enhance per needs
 - Sample traces
 - Statistics* and Visualization (Helios)

Condor team contributions:

- Sparta (multiple “original” members)
- Mavis (“original” members)
- Trace framework enhancements
- Fusion Specification Language
- Branch prediction (*TAGE adaptation)
- Creates democratization of high-performance compute research and product design with high quality



RISC-V in HPC and AI/ML: Ecosystem

RV Accelerators

- Clusters: large number of cores
 - Meta (MTIA), Nvidia,
 - Startups: InspireSemi, Calligo Tech
- RISC-V CPUs and/or GPUs
- RISC-V cores as command and control processors

Advantages

- Custom area, power efficient cores
- Standardized ISA extensions
- Custom functions/instructions
- Ecosystem for software stack (RISE)

SIG-Perf-Modeling Covers Single core systems

- Detailed modeling of single core, its memory system and finer components
- Stats generation of finer components, workload analysis (stf, stalls, visuals)

SIG-Perf-Modeling extensions needed for HPC

- Extend open source support from CPU level modeling and analysis to the system
- Performance and power modeling
- Memory system, on-chip/off-chip network design trade-offs
- Plug-and-play modules for the same
- Coherent and non-coherent protocol based analysis

Thanks: Nick Brown, RISC-V Summit North America, October 2025

FSL: Fusion/Fracture Specification Language

- FSL expressions define instruction transforms
 - FRACTURE A **one-to-many** instruction transformation
 - FUSION A **many-to-one** instruction transformation
 - BINARY General binary translation
- FSL provides a language, transform toolchain and C++ API
 - A Python interface is built from the shared library
- FSL based tools can generate RTL, perf model methods and compiler MDL
 - RTL (System Verilog): decoders, trace/uop structures, unit computation sites, etc.
 - Perf model methods (C++): documented uses are shown through Sparta/Olympia
 - Compiler output (LLVM): TableGen fusion patterns, predicates, scheduling hooks, etc.
- FSL is ISA agnostic
 - Current examples and documentation focus on RISC-V

FSL: Transform Syntax Example

FSL processes instructions windows in three phases: each phase has a syntactical expression in the grammar, a clause

1. Identify instruction **sequence**; 2. apply **constraints** to filter tuples; and 3. perform **transformation** to map to new instructions

```
transform uf10
{
    // Prolog elements
    isa rv64g      // ISA definition object
    uarch oly1      // uArch definition object
    ioput iop1      // API interface object

    // Variables at transform scope
    gpr g1,g2,g3,g4
    u5 c1
    s12 c2

    // Abstracted instruction sequence
    sequence seq_uf10 {
        sd g1,c1(g2)
        sd g3,c2(g4)
    }

    // continued ->
}

// -> continued

// Example of a constraint specification
constraints cns_uf10 {
    g1 != g2
    g3 != g4
    g1 != g4
}


// Conversion clause using abstract morphing
conversion cnv_uf10 {

    // merge sequence into 1 object
    instr instr_uf10.morph(seq_uf1)


    // insert transform into pipeline
    iop1.input.replace(seq_uf10,instr_uf1)
}
```

Cuzco Feature Overview

- RVA23 profile compliant, with Hypervisor
 - 64-bit, RV64GCBKV + CMO
- Innovative time-based microarchitecture
- 12-Stage Pipeline
- 8-Wide Frontend Decode
- 256-Entry Reorder Buffer (ROB)
- 8 Execution Pipelines
- RISC-V Vector 1.0 + Crypto, 256/512b VLEN
- Branch Target Buffer-way predicted (BTB)
- TAGE-SC-L & Tournament Branch Predictor
- 1K/2K/4K 4-Way L2 TLBs
- 64 KB, 8-Way Private I/D Caches
- Up to 8 MB, Private L2\$
- SECDED ECC error protection
- 8-Core Multiprocessor w/ Shared L3\$ (up to 256MB)
- 256/512-bit CHI and 64b/512b MMIO Buses



Cuzco CPU Core Block Diagram



Cuzco: Vector Everywhere

O-O-O Scalar/Vector:

- Compilers making more use of RVV:
 - vector/scalar intermingled code
- Gains from "vector everywhere":
 - frequent, tight lower latency interactions and power-conscious resource sharing
- Unlike traditional wide-vector, decoupled designs:
 - no need for dense, pure vector library code to be performant and cost effective

Excels at:

- Capturing enhanced data bandwidth opportunities general purpose applications
- AI transformer inferencing (quantization, dot product and matmul)
- Bursts of parallel element computation with scalar/vector ILP

Spec2K17, XZ benchmark:

```
...
1ed78: sub    x9, x31, x8
1ed7c: add    x22, x21, x8
1ed80: add    x23, x12, x8
1ed84: vsetvli x24, x9, e8, m2, ta, ma
1ed88: vle8.v  v8, (x22)
1ed8c: vle8.v  v10, (x23)
1ed90: vmsne.vv v12, v8, v10
1ed94: vfist.m x23, v12
1ed98: c.mv   x22, x24
1ed9a: blt    x23, x0, 0x1eda0 <bt_find_func+0xd4>
...
```


[bt_find_func\(\)](#): a byte-by-byte string comparison to find matches in a binary tree

Cuzco Core Pipeline

IF0	IF1	IF2	IF3	ID0	ID1	ID2	EX0	EX1/DC1	DC2	DC3	DC4
Calculate next address & access iTLB	Access IC tag array, IC(hit/miss) BTB, and GHT, TAGE-SC-L	Access IC data array access, BTB hit/miss and taken/non-taken	Write IC cache line to ICQ & bypass/ read N instructions to XIQ	Read N instrs from XIQ, 1 st instr decode, & access RFL & RAT	Access RSB to calculate execution times, for each instruction, and create dependence chains	Access TRM to issue instruction, secondary instruction decode, write issued instructions to XEQ	Read RF/fwd data, check RSB, send instr from XEQ to functional unit	Execute instrs, write result to PRF. AGU and access dTLB	Access DC tag array (DC hit/miss)	Access DC data array	Align and send load data to write back to PRF

Cuzco Time-Based Microarchitecture

- 1st CPU designed with hardware compilation for optimal instruction sequencing
- Schedule instruction execution with a perfect view of all past instructions
 - Register Scoreboard: record write time of an instruction to a register which becomes the read time of dependent instruction
 - Time Resource Matrix (TRM): busy indicators for read & write buses, and other resources
 - Issues instructions with precise predicted future times for execution
 - Reschedule/replay: Account for dynamic latencies, resource conflicts
- Reduces scheduling complexity, timing, area and power of typical OOO cpu mid-cores




Time Resource Matrix (TRM): Performance Model

Time	Scheduled	Read ports			Write ports		Br	ALU		LSU
83	O	P	A	A	J					
84	Q	R	L			J		A	J	
85	S	T	N			M		M		L
86	U	V	H	I	T	N		N		
87	W	X	K	K		H	I	H	I	T
88	Y		B	B		K		K		
89	a	b	C	C	D	D	B		B	
90	c	d	E	E	F	F	D	C	D	
91	e	f	G	G	P		F	E	F	
92	g	h	U					G		P
93	i	j	O	O	V	c				U
94	k	l	X			O	c	O	c	V

- Instructions are efficiently scheduled in the TRM to take advantage of available resources
- In this example snippet from Dhrystone, instruction O is scheduled in cycle 83, issued in cycle 93, and executed in cycle 94


Time-Based Microarchitecture: Design Selection

- Build execution schedule with known & projected dependencies and latencies
 - Activate operands, resources, and execution units only at scheduled times
 - No need to search or prioritize selection at reservation stations
 - Reduced scheduler complexity (logic & area)
 - Dynamic power reduction vs conventional scheduler
- Two instructions scheduled/slice/cycle, callable with slices
 - TRM2: Identify functional unit slots within 2 cycles of operand availability
 - TRM8: Within the next 8 cycles
 - Ideal: Idealized scheduling




Slice-based Microarchitecture: Config Selection

- Baseline: 4 slices, 2 instruction pipelines/slice
- Uniformly scalable design IP
 - Each slice implements a fully compatible RISC-V CPU
 - Each slice adds symmetric set of resources to the machine
- Static PPA control through IP configuration
- Dynamic power vs. throughput control with slice-based enables and scheduling



Future Improvements

- Cycle-accurate performance model allows projecting impact of features implementable in the next generation
- This chart shows the impact of successively applying the first 5 of these features



Comparison with Andes O-O-O Processor

CPU Cores	Andes AX65	Cuzco
Pipeline stages	13	12
Issue width	4	8
ROB	128	256
Branch prediction	TAGE-L	TAGE-SC-L-IT
Int ALU & FPU Load/Store Unit	4 & 2 2 L/S	8 & 4 4 L/S
L1 Caches	I\$: 64KB D\$: 64KB	I\$: 64KB D\$: 64KB
L2 Cache	8MB Shared	Up to 8MB Private
L3 Cache (Shared)	None	Up to 256MB
DMIPS/MHz	4.82	8.5
SPECint2k6/GHz	8.78*	>17.5*

*with 8MB L2 and no vector ISA usage

Cuzco Memory System Microarchitecture

Private L1/L2 Caches

- L1: 64KB I\$/D\$, 8-way
- I\$/D\$ L1/L2 prefetch and D\$ writearound
- L2: Up to 8MB, up to 16-way
- L2: Configurable multi-cycle SRAM accesses
- Configurable up to 64 outstanding requests

Privilege Modes

- Machine (M), hypervisor (H), supervisor (S) and user (U)

Memory Management Unit (MMU)


- Bare, Sv39, Sv48, and Sv57 VA

Shared L3 Cache, Coherent Cluster

- Up to 256MB, up to 16-way,
- Up to 8 cores+L2\$s in a cluster

Bus Interfaces (synchronous and asynchronous)

- 512-bit main memory CHI bus interface
- 256-bit memory mapped I/O (MMIO) interface



Cuzco in RISC-V Ecosystem

- **With RISC-V SIG Perf Modeling: open source tools/infra**
 - From scratch to full fledged cycle-accurate model and tools significantly impacting design within 2 years
 - Contributions back to open-source tools/infra
- **Cuzco: Industry standard RISC-V ISA Compliant Design**
 - Better performance / \$
 - Better performance / μW of power
 - Fully MP-ready: up to 8 HPC CPU cores per cluster
 - RVA23 profile
 - Extensible
- **Approach also used by multiple RISC-V teams**

Acknowledgments

The team at Condor Computing wishes to thank:

- Andes Technology Corporation, for their support:
 - Frankwell Lin, Chairman & CEO
 - Dr. Charlie Su, CTO & President
 - Emerson Hsiao, President Andes Technology USA
- Dr. Thang Tran, for his fundamental contributions to the invention of the time-based scheduling microarchitecture



Thank you!

Shashank Nemawarkar

shashank.nemawarkar@condorcomputing.com
www.condorcomputing.com

FSL: Processing and Syntax

- FSL processes instructions windows in three phases
 - Sequence, constraints and transform
- Each phase has a syntactical expression in the grammar, a clause
 - **Sequence clause**
 - Action: Identifies instruction sequences matching specified attributes
 - Eg: ADD/ADD or ADD/ADD/SUB or LD/ST/BR, etc., wild cards are supported
 - **Constraints clause**
 - Action: filters sequences for transformable tuples, returns sequences
 - Eg: Operand limits (PRF ports), machine state constraints, e.g. busy/available computation sites
 - **Transform clause**
 - Action: map instructions to new instruction(s), adjust operands, insert transform into pipe
 - Eg: Fuse ADD/ADD/SUB into MAGIC_AAS, adjust operands, and dispatch to the correct execution queue

FSL: Transform Syntax Example

```
transform uf10
{
    // Prolog elements
    isa    rv64g      // ISA definition object
    uarch oly1       // uArch definition object
    ioput iop1       // API interface object

    // variables at transform scope
    gpr  g1,g2,g3,g4
    u5   c1
    s12  c2

    // Abstracted instruction sequence
    sequence seq_uf10 {
        sd  g1,c1(g2)
        sd  g3,c2(g4)
    }

    // continued ->
}

// -> continued

// Example of a constraint specification
constraints cns_uf10 {
    g1 != g2
    g3 != g4
    g1 != g4
}

// Conversion clause using abstract morphing
conversion cnv_uf10 {

    // merge sequence into 1 object
    instr instr_uf10.morph(seq_uf1)

    // insert transform into pipeline
    iop1.input.replace(seq_uf10,instr_uf10)
}
```

FSL: More Information

- **We continue to develop FSL and utilities for analysis and generation**
 - STF and Simpoint based instruction stream mining utilities
 - Automation for performance model stats to fusion/fracture candidates expressed in FSL
 - Application examples for RISC-V vector ISA
 - Compiler build automation and output comparison utilities
 - Interaction/integration with other open-source analysis tools
- **FSL available at the Condor Computing git repo**
 - <https://github.com/condorcomputing/fsl>
- **FSL in RISC-V ecosystem**
 - Olympia usage example <https://github.com/riscv-software-src/riscv-perf-model>
 - FSL vs Python: <https://github.com/riscv-software-src/riscv-perf-model/discussions/121>

Cuzco Memory System Microarchitecture

Private L1 Caches


- 64KB I\$/D\$, 64B line size, 8-way, pLRU
- PIPT (Physical Index and Physical Tag)
- 64-byte cache line size
- I\$/D\$ prefetch and D\$ writearound
- SECDED ECC error protection
- Up to 64 pending miss requests
- 4 cycle load->use penalty

Private L2 Cache

- Up to 8MB, 64B line size, up to 16-way, pseudo-random replacement
- I/D prefetch
 - Preset and configurable prefetch policies
- Configurable multi-cycle SRAM accesses
- SECDED ECC error protection
- +14 cycle delay on L2 hit

Privilege and Memory Management

- Machine (M), hypervisor (H), supervisor (S) and user (U) privilege modes
- Memory management unit (MMU)
 - Bare, Sv39, Sv48, and Sv57 VA translations
 - Svnapot, Svpbmt, Svinval VM extensions
 - L1 I/D TLBs: 64-entry, fully associative
 - L2 TLB: up to 4K-entry, 4-way
 - PMP and ePMP support with 16 PMP entries
- 16 PMA regions



Cuzco: Cluster Microarchitecture

Shared L3 Cache


- Up to 256MB, 64B line size, up to 16-way, pseudo-random replacement
- I & D prefetch, configurable policies
- Max outstanding reads/writes (cacheable and uncachable): 32-128
- Max 64 outstanding snoop transactions
- Configurable multi-cycle SRAM accesses
- SECDED ECC error protection

Cluster with Multicore Cache Coherency

- Up to 8 cores+L2\$ in a cluster
- Coherence Manager and L3\$

Bus Interfaces

- 512-bit main memory CHI bus interface
- 256-bit memory mapped I/O (MMIO) interface
- Core+L2 vs. external-bus clock
- Asynchronous, and Synchronous N:1 clock ratios



Analysis Example: Identifying the Workload

Identify the workload: feature **ir1**:

xalancbmk (SpecInt2K6) and **x264_s (SpecInt2K17)**

Benchmark	ir1	Benchmark	ir1
si06	-0.22%	si17	-0.26%
astar	-0.00%	deepsjeng_s	-0.00%
bzip2	-0.22%	exchange2_s	0.02%
gcc	0.20%	gcc_s	-0.13%
gobmk	-0.16%	leela_s	-0.02%
h264ref	-0.05%	mcf_s	-0.00%
hmmer	-0.14%	omnetpp_s	-0.03%
libquantum	-0.33%	perlbench_s	-0.02%
mcf	0.01%	x264_s	-2.29%
omnetpp	-0.02%	xalancbmk_s	-0.14%
perlbench	0.04%	xz_s	0.02%
sjeng	-0.02%		
xalancbmk	-1.96%		

One workload each in
SpecInt2K6 and SpecInt2K17

Benchmark	ir1
xalancbmk	-1.96%
0001558	-8.48%
0001559	-0.03%
0001560	-18.78%
0001561	0.01%
0001562	-0.88%
0001563	-0.56%
0001564	7.93%
0001565	-3.76%
0001566	0.02%
0001567	-0.04%
0001568	-1.03%
0001569	-0.15%
0001570	-0.11%
0001571	0.05%
0001572	-2.12%
0001573	-0.07%
0001574	-0.24%

Single outlier trace
1560 in xalancbmk

Look for stalls at interfaces:

Front end, mid-core, LS, MemSys

stats	% diff
-----	-----
rob.ipc	-18.78%
fetch.fet_stall_pct_not_stalled	-8.00%
fetch.fet_stall_pct_no_xiq_credits	-11.39%
fetch.fet_stall_pct_fg_br_mispred	19.12%
fetch.pred_ind_correct	-37.86%
fetch.pred_ind_incorrect	17701300.00%

- Front end stall analysis
- ipc -19% drop,
- indirect br mispred +177K times!

- Feature “ir1” is midcore design alternative is unrelated to branches or predictions
- A perfect indirect predictor recovered the ipc loss completely

Analysis Example Part I: Identifying the Workload

- We see that feature **ir1** has small negative impact overall (**si06** and **si17** lines), but this is almost entirely due to two benchmarks, **xalancbmk** and **x264_s**
- We examine the SimPointed traces that make up **xalancbmk** and see that there is a single trace, **1560**, which is a large outlier.
 - When we examine **x264_s** traces (not shown here), we see that the performance impact is more broad-based and affects all traces in this benchmark
- We therefore start by examining trace **1560**

Benchmark	ir1	Benchmark	ir1
si06	-0.22%	si17	-0.26%
astar	-0.00%	deepsjeng_s	-0.00%
bzip2	-0.22%	exchange2_s	0.02%
gcc	0.20%	gcc_s	-0.13%
gobmk	-0.16%	leela_s	-0.02%
h264ref	-0.05%	mcf_s	-0.00%
hmmer	-0.14%	omnetpp_s	-0.03%
libquantum	-0.33%	perlbench_s	-0.02%
mcf	0.01%	x264_s	-2.29%
omnetpp	-0.02%	xalancbmk_s	-0.14%
perlbench	0.04%	xz_s	0.02%
sjeng	-0.02%		
xalancbmk	-1.96%		

Benchmark	ir1
xalancbmk	-1.96%
0001558	-8.48%
0001559	-0.03%
0001560	-18.78%
0001561	0.01%
0001562	-0.88%
0001563	-0.56%
0001564	7.93%
0001565	-3.76%
0001566	0.02%
0001567	-0.04%
0001568	-1.03%
0001569	-0.15%
0001570	-0.11%
0001571	0.05%
0001572	-2.12%
0001573	-0.07%
0001574	-0.24%

Analysis Example Part 2: Looking at Stalls

- For trace 1560, we compare the stats between the baseline and the feature in question
- We confirm that the IPC has gone down nearly 19% for this trace.
- We first look at the fetch stall statistics, which shows how the fetch slots are distributed
 - **not_stalled** – this is the number of fetch slots that are not stalled (used for sending instructions to the midcore)
 - The higher this number, the better
 - **no_xiq_credits** – this is the number of fetch slots that are stalled due backpressure from the midcore
 - Note that midcore backpressure has gone down 11%
 - **fg_br_mispred** – this is the number of fetch slots that are stalled due to branch mispredicts
 - The stalls due to branch mispredicts has gone up 19%, which is surprising because our feature should not affect branch prediction rates
- We then examine our prediction rates. We see that the number of indirect target predictions has gone through the roof. We have almost no indirect target mispredicts in the baseline, but a significant number when the feature is enabled
- Conclusions:
 - We must investigate further to determine the source of instability in the modeled indirect target predictor, which is unrelated to our feature
 - In the meantime, we experiment by setting indirect target prediction to **perfect** in the model. We see that the **xalancbmk** performance difference has disappeared, although **x264_s** is unaffected and must still be analyzed

stats	% diff
rob.ipc	-18.78%
fetch.fet_stall_pct_not_stalled	-8.00%
fetch.fet_stall_pct_no_xiq_credits	-11.39%
fetch.fet_stall_pct_fg_br_mispred	19.12%
fetch.pred_ind_correct	-37.86%
fetch.pred_ind_incorrect	17701300.00%



Analysis Example Part 3: Other Stalls

- In a different example, we analyze a 67% drop in IPC
 - Used fetch slots has gone down, but this is almost entirely due to downstream backpressure
 - fet_stall_pct_no_xiq_credits +23%
 - We see that used rename slots has gone down, almost entirely due to backpressure from the LSU
 - ren_stall_pct_no_lsu_ldb_credits +22%
 - Downstream of rename, we see that the used dispatch slots has gone down, but it's primarily due to upstream stalls (from rename)
 - dp_stall_pct_d1q_empty +32%
- Our problem is that we are mis-scheduling instructions, filling up the execution units with poorly-scheduled instructions that must be rescheduled
 - This is confirmed by the number of TRM reschedules
 - num_trm_rescheduled +1745%
 - Our baseline configuration avoids this problem

stats	% diff
-----	-----
rob.ipc	-67.22%
fetch.fet_stall_pct_not_stalled	-21.67%
fetch.fet_stall_pct_no_xiq_credits	22.90%
rename.ren_stall_pct_not_stalled	-21.67%
rename.ren_stall_pct_no_lsu_ldb_credits	22.35%
sl_mgr.dp_stall_pct_not_stalled	-21.67%
sl_mgr.dp_stall_pct_trmc_busy_overflow	-23.17%
sl_mgr.dp_stall_pct_alu_xeq_full	10.98%
sl_mgr.dp_stall_pct_br_xeq_full	10.75%
sl_mgr.dp_stall_pct_d1q_empty	32.15%
sl_mgr.slice0.num_trm_rescheduled	1745.00%